

Anatomia de um Garbage Collector Moderno



Renê Cardozo

- Formado em Ciência da Computação - USP
- 6 anos desenvolvendo em Go, Java e no que tiver.



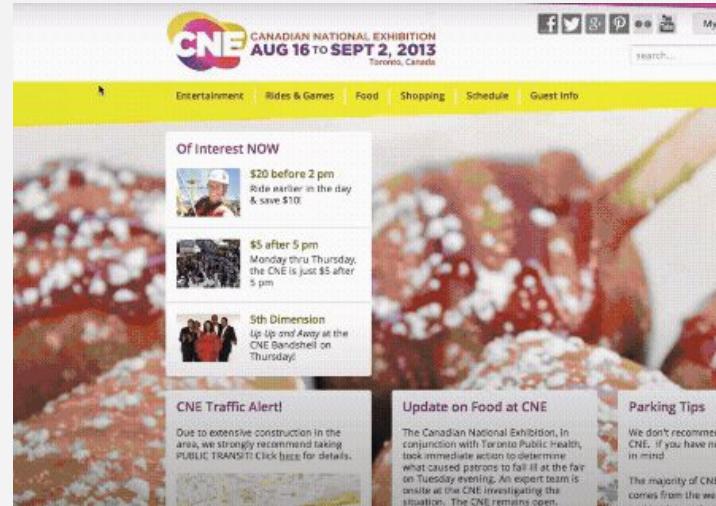
@reneepc



Jank

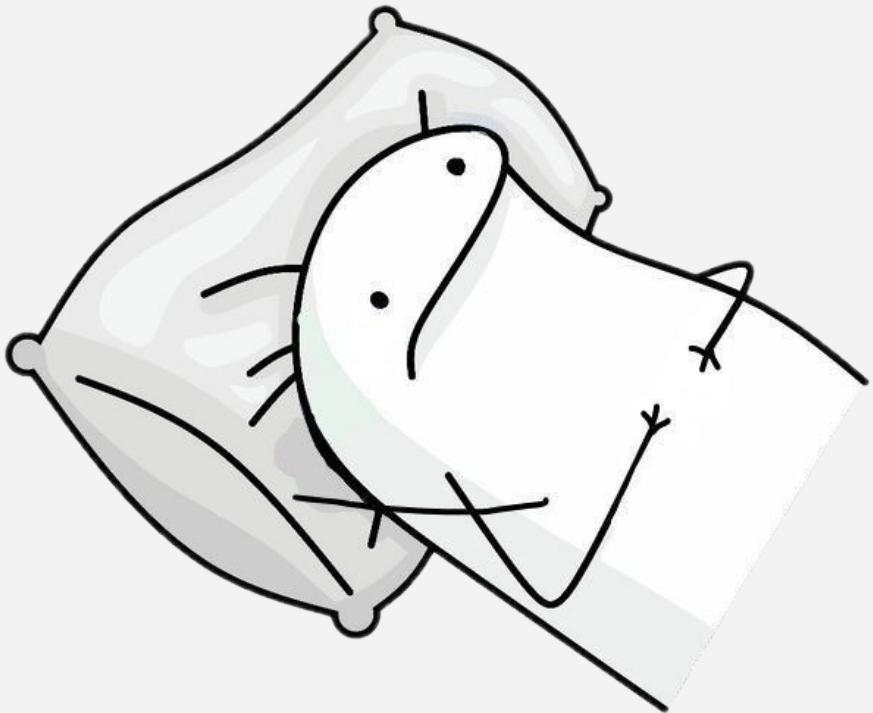


Janky



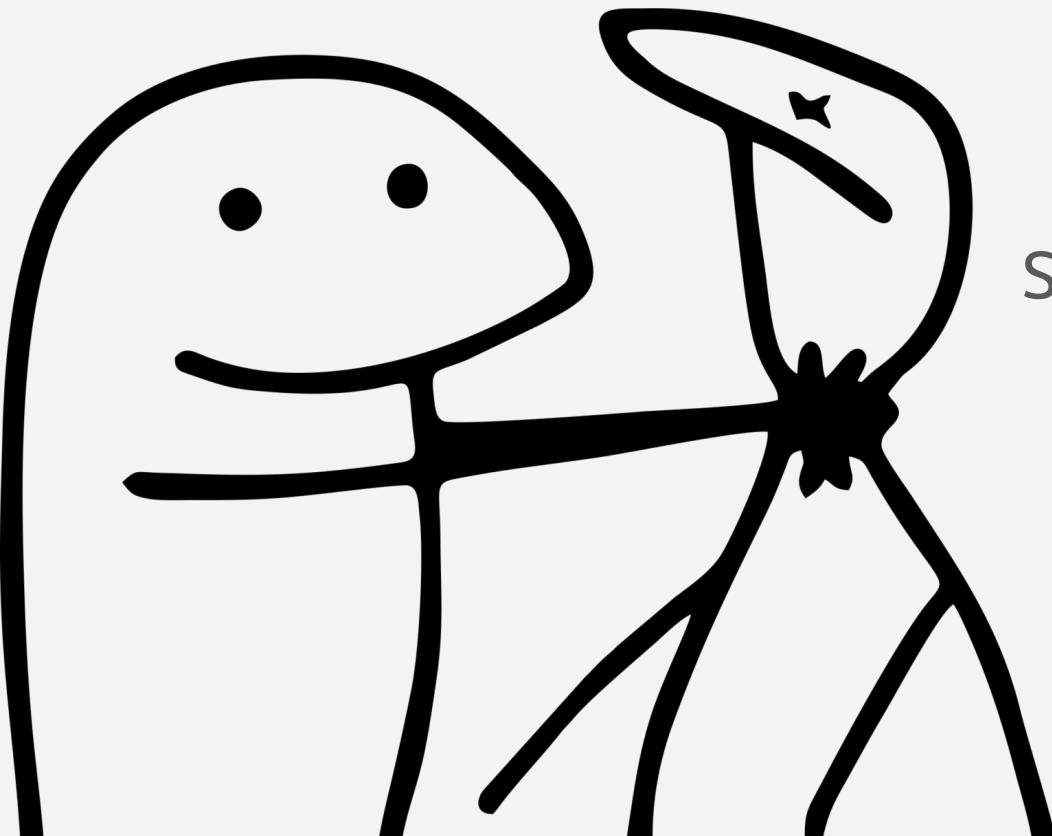
Fluido

Mas o real motivo



Ah, mano. É legal.
**Não tinha ninguém melhor
para palestrar**

Se você achar ruim



Submeta sua palestra para
o C4P da GolangSP e da
GitTogether



SP

Stack trace

```
Starting charge of $49.99 for user user123
panic: runtime error: invalid memory address or nil pointer dereference
[signal SIGSEGV: segmentation violation code=0x2 addr=0x0
pc=0x100782e30]
goroutine 1 [running]:
main.(*PaymentGateway).Charge(0x0, "user123", 49.99)
    payment.go:34 +0x30

main.chargeCard("user123", 49.99)
    payment.go:26 +0xa4

main.processPayment("user123", 49.99)
    payment.go:14 +0x20

main.main()
    payment.go:6 +0x30

exit status 2
```

Stack

```
package main

import "fmt"

func main() {
    userID := "user123"
    amount := 49.99

    processPayment(userID, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
userID)
    // Aciona o gateway de pagamentos
}
```

Stack

```
package main

import "fmt"

func main() {
    userID := "user123"
    amount := 49.99

    processPayment(userID, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
userID)
    // Aciona o gateway de pagamentos
}
```

Cada thread
da sua
aplicação tem
uma

Stack

```
package main

import "fmt"

func main() {
    userID := "user123"
    amount := 49.99

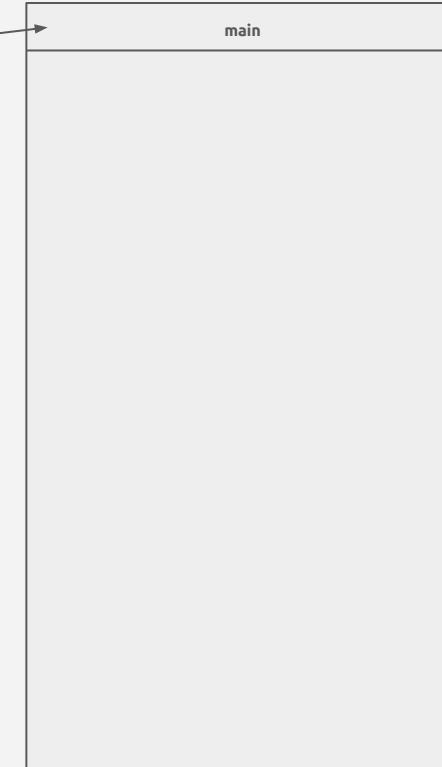
    processPayment(userID, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
userID)
    // Aciona o gateway de pagamentos
}
```

Cria um
stack frame



Stack

```
package main

import "fmt"

func main() {
    userId := "user123"
    amount := 49.99
    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```

Variáveis locais
são
armazenadas
na Stack

main
userId amount

Stack

```
package main

import "fmt"

func main() {
    userId := "user123"
    amount := 49.99

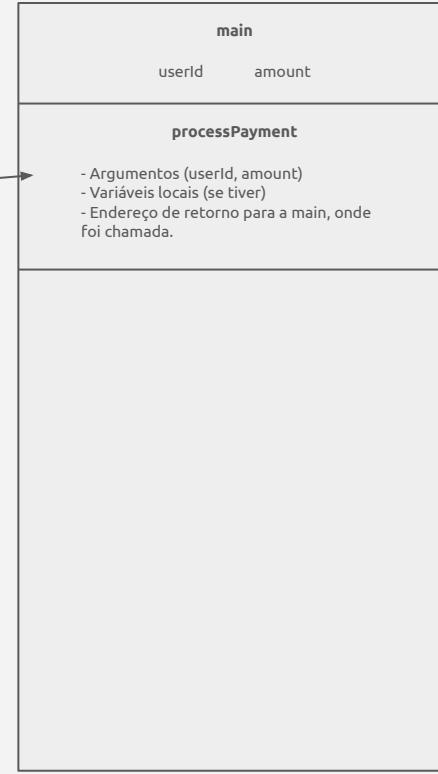
    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

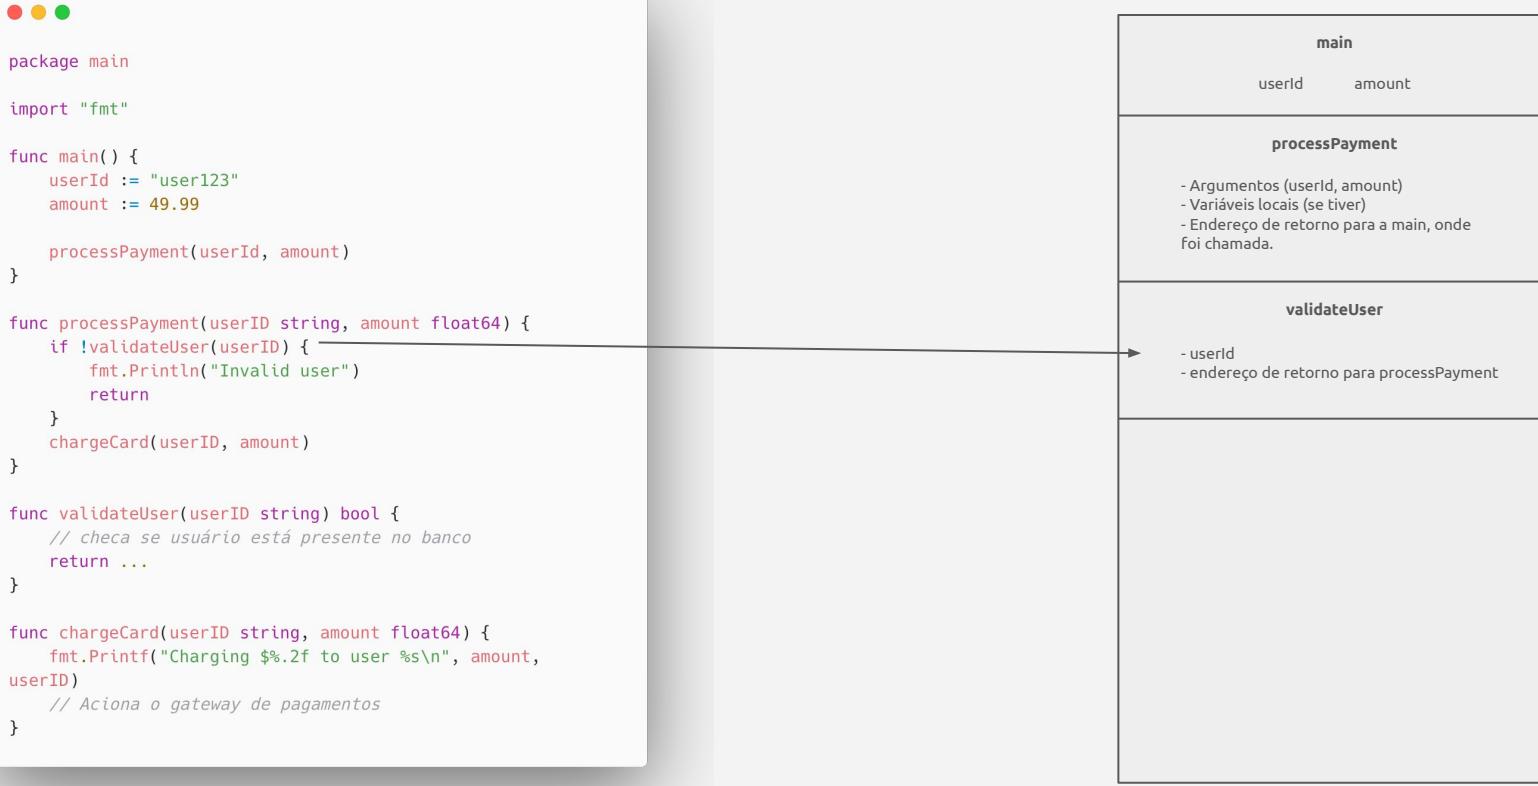
func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```

Cada stack frame mantém o contexto de execução.



Stack



Stack

```
package main

import "fmt"

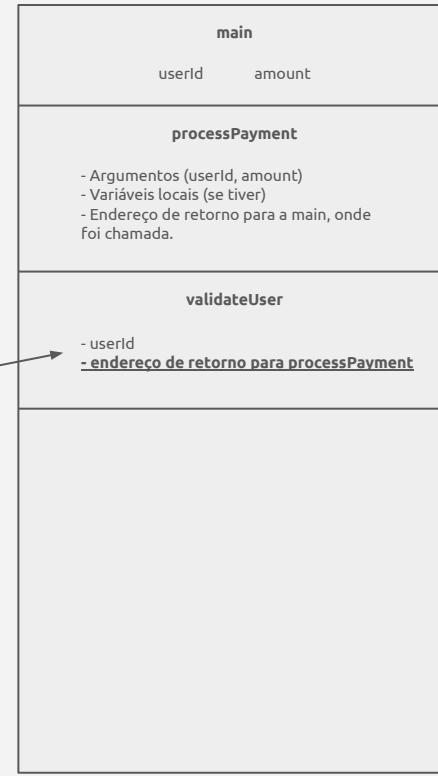
func main() {
    userId := "user123"
    amount := 49.99

    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```



Stack

```
package main

import "fmt"

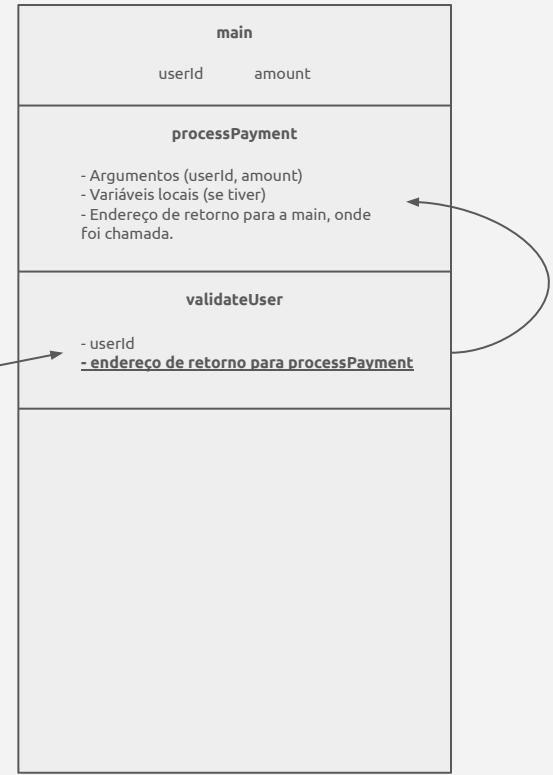
func main() {
    userId := "user123"
    amount := 49.99

    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```



Stack

```
package main

import "fmt"

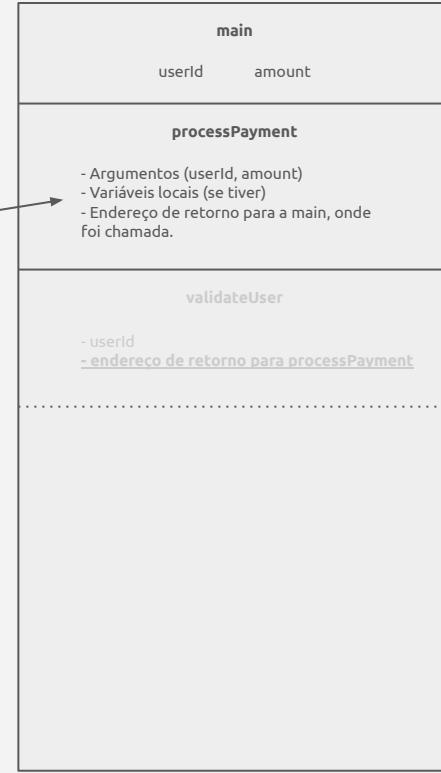
func main() {
    userId := "user123"
    amount := 49.99

    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```



Stack

```
package main

import "fmt"

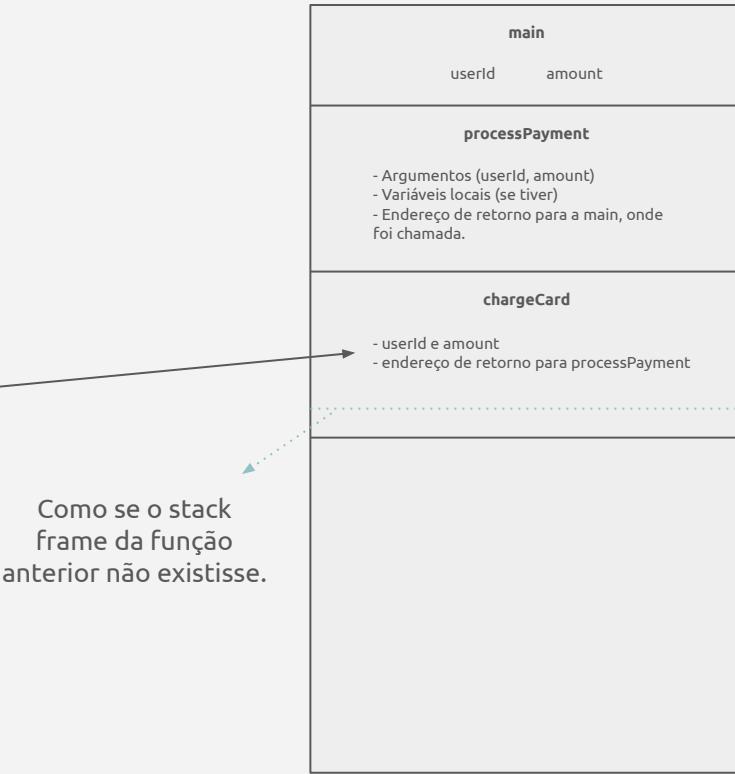
func main() {
    userId := "user123"
    amount := 49.99

    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```



Stack

```
package main

import "fmt"

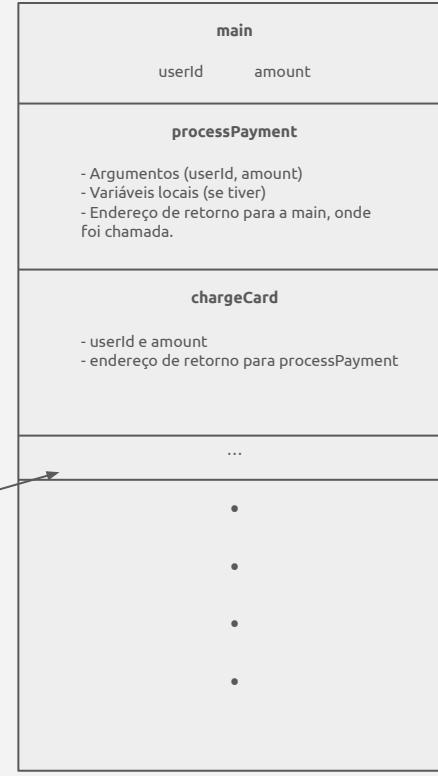
func main() {
    userId := "user123"
    amount := 49.99

    processPayment(userId, amount)
}

func processPayment(userID string, amount float64) {
    if !validateUser(userID) {
        fmt.Println("Invalid user")
        return
    }
    chargeCard(userID, amount)
}

func validateUser(userID string) bool {
    // checa se usuário está presente no banco
    return ...
}

func chargeCard(userID string, amount float64) {
    fmt.Printf("Charging %.2f to user %s\n", amount,
    userID)
    // Aciona o gateway de pagamentos
}
```



Stack

```
Starting charge of $49.99 for user user123
panic: runtime error: invalid memory address or nil pointer dereference
[signal SIGSEGV: segmentation violation code=0x2 addr=0x0
pc=0x100782e30]
goroutine 1 [running]:
main.(*PaymentGateway).Charge(0x0, "user123", 49.99)
    payment.go:34 +0x30

main.chargeCard("user123", 49.99)
    payment.go:26 +0xa4

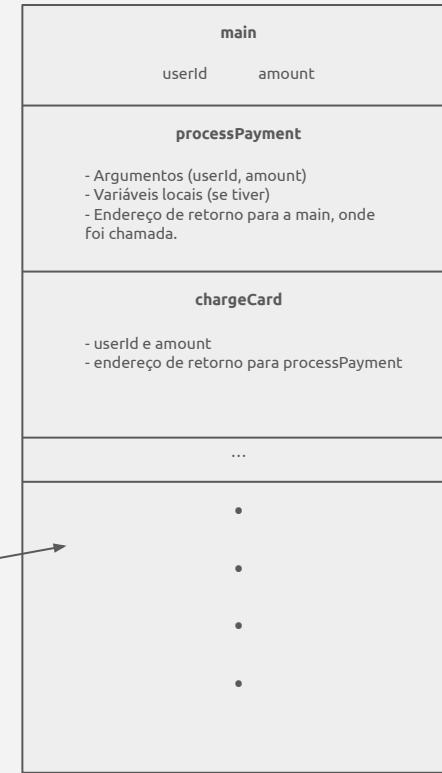
main.processPayment("user123", 49.99)
    payment.go:14 +0x20

main.main()
    payment.go:6 +0x30

exit status 2

userID)
    // Aciona o gateway de pagamentos
}
```

Caminha os endereços de retorno da stack para montar a stack trace.



Quando a stack não é suficiente

```
package main

import (
    "fmt"
)

type Notification struct {
    UserID string
    Channel string
    Message string
}

func main() {
    // ...
    // processPayment
    // ...
    notifications := buildNotifications("user123", 49.99)

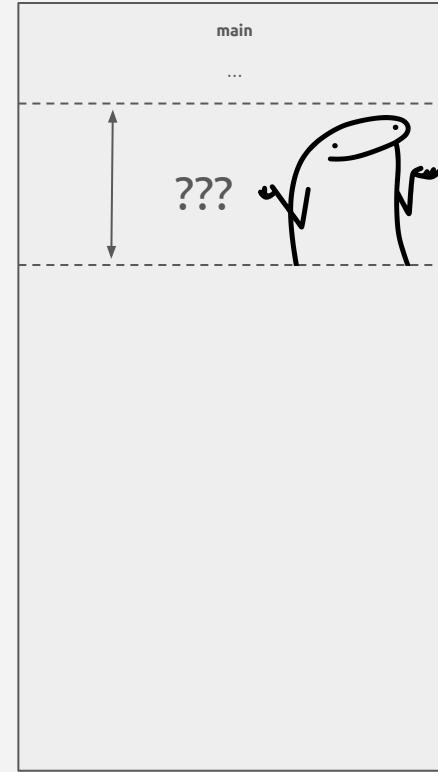
    for _, n := range notifications {
        // Persiste em outbox ou envia notificação
    }
}

func buildNotifications(userID string, amount float64) []Notification {
    notifications := make([]Notification, 0)

    for _, channel = getUserNotificationChannels(userID) {
        notifications = append(notifications, Notification{
            UserID: userID,
            Channel: channel.Type,
            Message: channel.paymentNotification(amount)
        })
    }

    return notifications
}
```

→ Não é possível saber quantas notificações serão geradas, porque depende das configurações do usuário.



Quando a stack não é suficiente

```
package main

import (
    "fmt"
)

type Notification struct {
    UserID string
    Channel string
    Message string
}

func main() {
    // ...
    // processPayment
    // ...
    notifications := buildNotifications("user123", 49.99)

    for _, n := range notifications {
        // Persiste em outbox ou envia notificação
    }
}

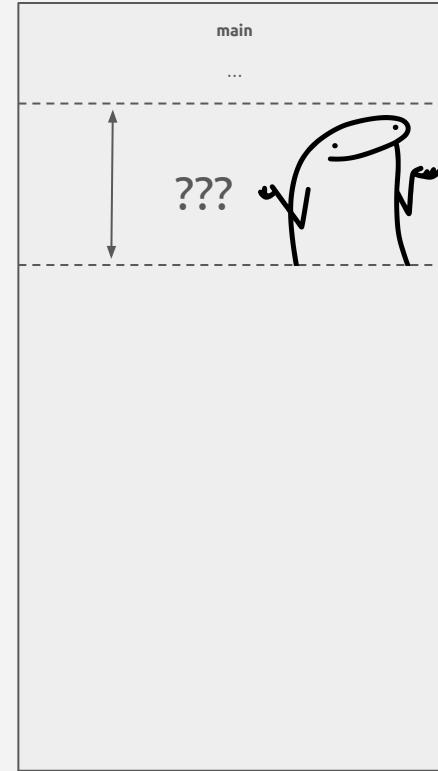
func buildNotifications(userID string, amount float64) []Notification {
    notifications := make([]Notification, 0)

    for _, channel = getUserNotificationChannels(userID) {
        notifications = append(notifications, Notification{
            UserID: userID,
            Channel: channel.Type,
            Message: channel.paymentNotification(amount)
        })
    }

    return notifications
}
```

Além disso, a stack também tem tamanho limitado.

Não é possível saber quantas notificações serão geradas, porque depende das configurações do usuário.



Heap

```
package main

import (
    "fmt"
)

type Notification struct {
    UserID string
    Channel string
    Message string
}

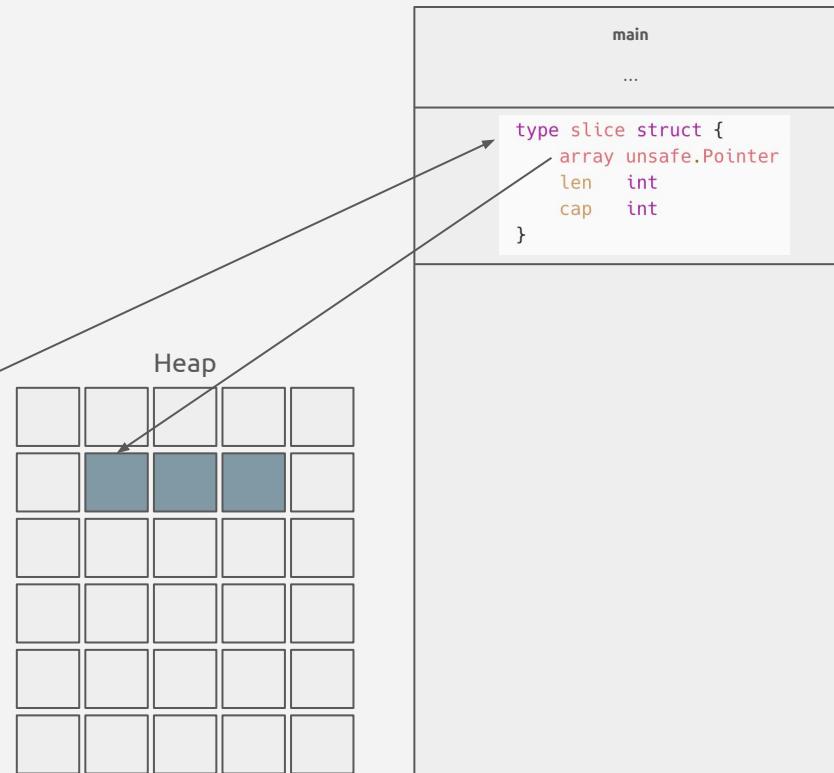
func main() {
    // ...
    // processPayment
    // ...
    notifications := buildNotifications("user123", 49.99)

    for _, n := range notifications {
        // Persiste em outbox ou envia notificação
    }
}

func buildNotifications(userID string, amount float64) []Notification {
    notifications := make([]Notification, 0)

    for _, channel = getUserNotificationChannels(userID) {
        notifications = append(notifications, Notification{
            UserID: userID,
            Channel: channel.Type,
            Message: channel.paymentNotification(amount)
        })
    }

    return notifications
}
```



Heap

```
package main

import (
    "fmt"
)

type Notification struct {
    UserID string
    Channel string
    Message string
}

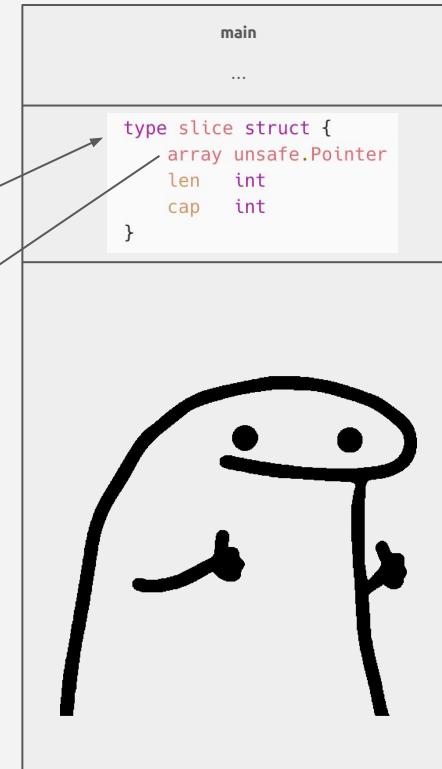
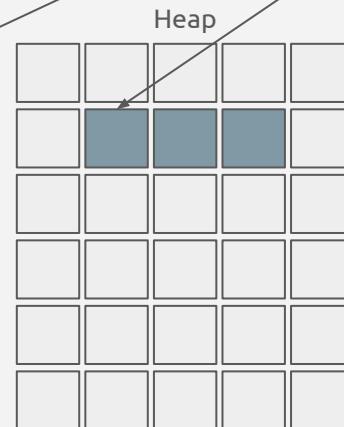
func main() {
    // ...
    // processPayment
    // ...
    notifications := buildNotifications("user123", 49.99)

    for _, n := range notifications {
        // Persiste em outbox ou envia notificação
    }
}

func buildNotifications(userID string, amount float64) []Notification {
    notifications := make([]Notification, 0)

    for _, channel = getUserNotificationChannels(userID) {
        notifications = append(notifications, Notification{
            UserID: userID,
            Channel: channel.Type,
            Message: channel.paymentNotification(amount)
        })
    }

    return notifications
}
```



Gerenciamento Manual

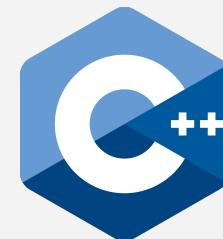
Alocar e desalocar memória são
comandos explícitos.

Gerenciamento Manual



malloc/free

Alocar e desalocar memória são
comandos explícitos.



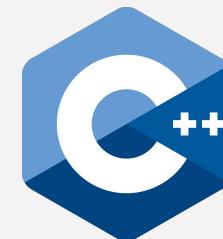
new/delete

Gerenciamento Manual



malloc/free

Alocar e desalocar memória são comandos explícitos.



new/delete

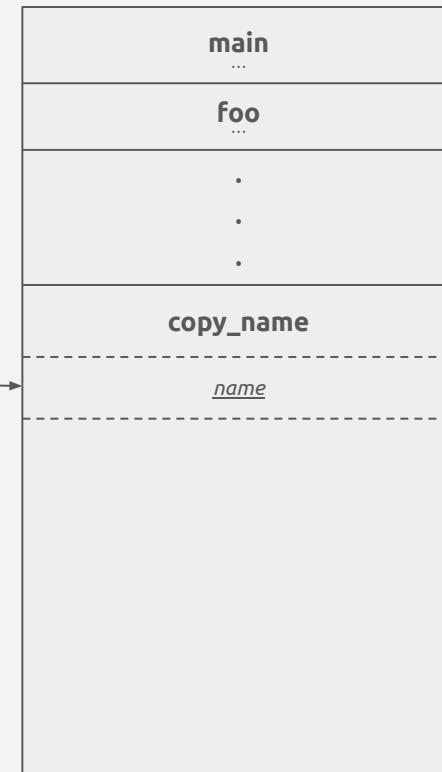


Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char name[100]; ——————
    strcpy(name, input);
    return name;
}
```

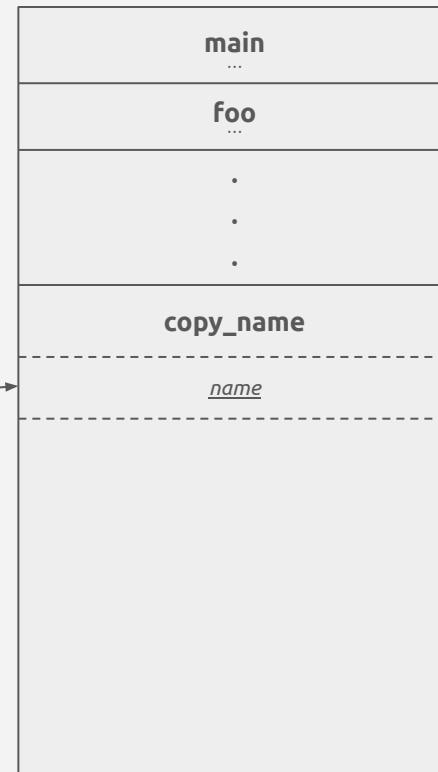


Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char name[100];
    strcpy(name, input);
    return name;
}
```

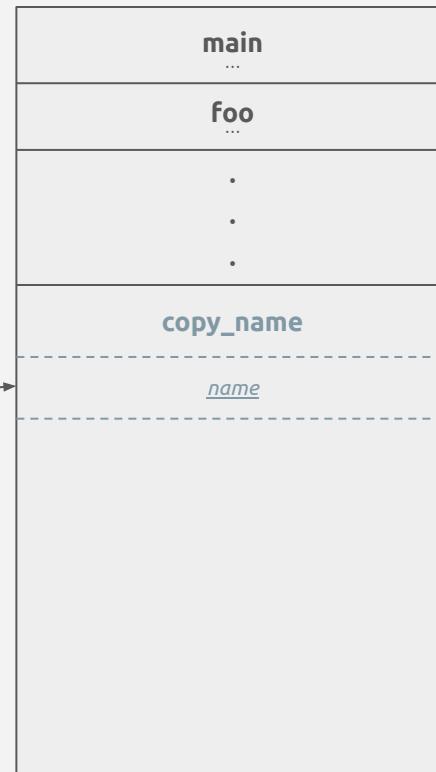


Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char name[100];
    strcpy(name, input);
    return name;
}
```



Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char name[100];
    strcpy(name, input);
    return name;
}
```

main

...

foo

.

.

.

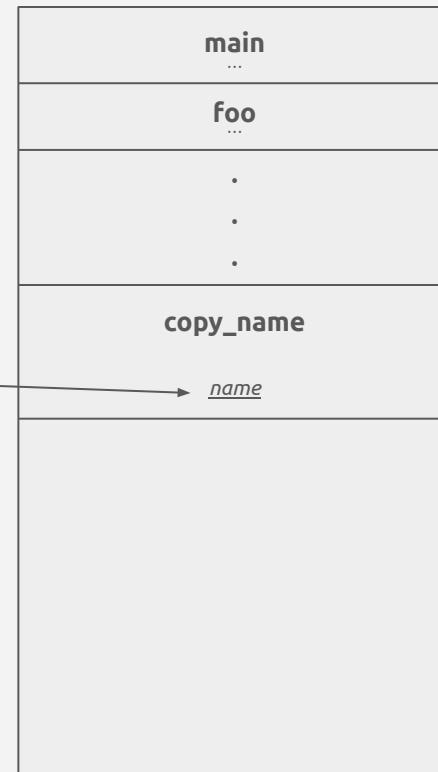
outra_função

Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```

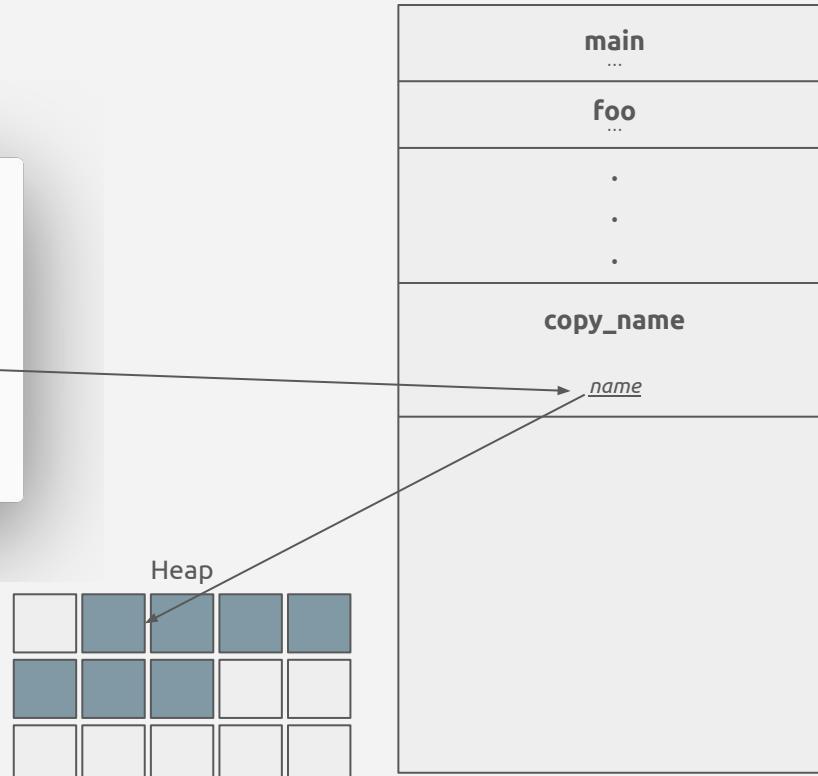


Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



Gerenciamento Manual



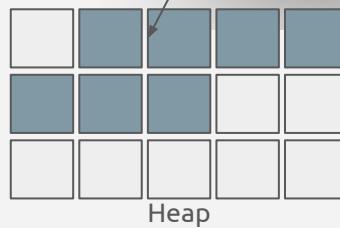
```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname);
    ...
}
```



Gerenciamento Manual



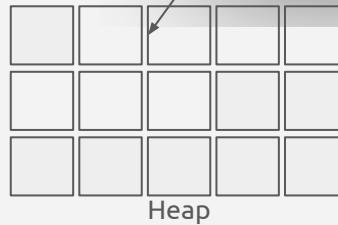
```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname);
    ...
}
```



Gerenciamento Manual



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname);
    ...
}
```

Caso isso não seja feito corretamente...



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname);
    ...
}
```

Memory leak



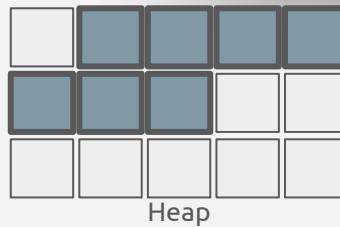
```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname);
    ...
}
```



Double free



```
#include <stdlib.h>
#include <string.h>

char* copy_name(const char* input) {
    char* name = malloc(strlen(input) + 1);
    if (!name) return NULL;
    strcpy(name, input);
    return name;
}
```



```
#include <stdio.h>

int função_chama_copia() {
    ...
    char* myname = copy_name("Renê");
    printf("My name is %s\n", myname);
    free(myname); ←————
    ...
    ...
    free(myname); ←————
}
```

Porque alguém faria isso?



```
#include <stdio.h>

int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    free(myname); ← 1
    ...
}
```

Thread 1



```
#include <stdio.h>

int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    free(myname); ← 2
    ...
}
```

Thread 2

Use-after-free



```
#include <stdio.h>

int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    free(myname); ← 1
    ...
}
```

1



```
#include <stdio.h>

int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    ...
}
```

2

Thread 1

Thread 2

Use-after-free



```
#include <stdio.h>

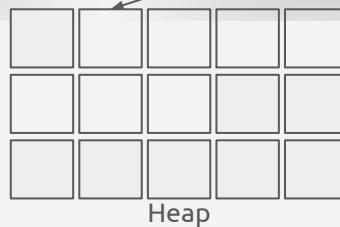
int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    free(myname); ← 1
    ...
}
```



```
#include <stdio.h>

int função_usa_nome(char* name) {
    printf("My name is %s\n", myname);
    ...
}
```

2



Thread 1

Thread 2

**O gerenciamento de memória vira parte
da API**

Gerenciamento Manual



70% dos bugs são relacionados
a gerenciamento de memória.

Qual a solução?



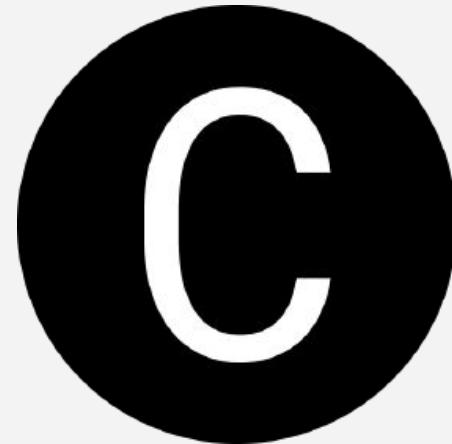
Qual a solução?



ZIG



Rust



Garbage Collection



Gerenciamento de memória automático

Limpar a memória não é o problema

Gerenciamento de memória automático

Limpar a memória não é o problema

O problema é achar aquilo que eu preciso remover

Contagem de referências



```
endereço = Endereço("Rua A", "Bairro B", ...)  
pessoa = Pessoa("Alice", endereço)  
return pessoa
```

Contagem de referências



```
endereço = Endereço("Rua A", "Bairro B", ...) ←
pessoa = Pessoa("Alice", endereço)
return pessoa
```

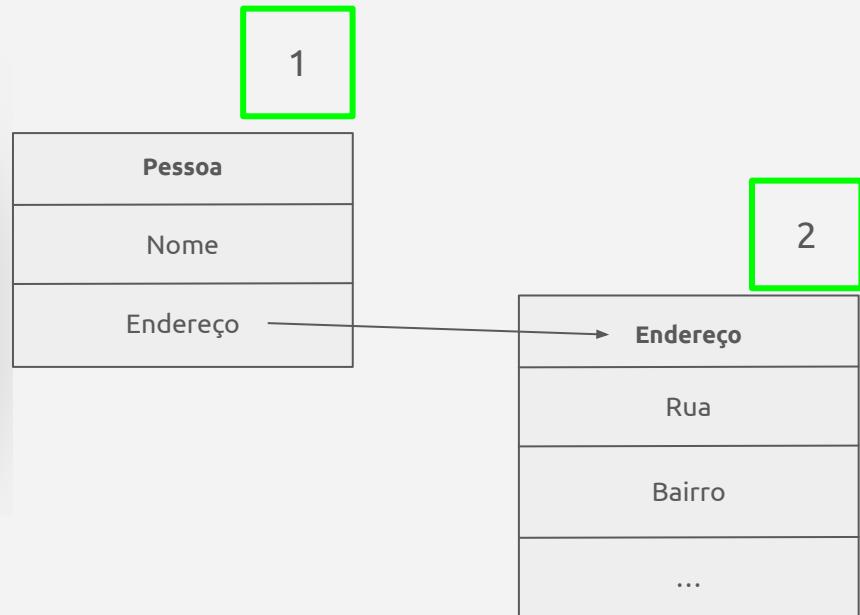
1

Endereço
Rua
Bairro
...

Contagem de referências



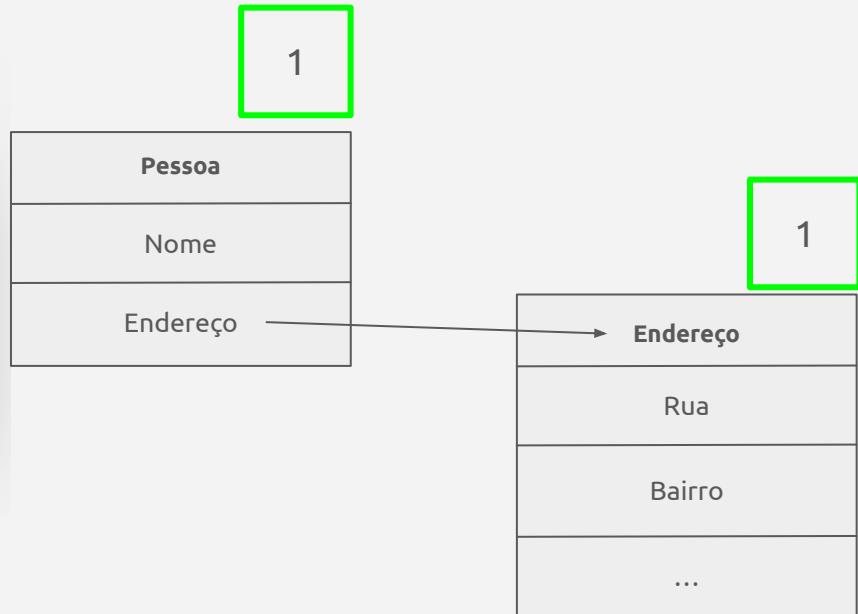
```
endereço = Endereço("Rua A", "Bairro B", ...)  
pessoa = Pessoa("Alice", endereço) ←  
return pessoa
```



Contagem de referências



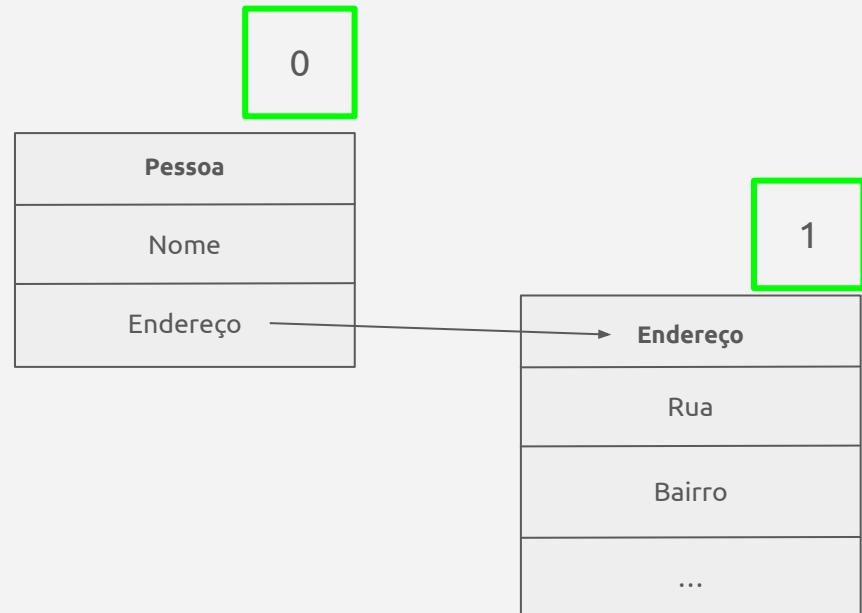
```
endereço = Endereço("Rua A", "Bairro B", ...)  
pessoa = Pessoa("Alice", endereço)  
return pessoa
```



Contagem de referências



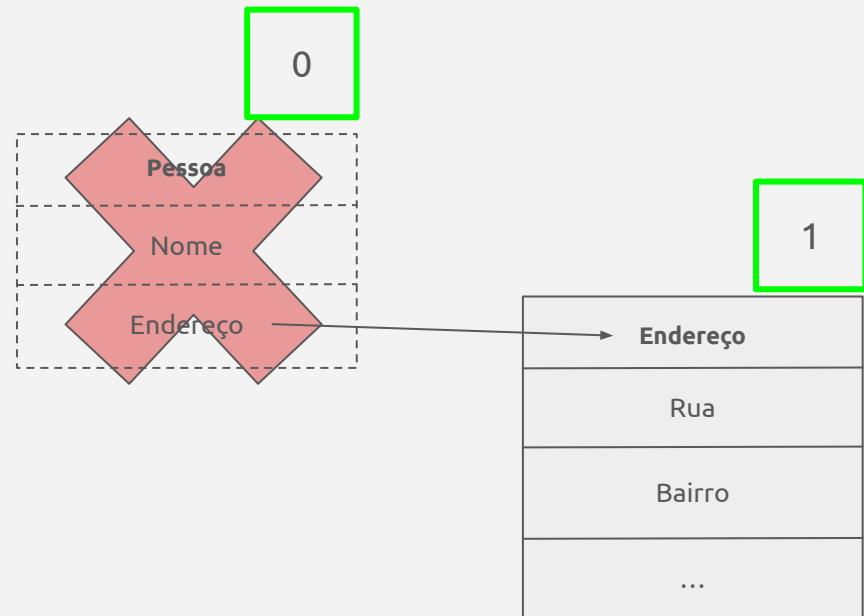
pessoa = None



Contagem de referências



pessoa = None



Contagem de referências



pessoa = None

0

Endereço

Rua

Bairro

...

Contagem de referências



pessoa = None

Contagem de referências



```
def virar_melhores_amigos(p1: Pessoa, p2: Pessoa) -> None:
    p1.melhor_amigo = p2
    p2.melhor_amigo = p1

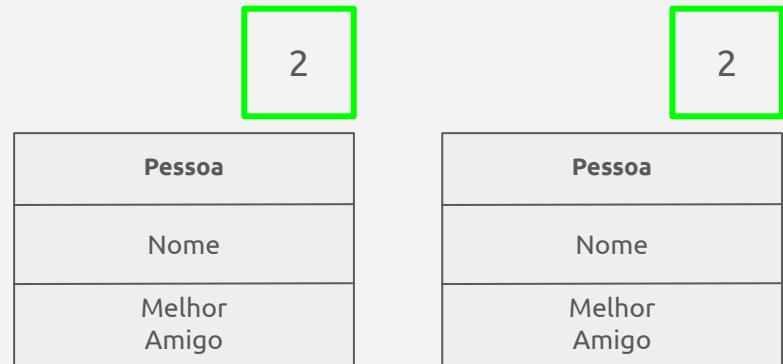
alice = Pessoa("Alice")
bob   = Pessoa("Bob")

virar_melhores_amigos(alice, bob)
```

Contagem de referências



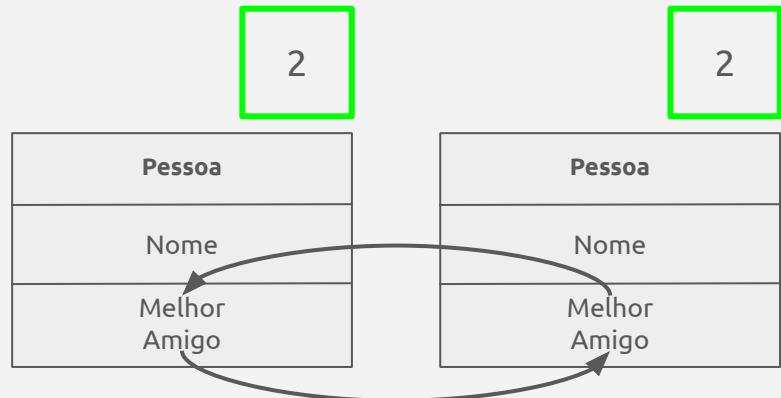
```
def virar_melhores_amigos(p1: Pessoa, p2: Pessoa) -> None:  
    p1.melhor_amigo = p2  
    p2.melhor_amigo = p1  
  
alice = Pessoa("Alice")  
bob   = Pessoa("Bob")  
  
virar_melhores_amigos(alice, bob) ←
```



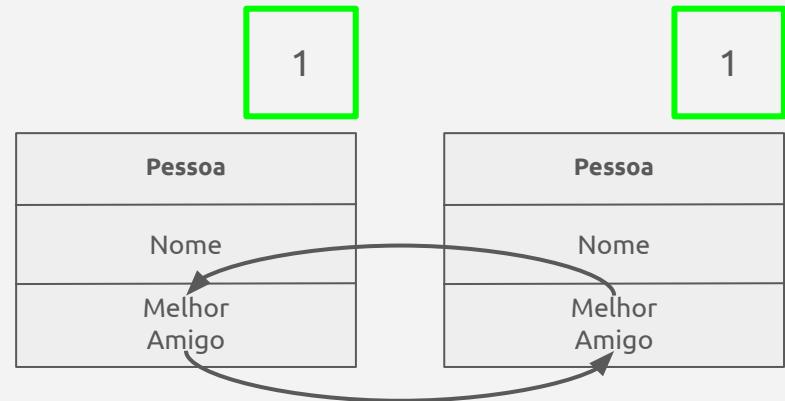
Contagem de referências



```
def virar_melhores_amigos(p1: Pessoa, p2: Pessoa) -> None:  
    p1.melhor_amigo = p2  
    p2.melhor_amigo = p1  
  
alice = Pessoa("Alice")  
bob   = Pessoa("Bob")  
  
virar_melhores_amigos(alice, bob)
```



Contagem de referências falha em ciclos

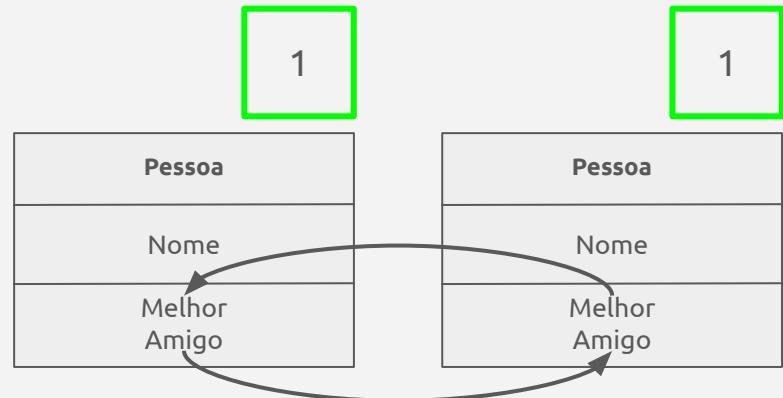


Desvantagens

Necessita de contador para cada objeto na memória

Cada atribuição tem custo

Ciclos necessitam de tratamento especial



Apesar das desvantagens



Algumas libs
e frameworks

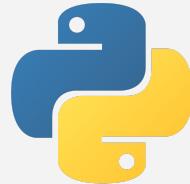
Apesar das desvantagens



Algumas libs
e frameworks



Swift



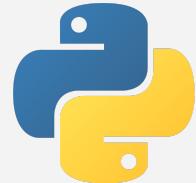
Apesar das desvantagens



Algumas libs
e frameworks



Swift



A resolução de ciclos é
responsabilidade do
programador

Apesar das desvantagens

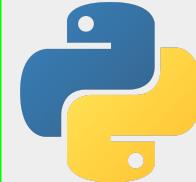


Algumas libs
e frameworks



Swift

A resolução de ciclos é
responsabilidade do
programador



Ainda precisam de
um GC para ciclos

De onde vem o Garbage Collector



Al e Garbage Collector

?

Al e Garbage Collector



John McCarthy

Al e Garbage Collector



John McCarthy



O IBM704 já tinha Fortran

```
C  PROGRAM FOR FINDING THE LARGEST VALUE IN A SET
C  (ARRAY A CONTAINS N NUMBERS; RESULT IS LEFT IN BIGA)

      BIGA = A(1)

      DO 20 I = 2, N          C  LOOP OVER THE REST OF THE ARRAY

10      IF (BIGA - A(I)) 15, 20, 20  C  ARITHMETIC IF:
           C    <0  → 15
           C    =0  → 20
           C    >0  → 20

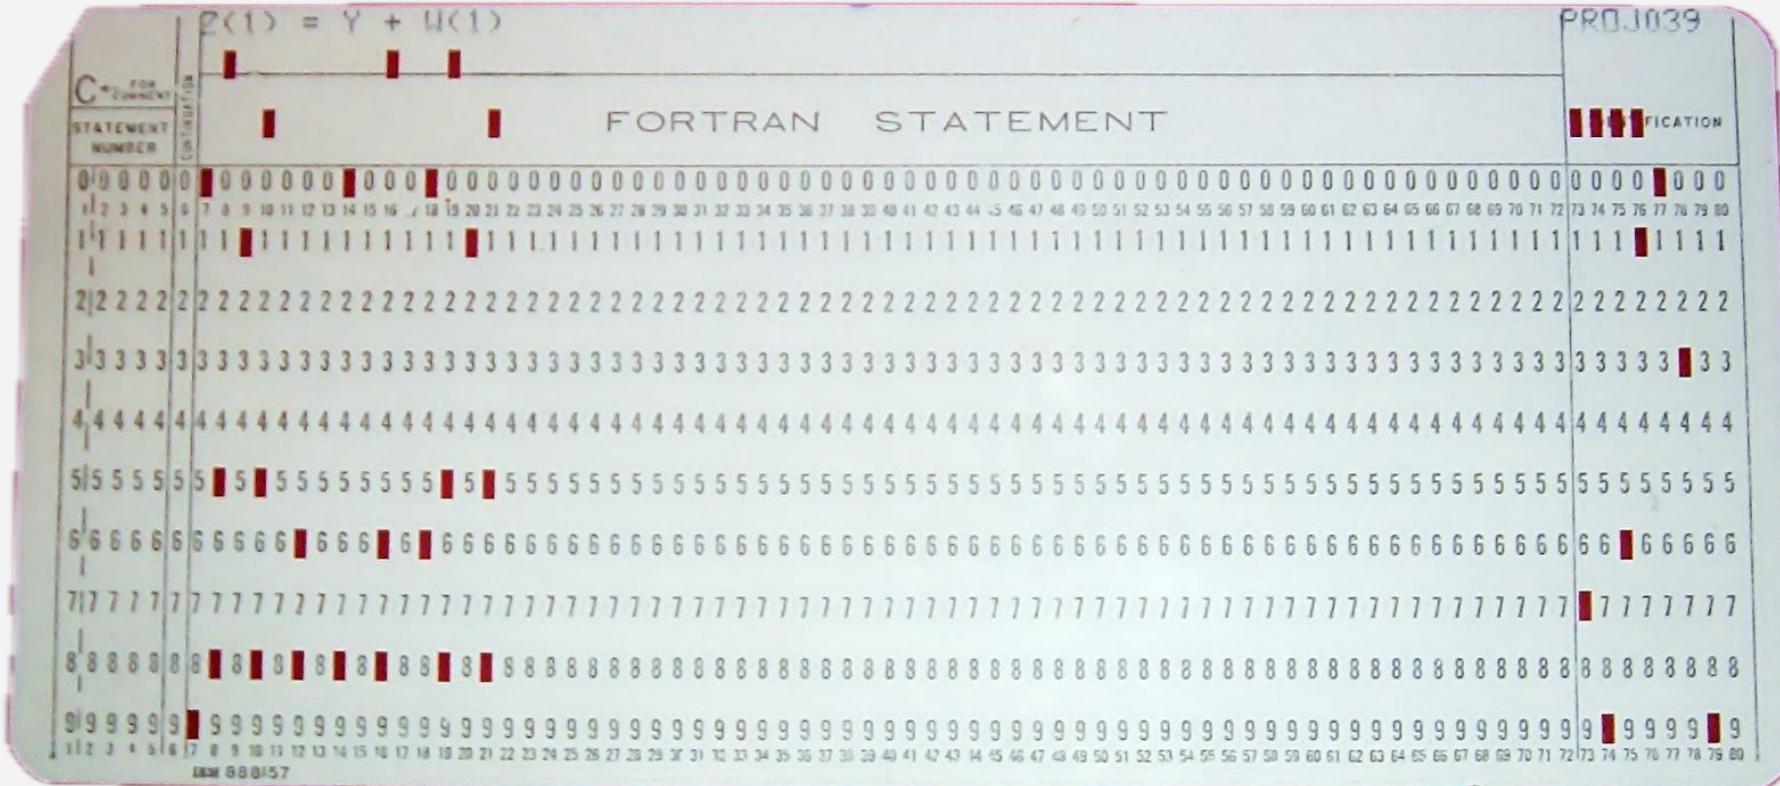
15      BIGA = A(I)          C  NEW MAXIMUM FOUND
      GOTO 20

20      CONTINUE          C  END OF LOOP

      WRITE (6,30) BIGA      C  PRINT THE ANSWER
30      FORMAT (1H , 'MAX VALUE =', 1PE14.7)

      STOP
      END
```

O IBM704 já tinha Fortran



A large black silhouette of a person's arm and hand reaches across the page from the left side, pointing its index finger towards the right edge of the document. The silhouette is dark and solid, casting a long shadow across the paper.

FORTRAN STATEMENT

FORTRAN STATEMENT

FOR
COMMENT
EMENT
NUMBER

CONTINUATION

FOR
COMMENT
EMENT
NUMBER

CONTINUATION

Fortran utilizava apenas stack

A screenshot of a comment section from a dark-themed platform. At the top left, it says "1,071 Comments". To its right is a "Sort by" button with a dropdown arrow. Below this is a search bar with the placeholder "Add a comment..." and a blue "GO" button with a magnifying glass icon. The main content area shows a single comment by a user with a purple profile picture containing a white letter "F". The comment text is "Fortran 68 years ago First". Below the text are interaction icons: a thumbs-up icon with the number "363", a reply icon, and a "Reply" link. A blue downward arrow icon followed by "7 replies" indicates there are more replies to this comment. On the far right of the comment card is a vertical ellipsis ("...").

John McCarthy decide criar uma linguagem para IA

1,071 Comments Sort by

 Add a comment... 

 **Fortran** 68 years ago 
First
 363  [Reply](#)
 [7 replies](#)

 **LISP** 65 years ago 
Pau no ** do first
 835  [Reply](#)
 [10 replies](#)

 **ALGOL** 65 years ago 
Aff. Mó babaca. Cheguei antes.
 1.7K  [Reply](#)
 [17 replies](#)

LISP



```


$$\begin{aligned}
& \text{DEFINe DIFF} \\
& (\text{LAMBDA } (E X) \\
& (\text{COND } ((\text{ATOM } E) \quad ; \text{ number or variable} \\
& \quad (\text{COND } ((\text{EQ } E X) 1) (\text{T } 0))) \\
& \quad ((\text{EQ } (\text{CAR } E) \text{ 'PLUS}) \quad ; \text{ (PLUS } u v) \\
& \quad (\text{LIST } \text{'PLUS} \\
& \quad \quad (\text{DIFF } (\text{CADR } E) X) \\
& \quad \quad (\text{DIFF } (\text{CADDR } E) X))) \\
& \quad ((\text{EQ } (\text{CAR } E) \text{ 'TIMES}) \quad ; \text{ (TIMES } u v) \\
& \quad (\text{LIST } \text{'PLUS} \\
& \quad \quad (\text{LIST } \text{'TIMES} \\
& \quad \quad \quad (\text{DIFF } (\text{CADR } E) X) \\
& \quad \quad \quad (\text{CADDR } E)) \\
& \quad \quad (\text{LIST } \text{'TIMES} \\
& \quad \quad \quad (\text{CADR } E) \\
& \quad \quad \quad (\text{DIFF } (\text{CADDR } E) X)))) \\
& \quad (\text{T } \text{'UNSUPPORTED})))) \\
\\
& ; \text{ Differentiate } (x+3)*(x+4) \text{ with respect to } x: \\
& (\text{DIFF } '(\text{TIMES } (\text{PLUS } X 3) (\text{PLUS } X 4)) \text{ 'X})
\end{aligned}$$


```

```
(((((define fact))))))  
((((((= n 0))))))  
(((((((((t)))))))  
((((((fact )))))  
((((((fact 1))))  
((((((fact )))))  
(((((( ))))))  
(((((( )))))  
(((( )))
```

Lisp stands for “Lots of Irritating Superfluous Parentheses”

LISP

COBOL

YOU BUILT THE HORSE
IN 1962



IT CAN ONLY BE TAMED BY
THE ORIGINAL CREATOR.

~~FOR ALL OTHER PURPOSES
IT'S A DRAGON.~~



LISP

Clojure é um dialeto de Lisp



Clojure é um dialeto de Lisp



nu

Clojure é um dialeto de Lisp

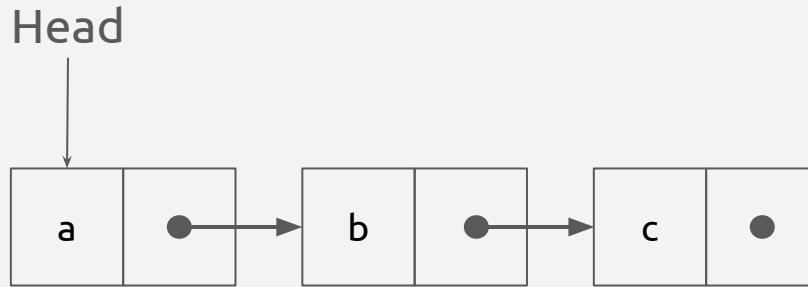


nu

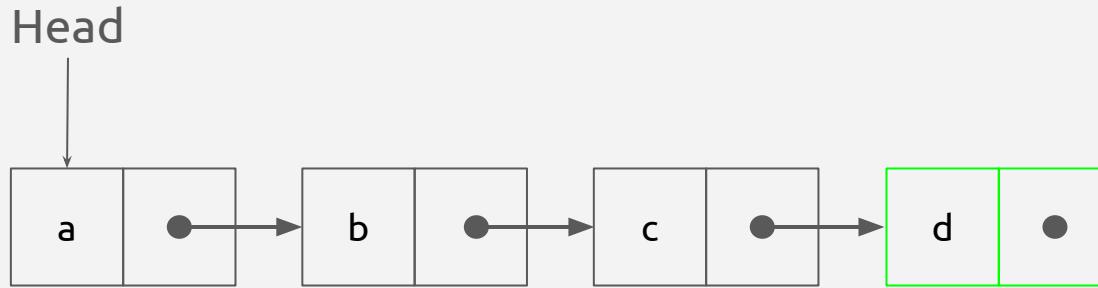
Walmart

NETFLIX

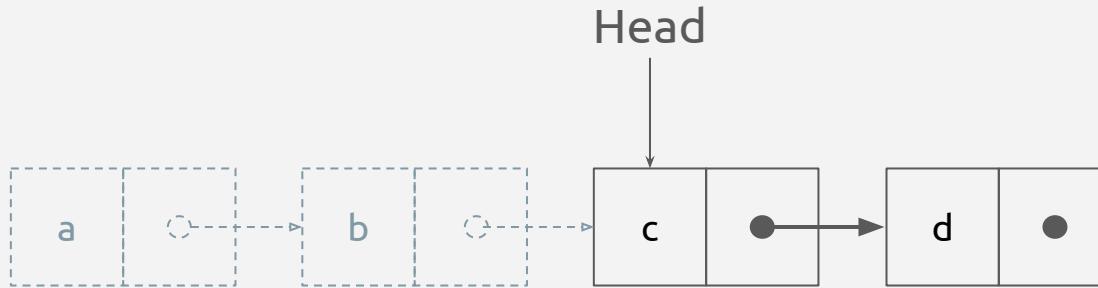
LISt Processor



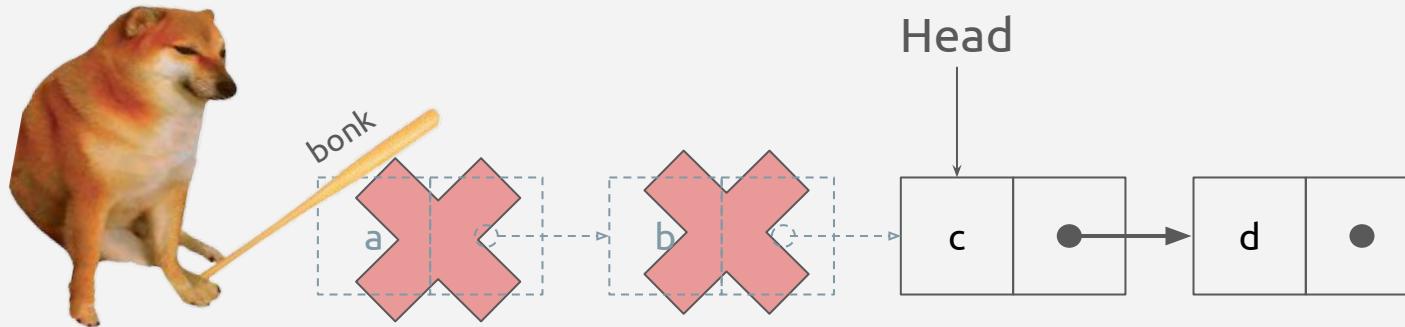
LISt Processor



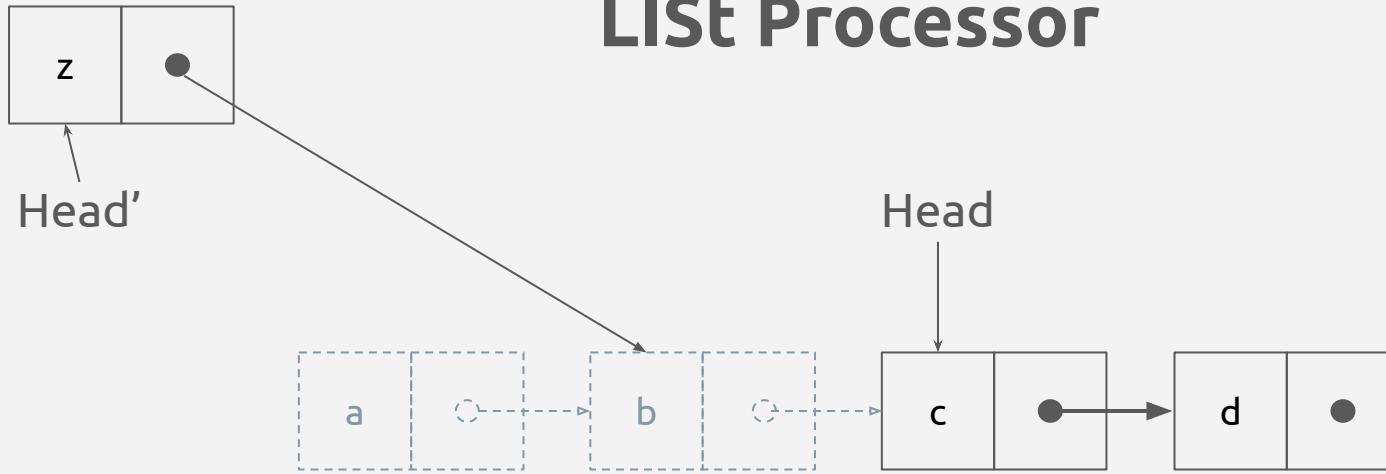
LISt Processor



LIS Processor



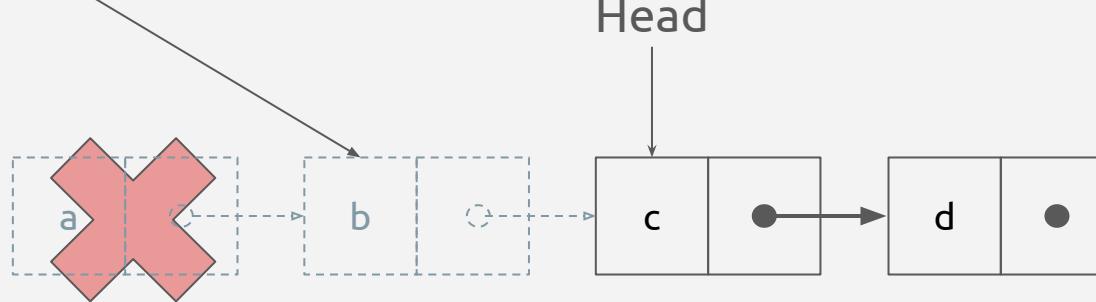
LISt Processor



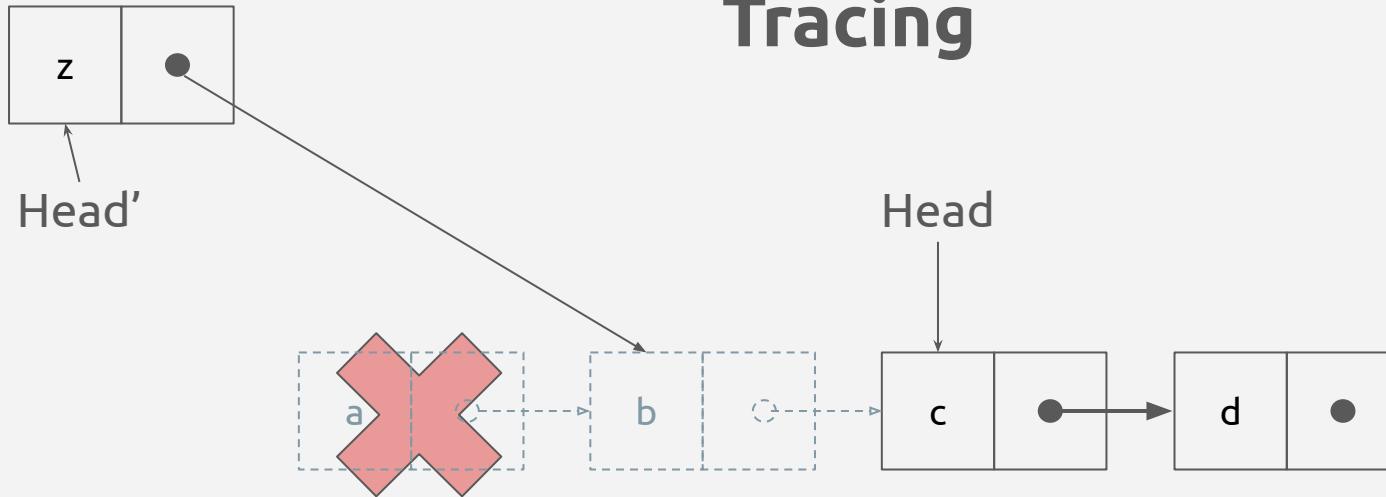
Como saber o que não é utilizado?



Head'



Tracing

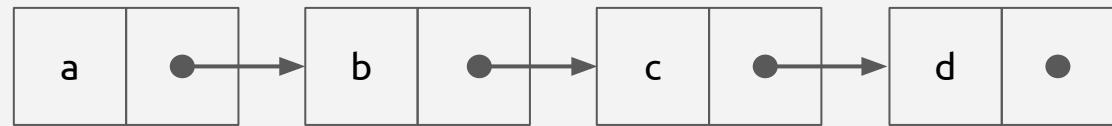


Roots

Head

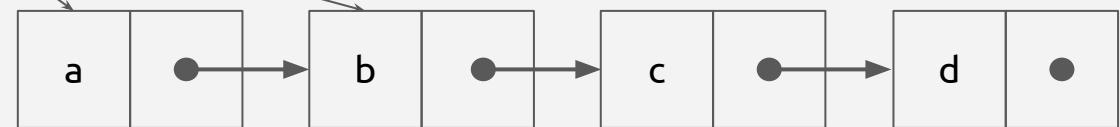
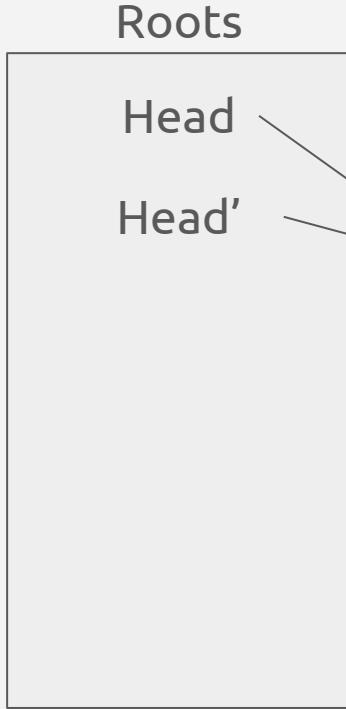
Head'

Tracing



Variáveis
declaradas
(Stack)

Tracing

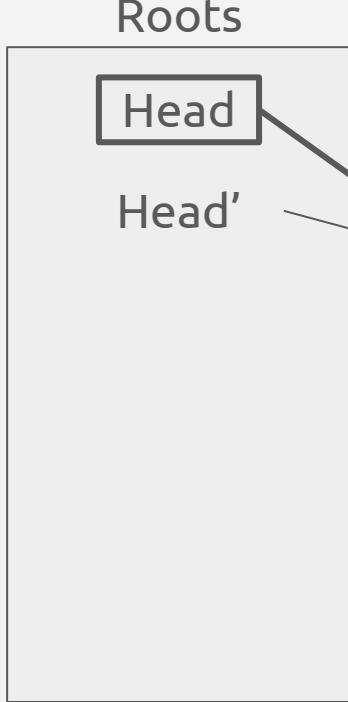


Em uso



Não Utilizada

Tracing

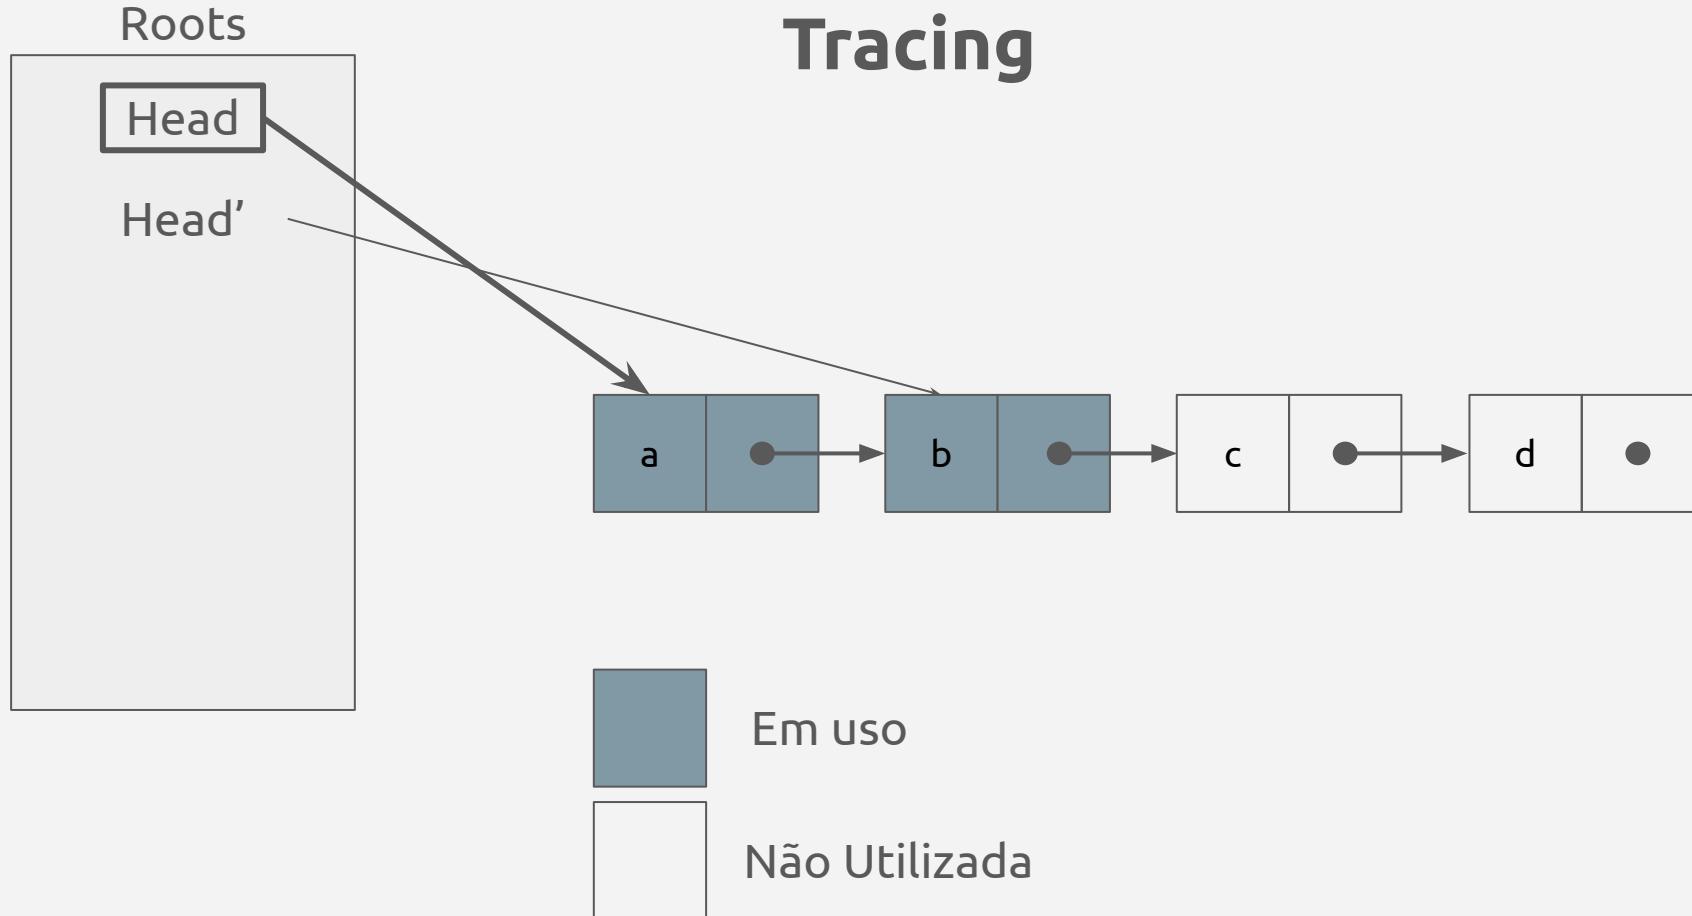


Em uso

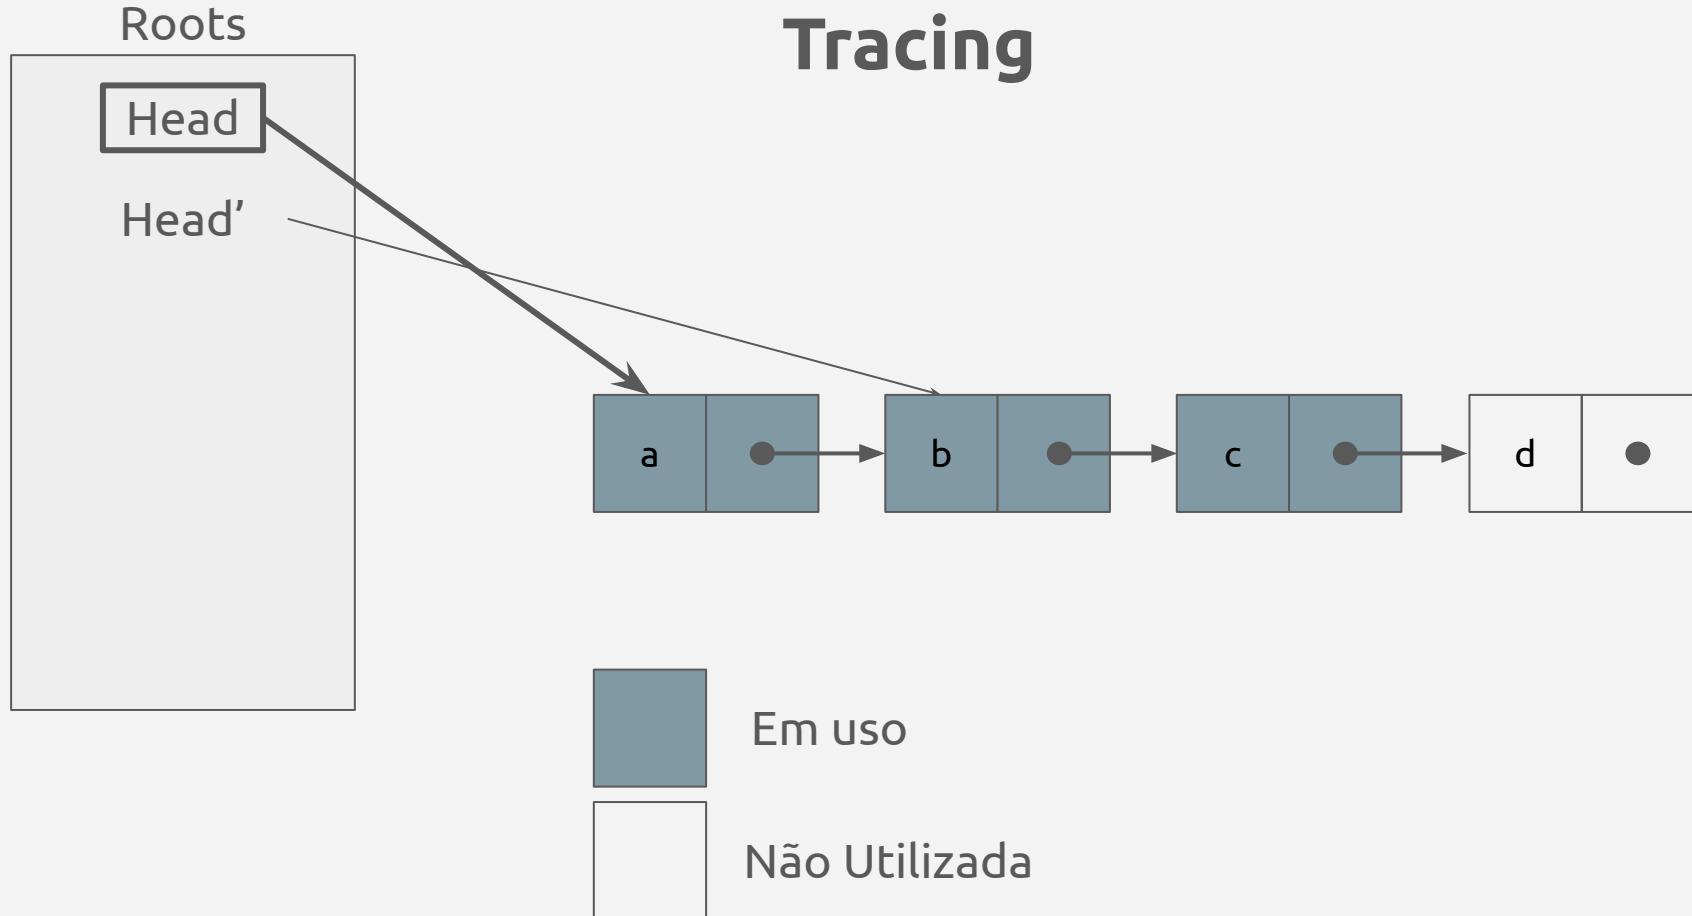


Não Utilizada

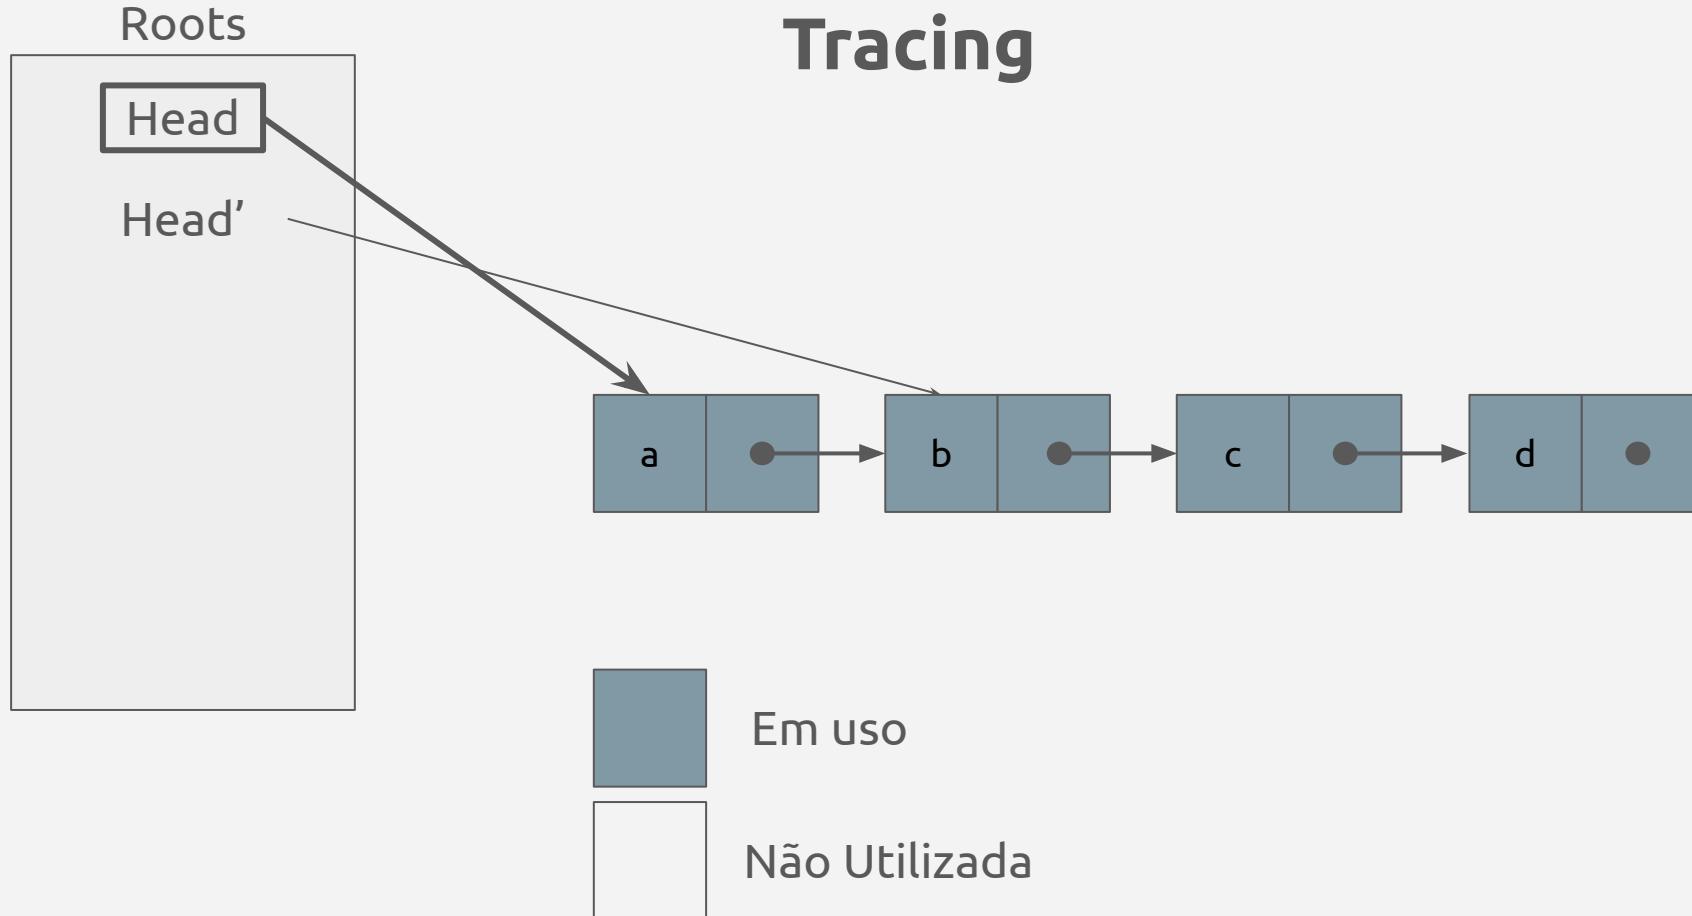
Tracing



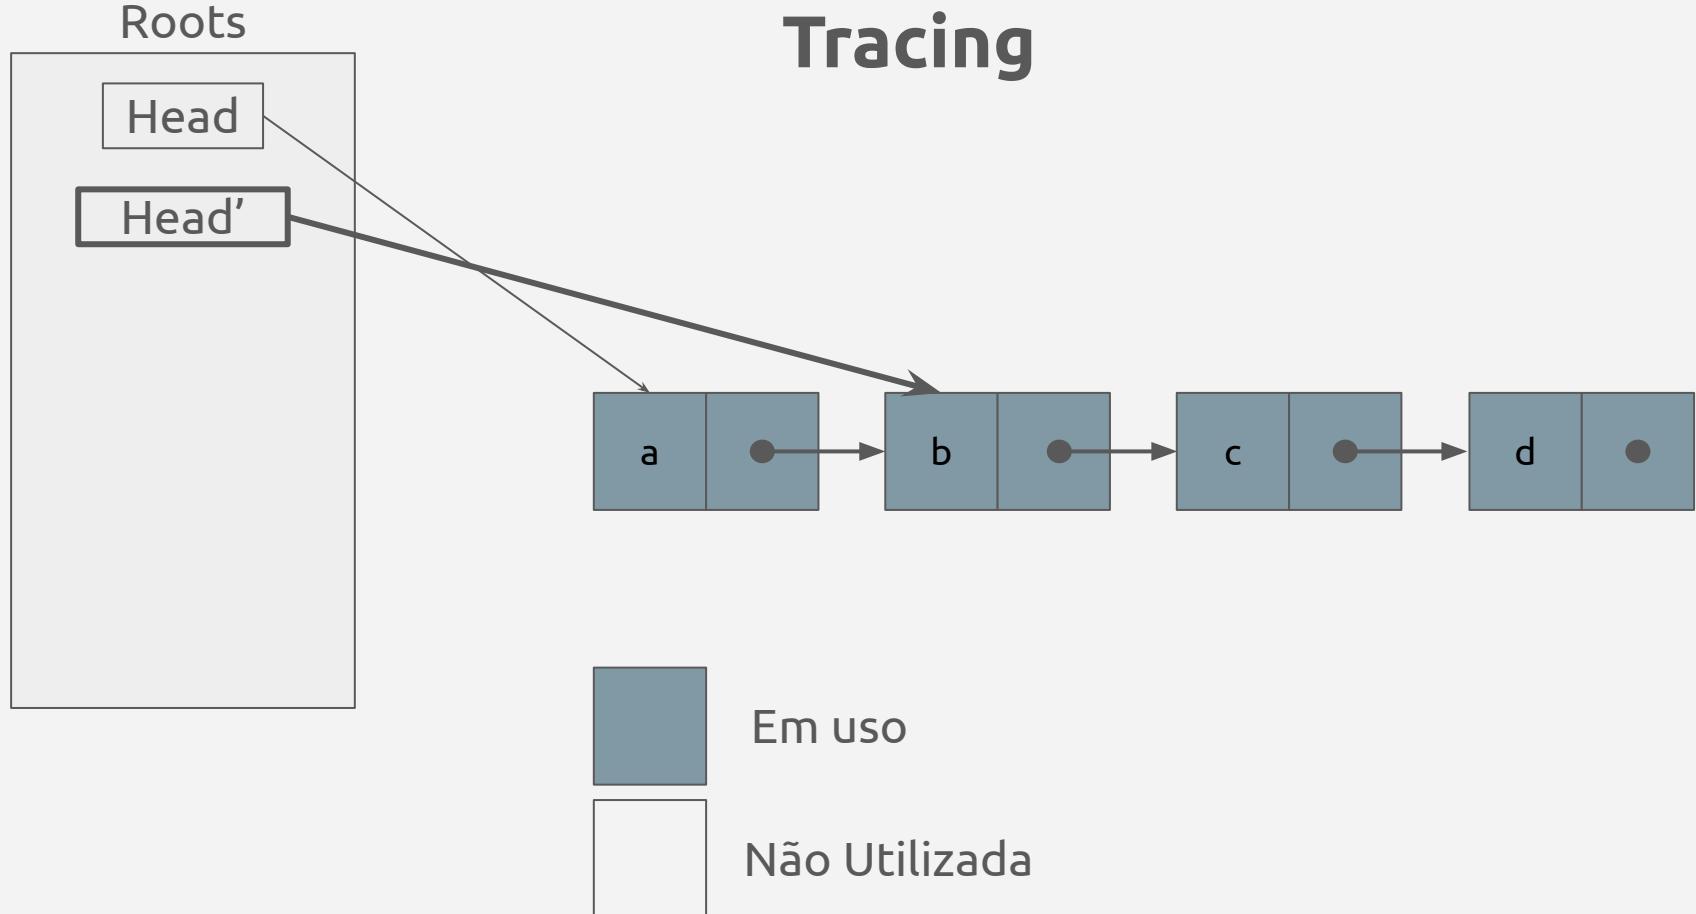
Tracing



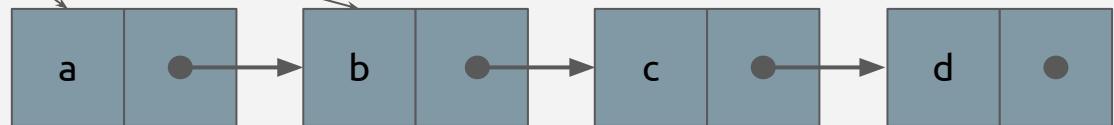
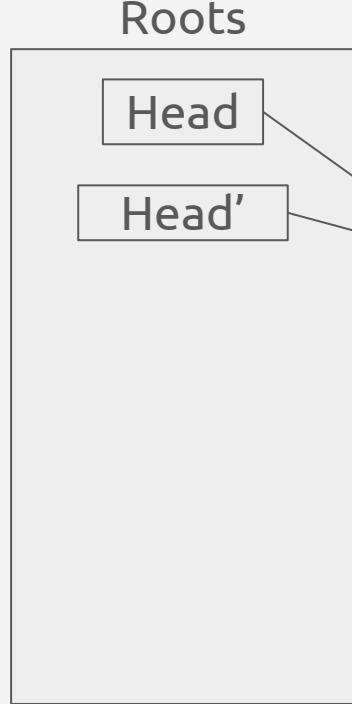
Tracing



Tracing



Tracing

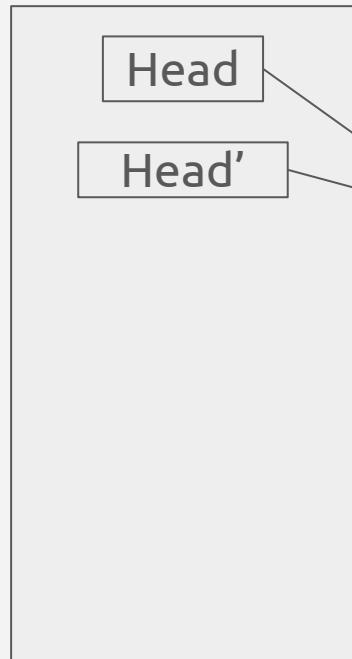


Em uso



Não Utilizada

Roots



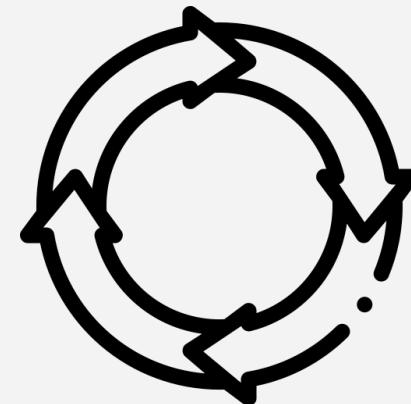
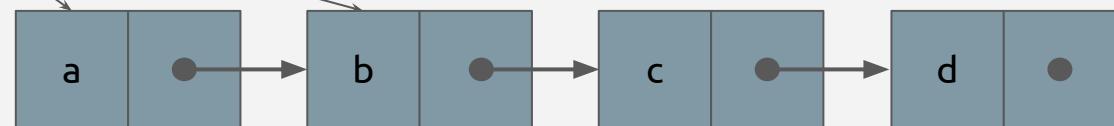
Tracing



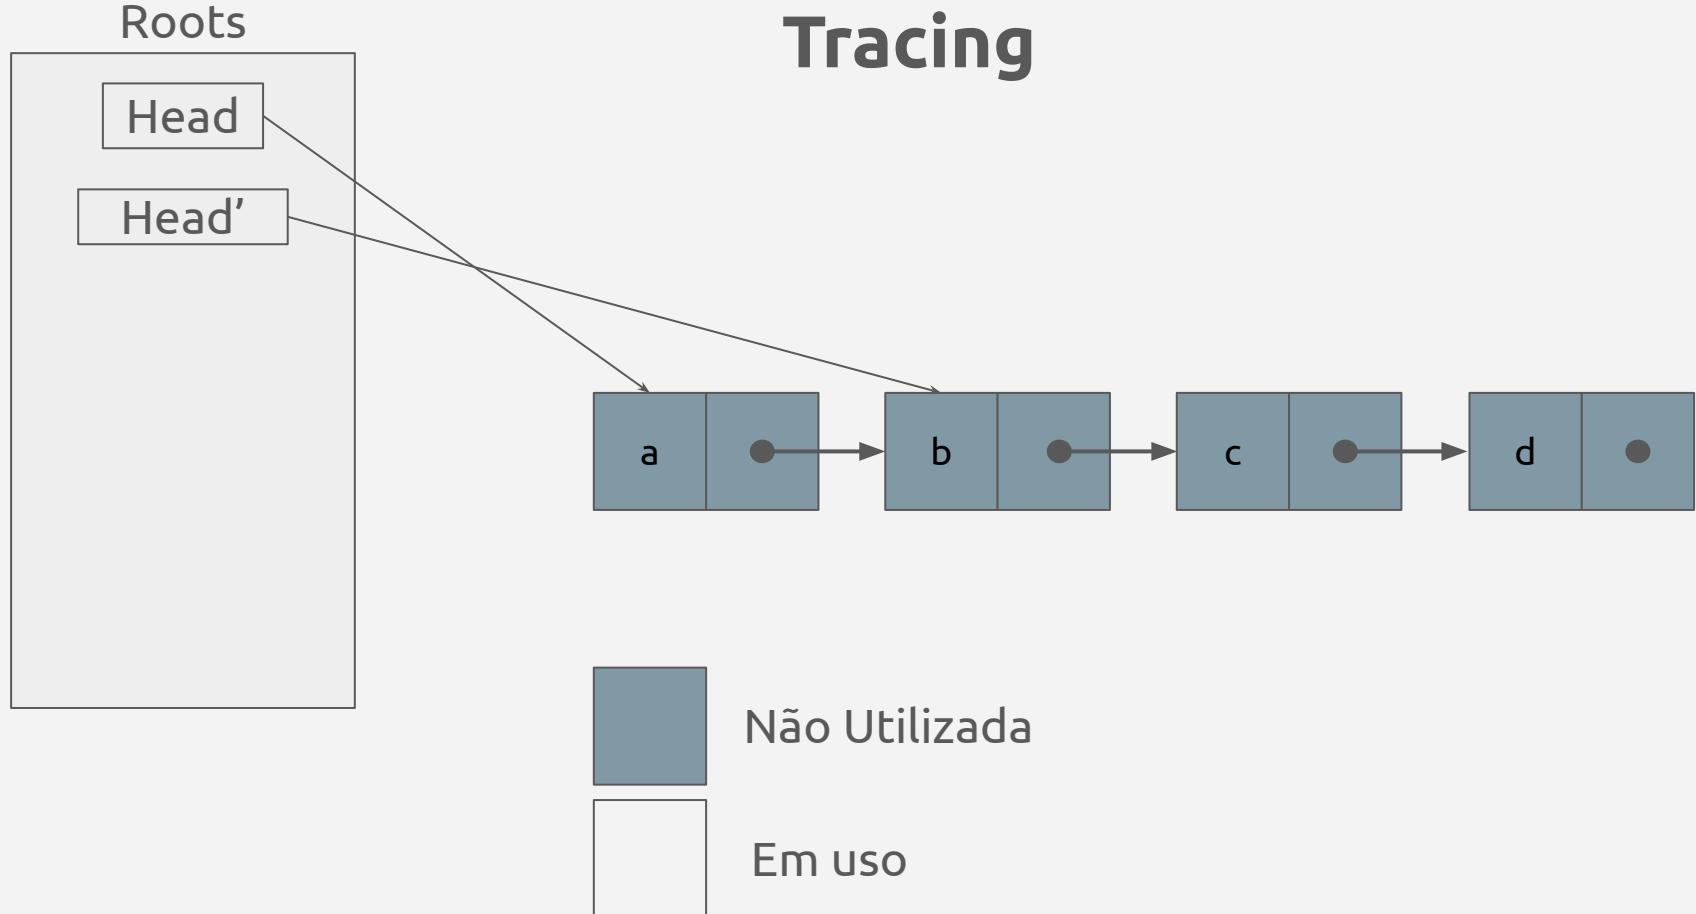
Em uso



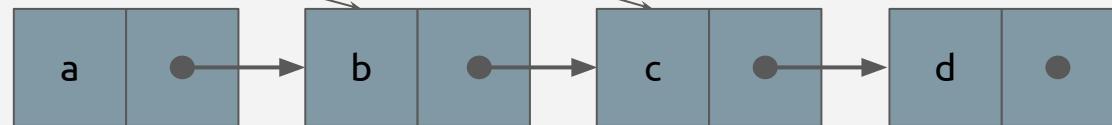
Não Utilizada



Tracing



Tracing

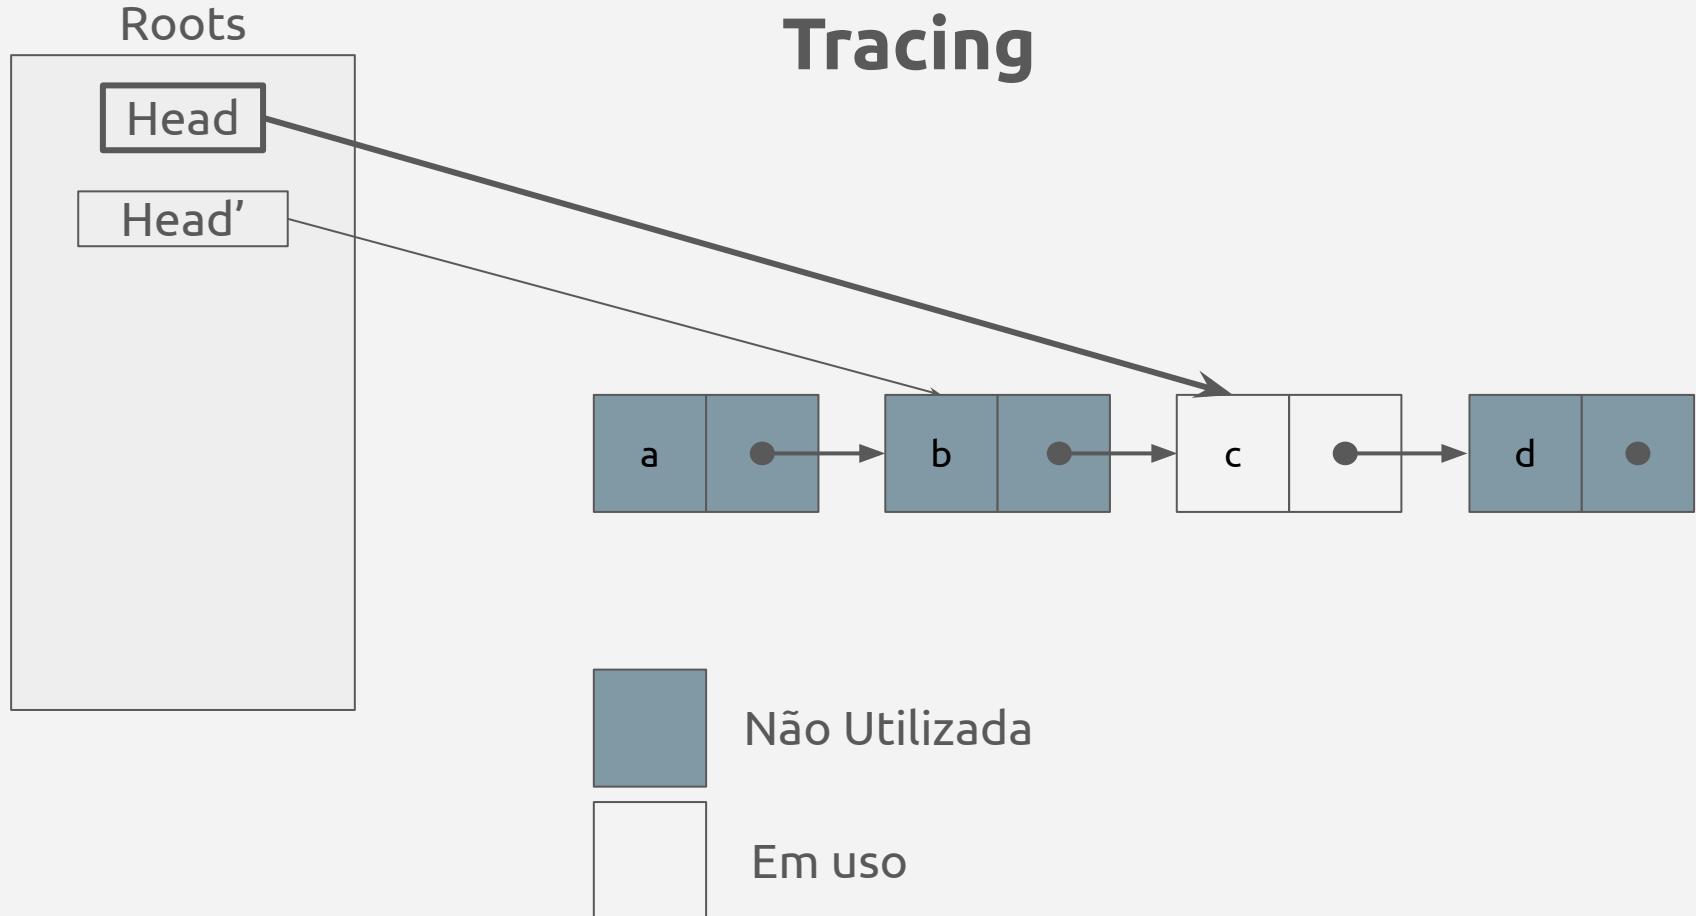


Não Utilizada

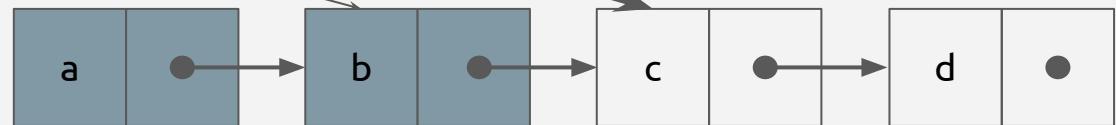
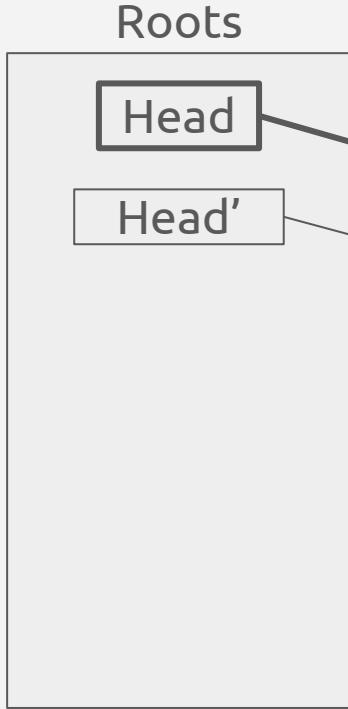


Em uso

Tracing



Tracing

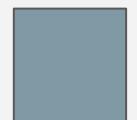
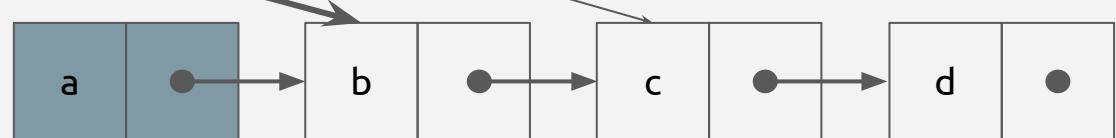
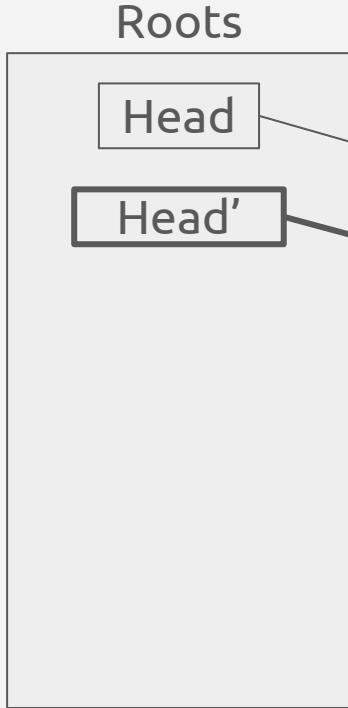


Não Utilizada



Em uso

Tracing

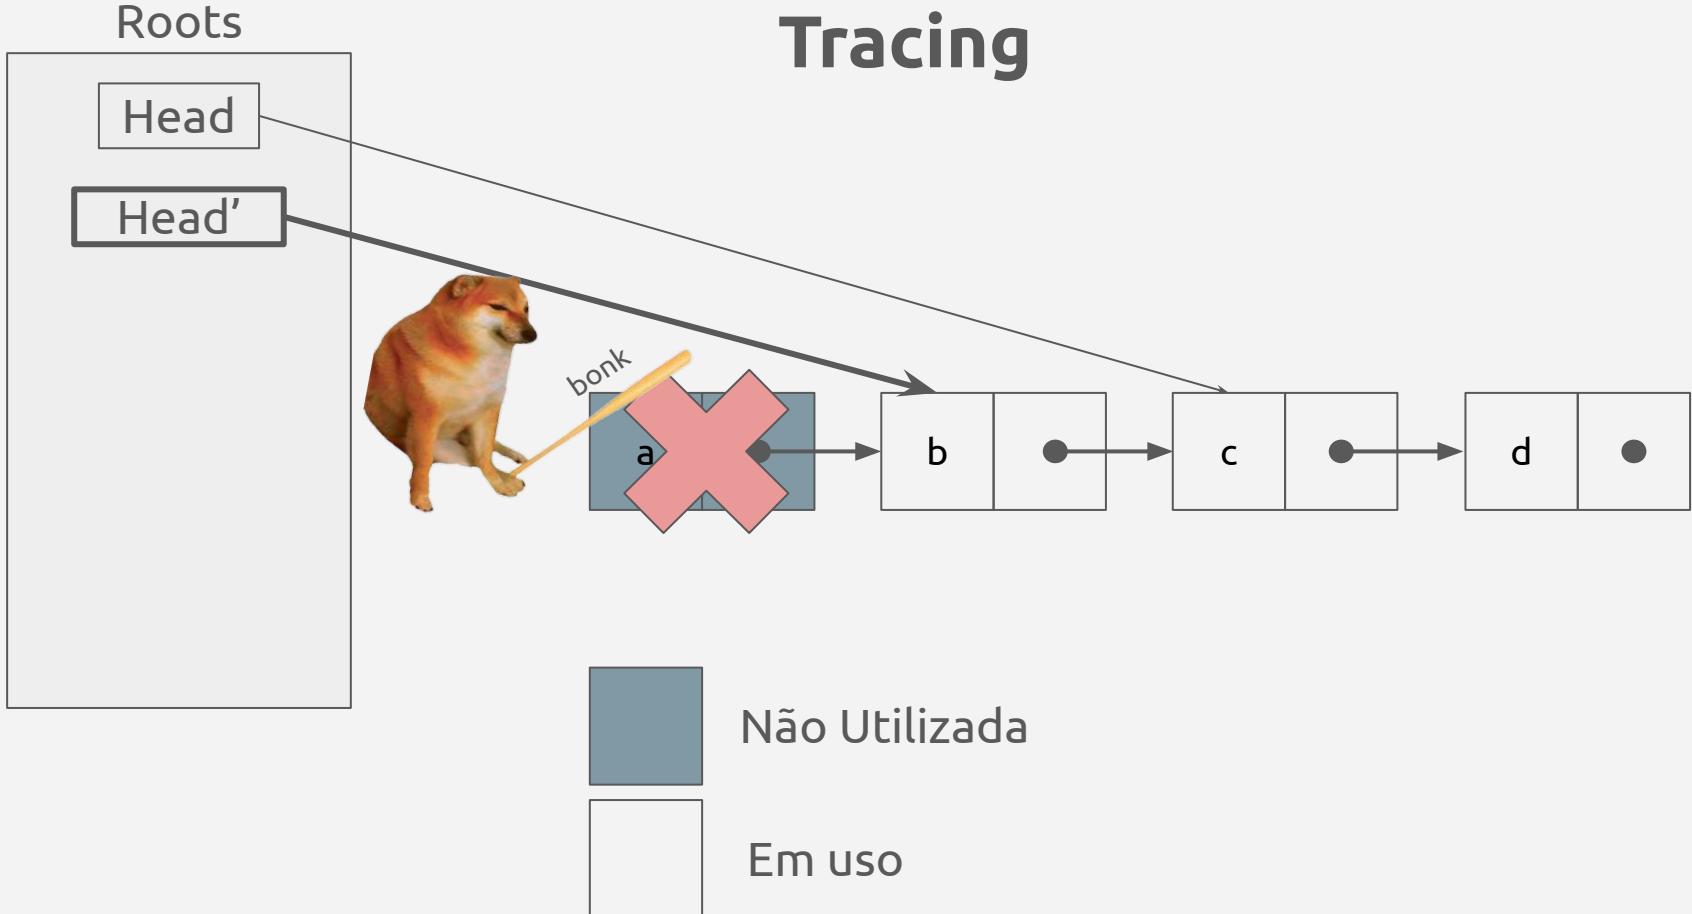


Não Utilizada

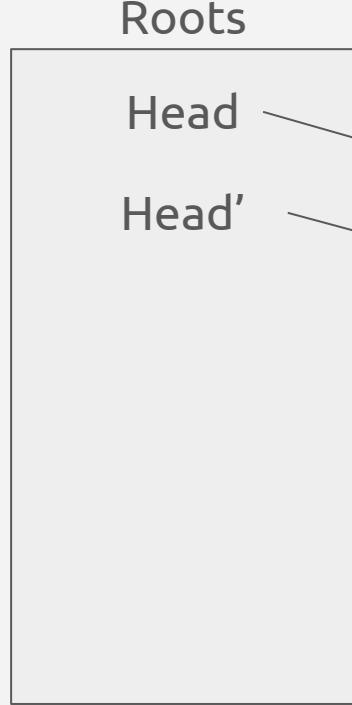


Em uso

Tracing



Tracing

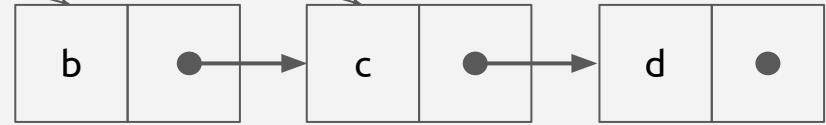


Não Utilizada



Em uso

Tracing

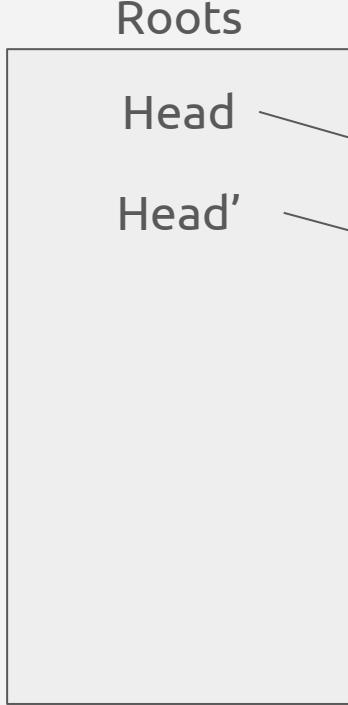


Em uso



Não Utilizada

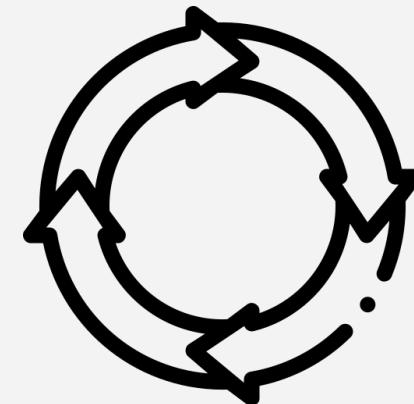
Tracing



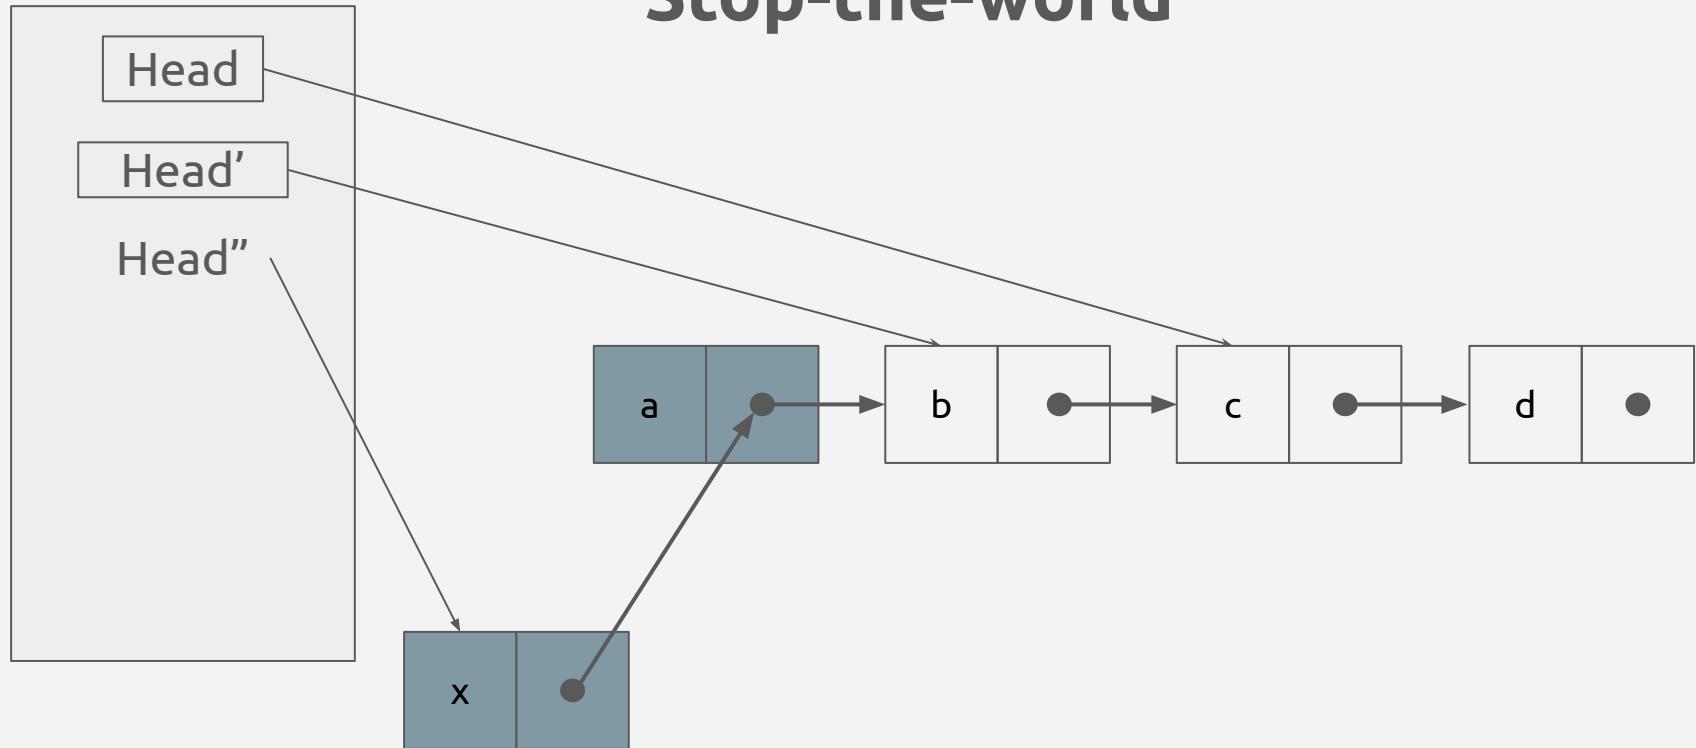
Em uso



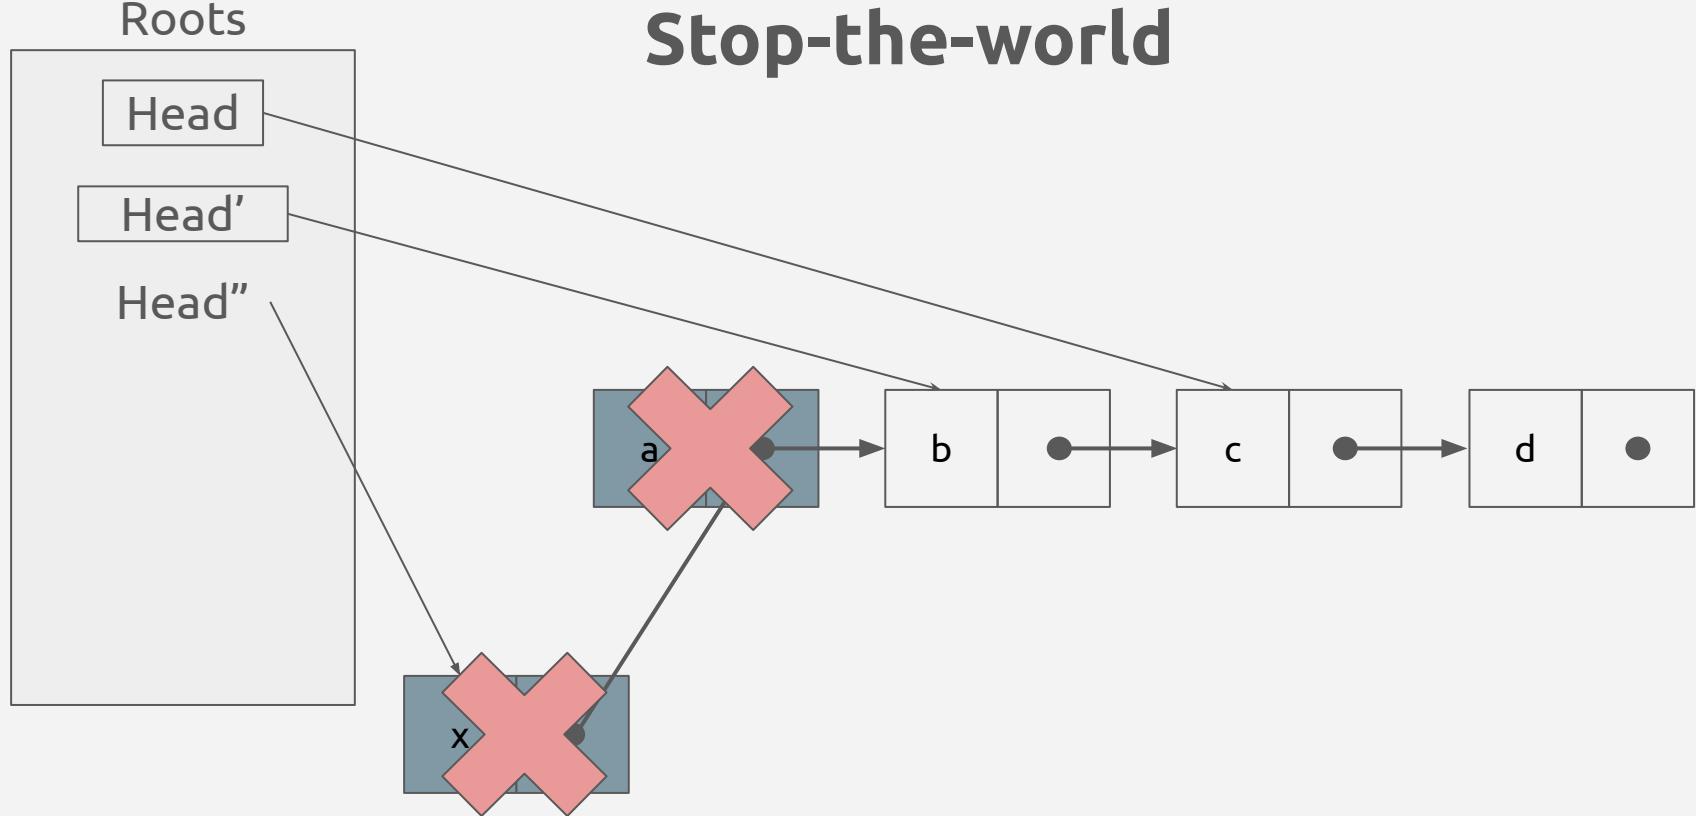
Não Utilizada



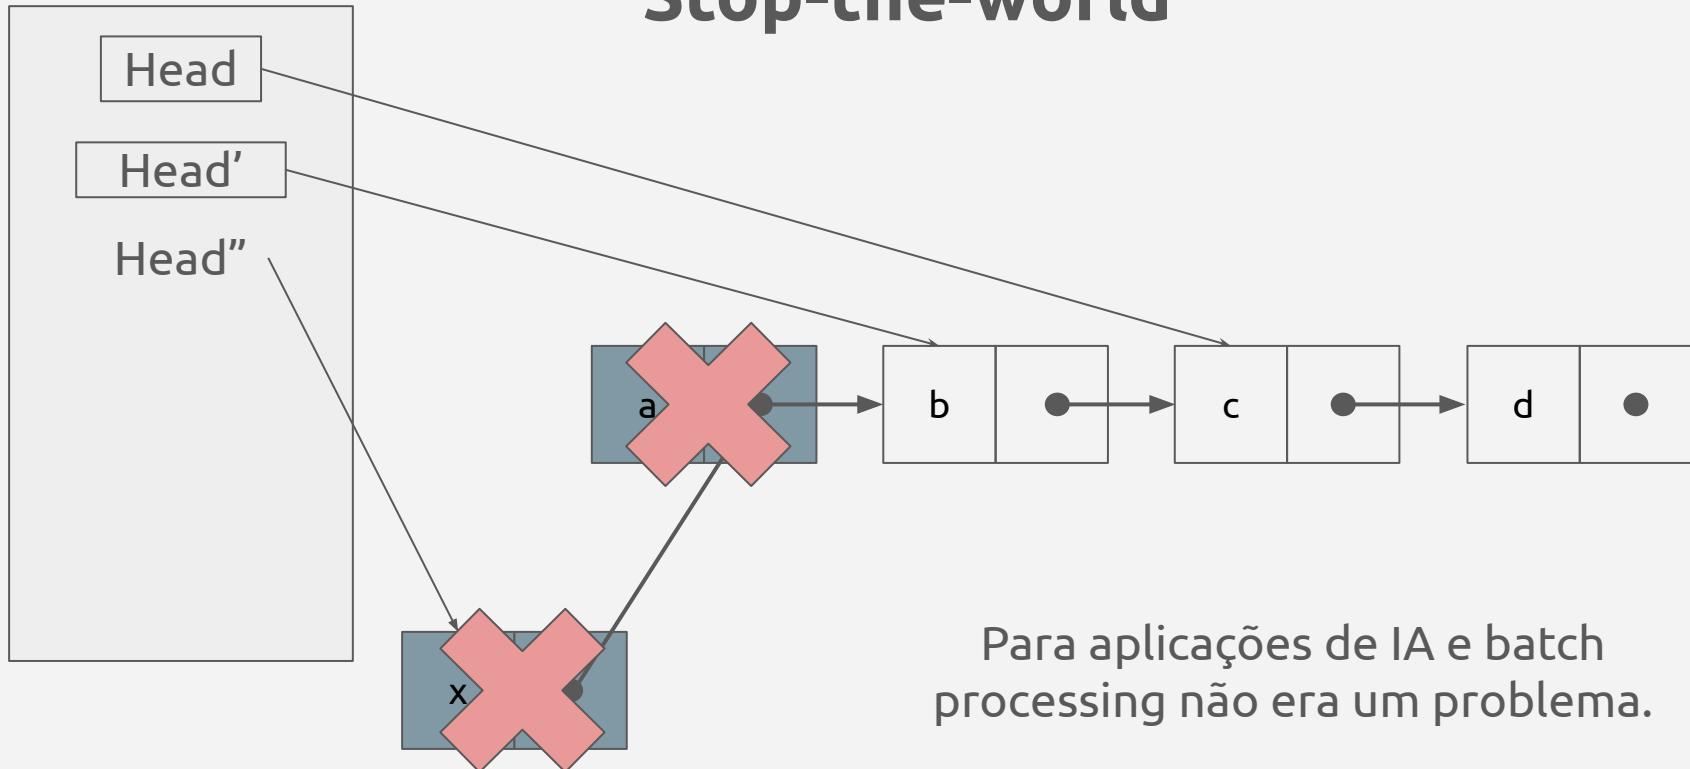
Stop-the-world



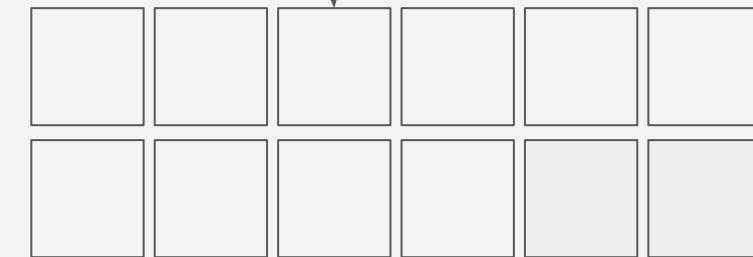
Stop-the-world



Stop-the-world

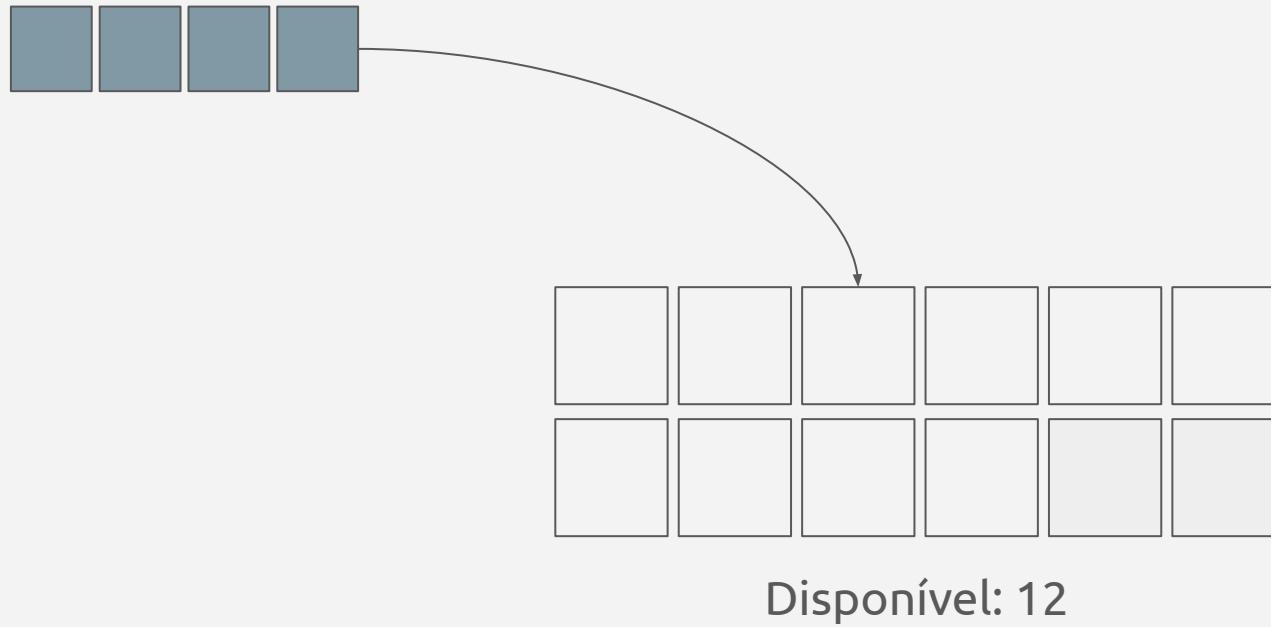


Outro problema...

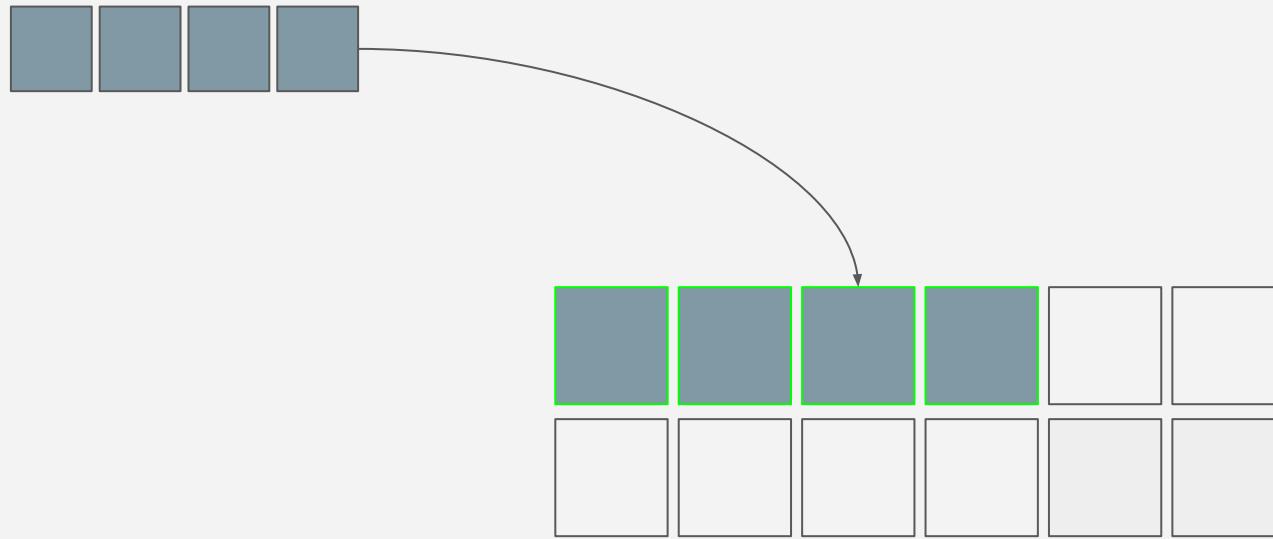


Disponível: 12

Alocando memória

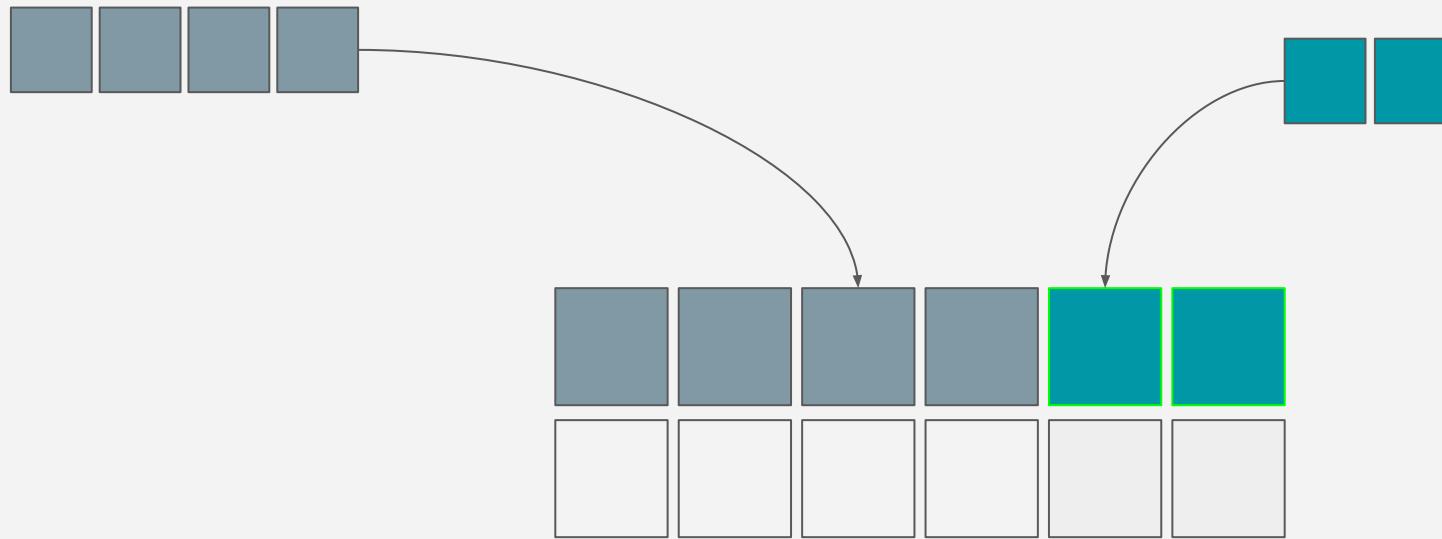


Alocando memória



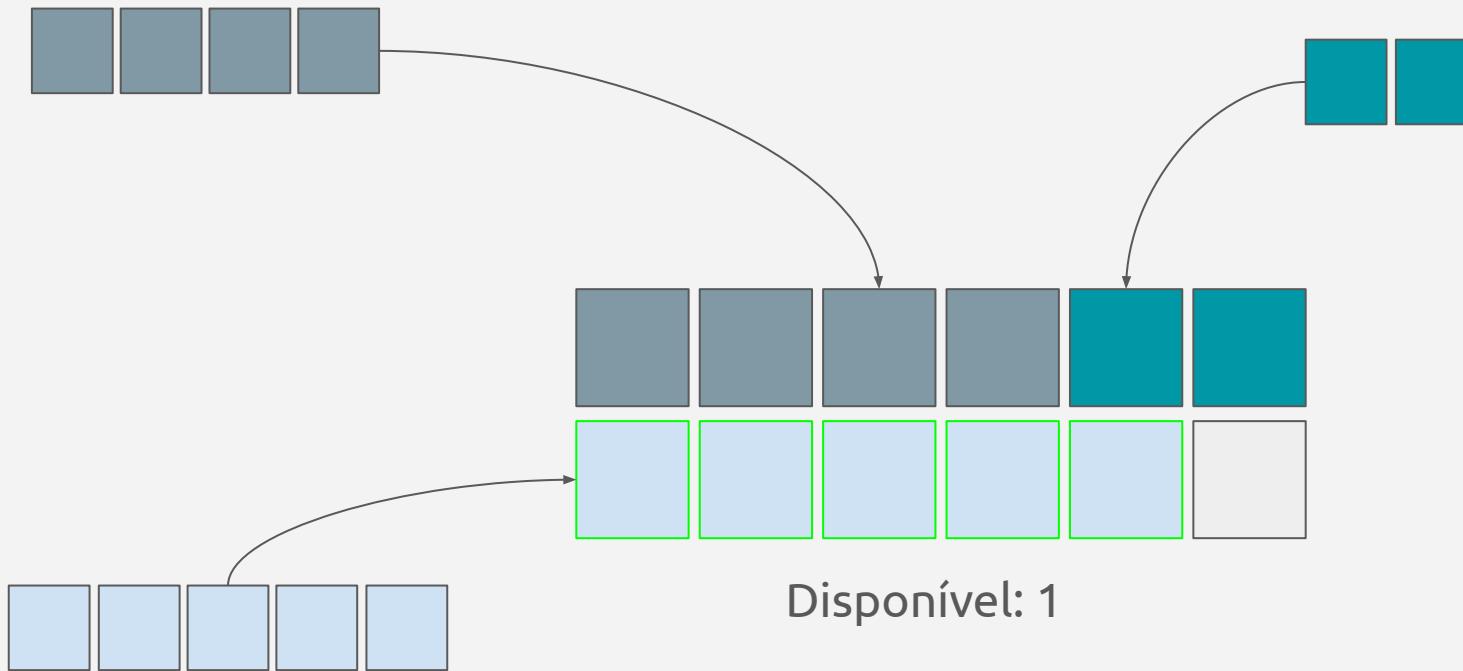
Disponível: 8

Alocando memória

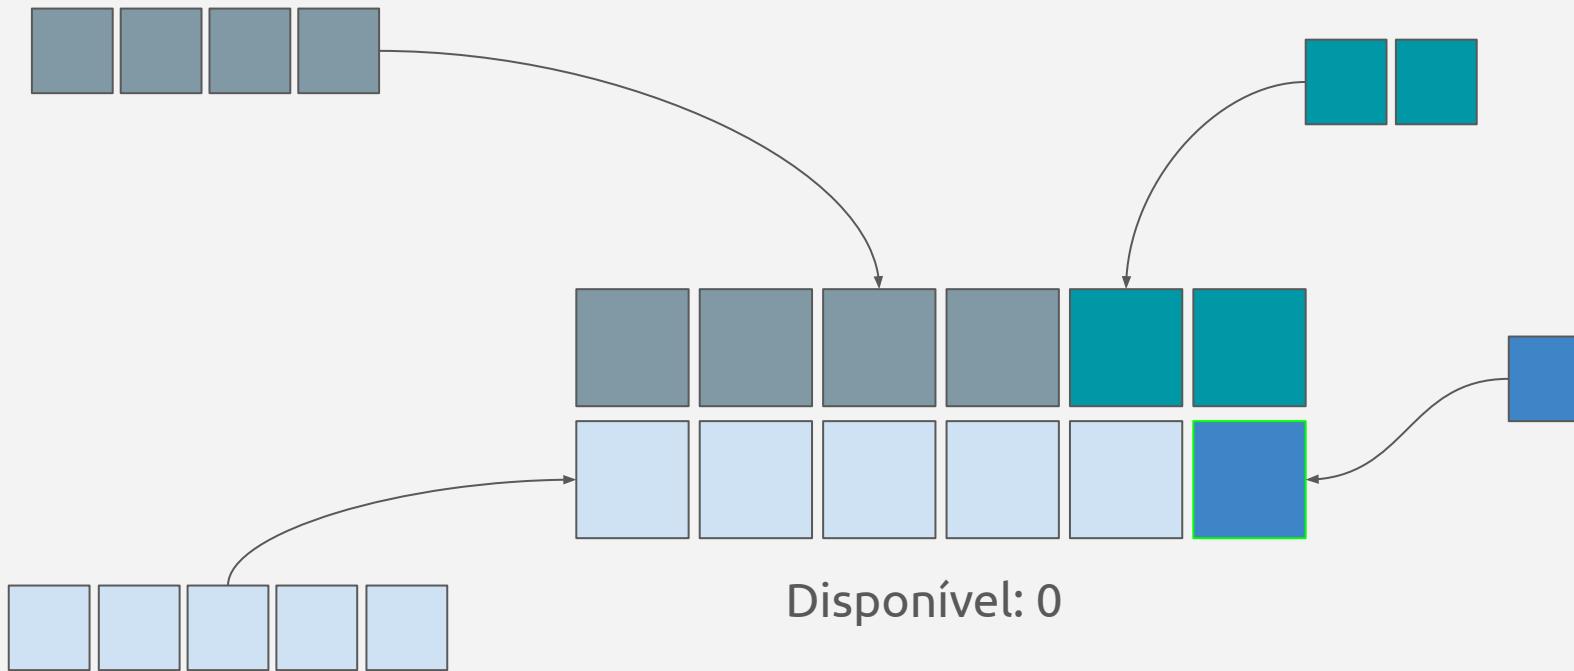


Disponível: 6

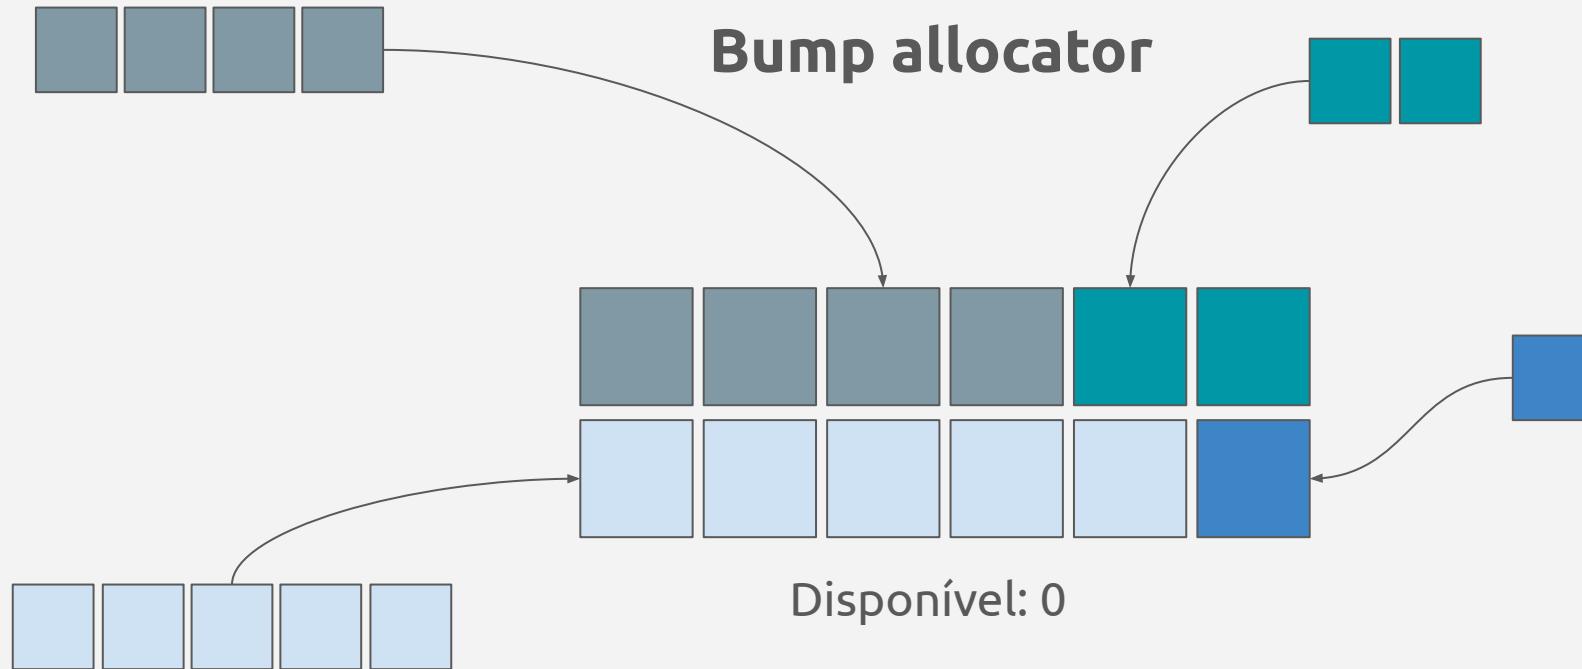
Alocando memória



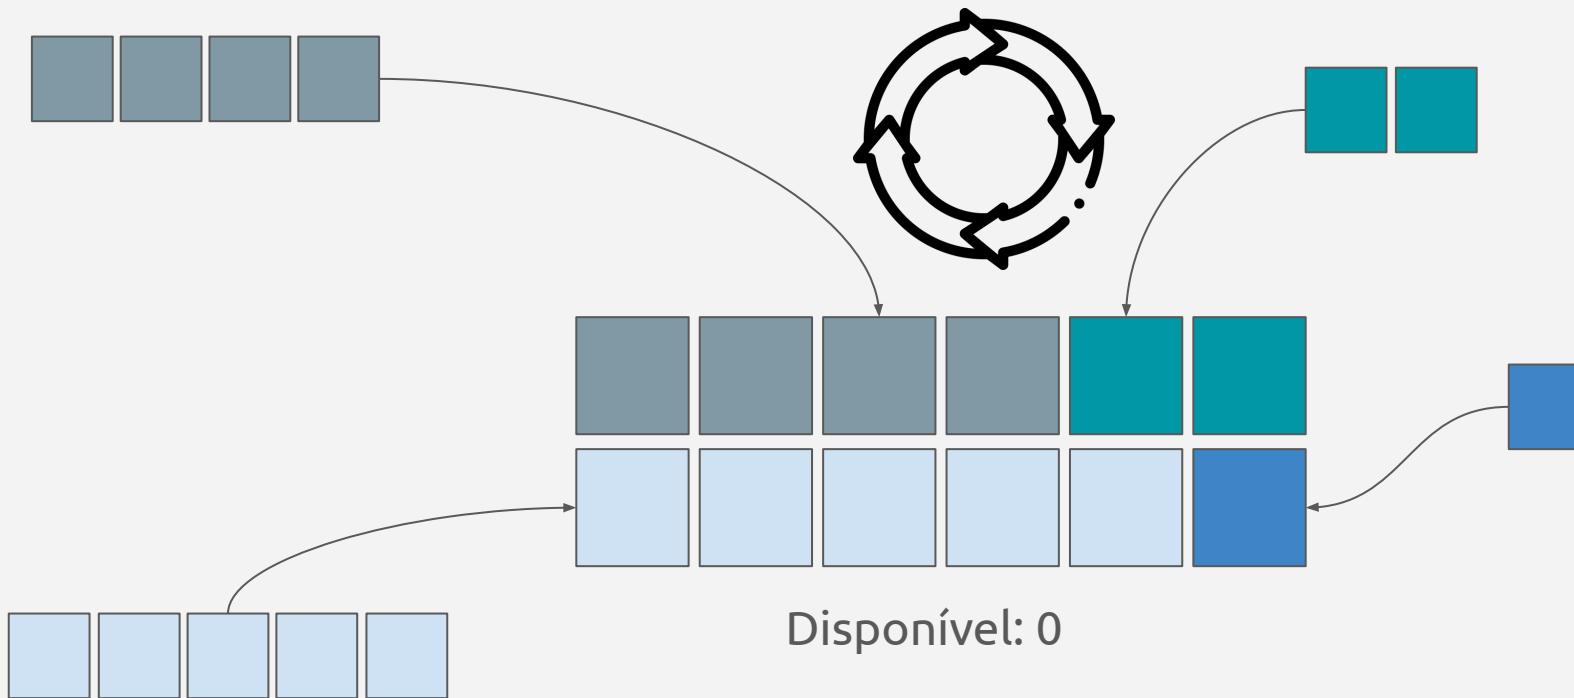
Alocando memória



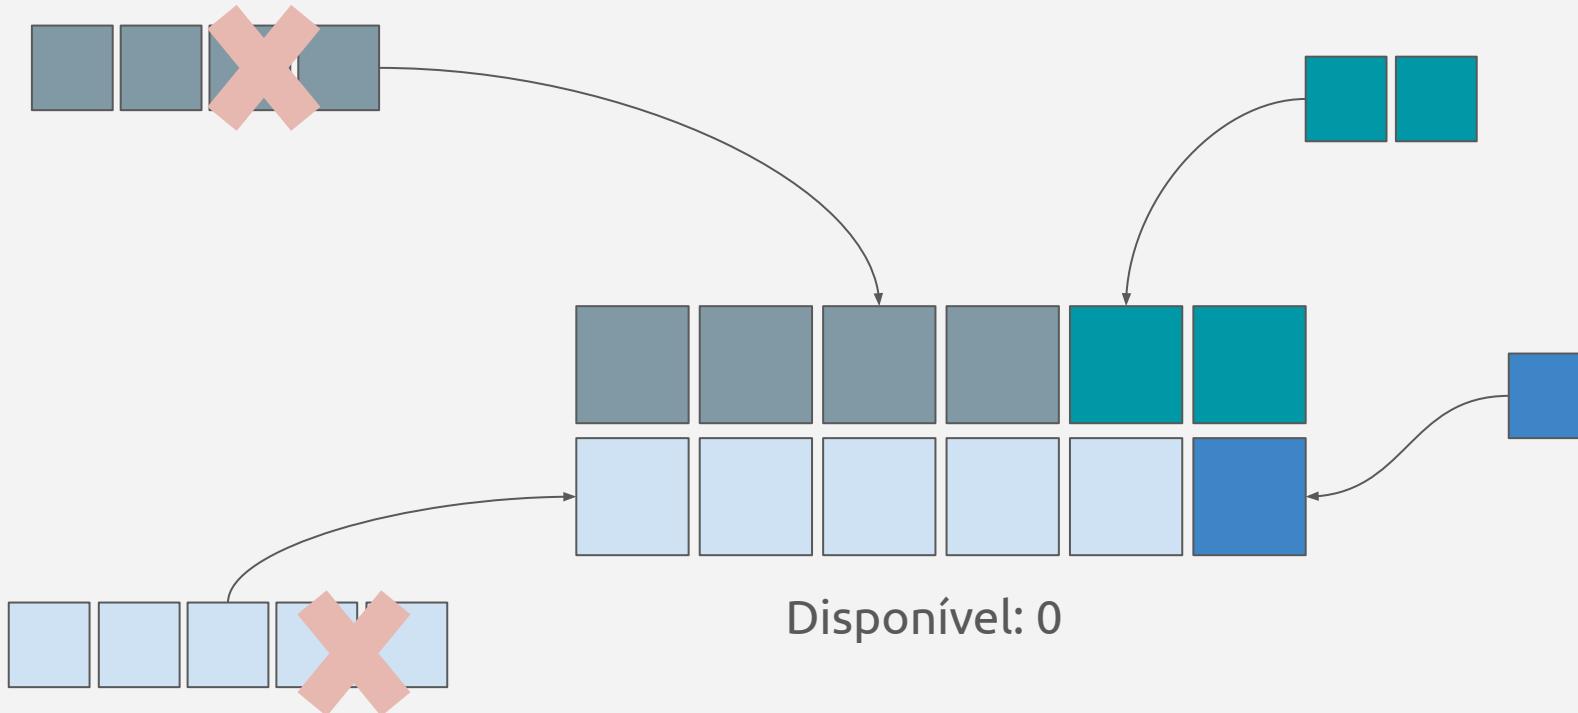
Alocando memória



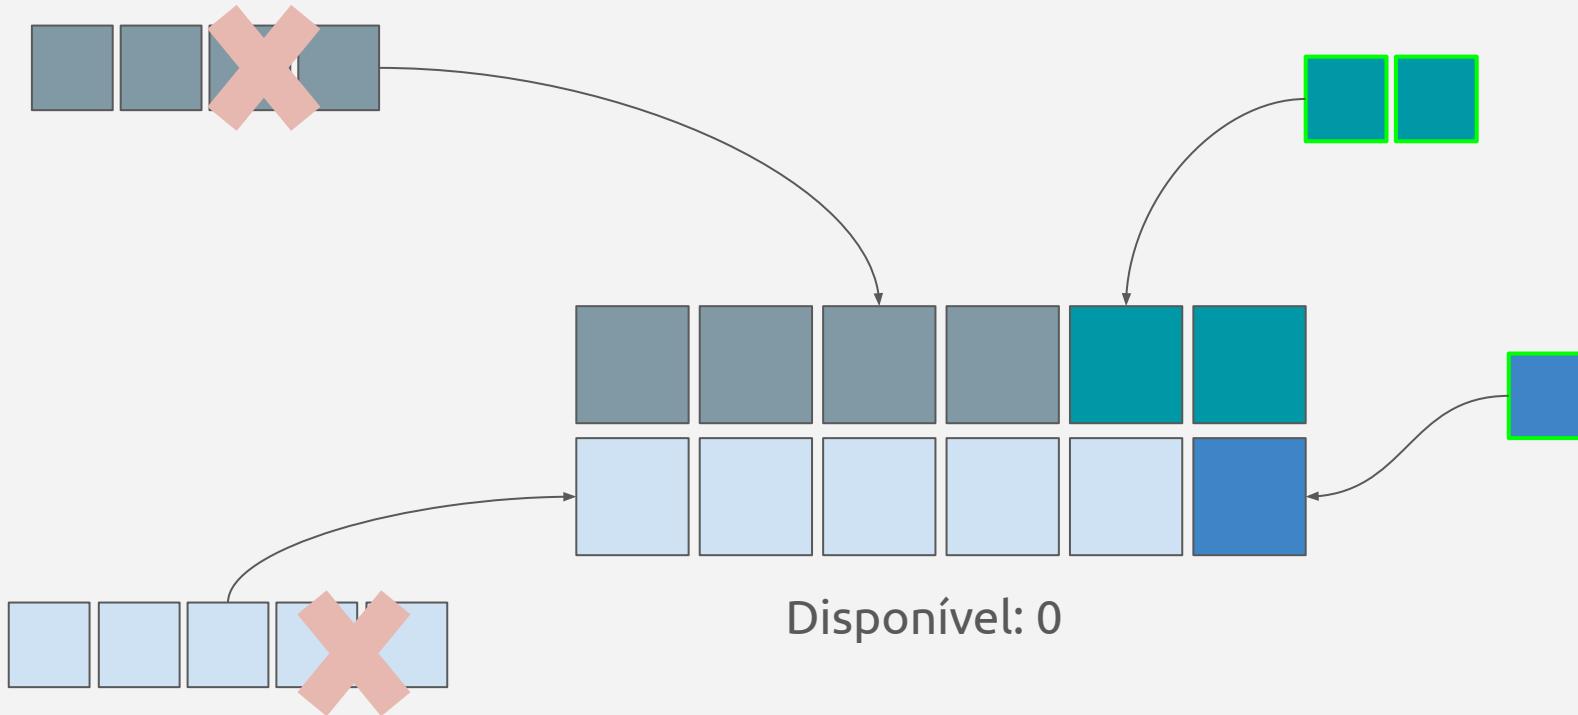
Coletando a memória



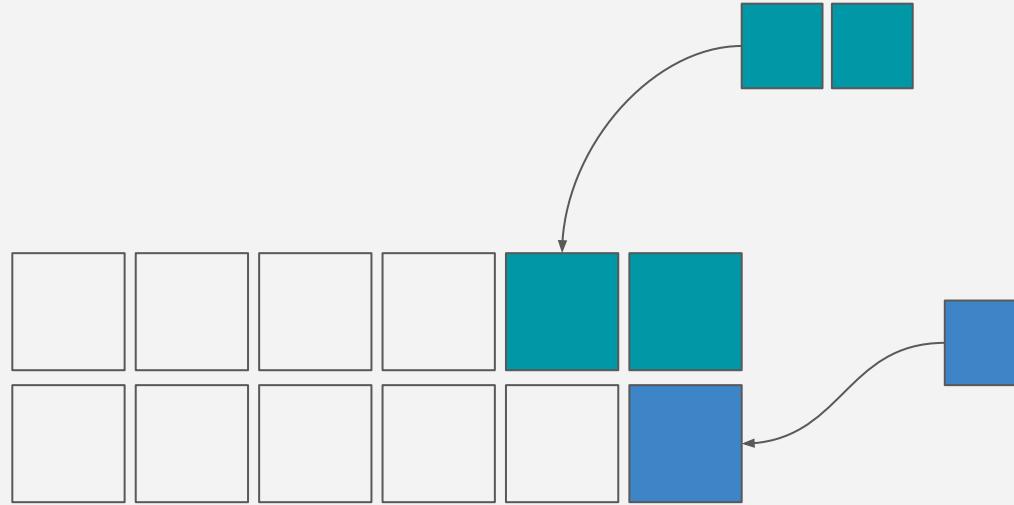
Coletando a memória



Mark

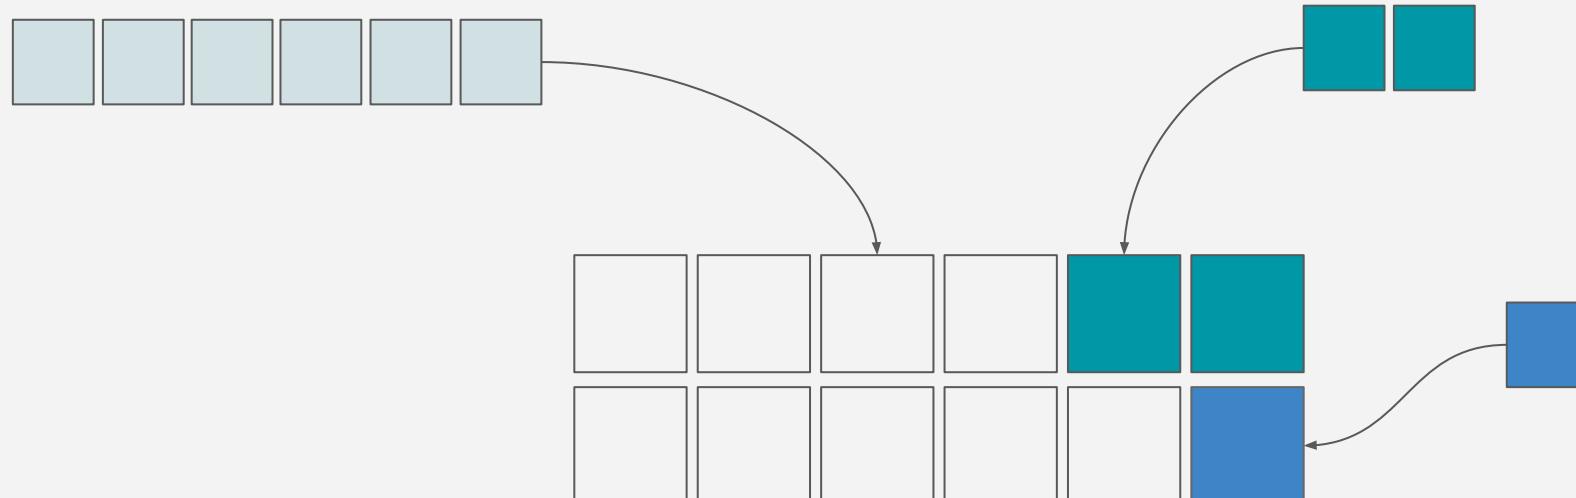


Sweep



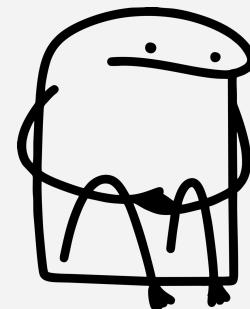
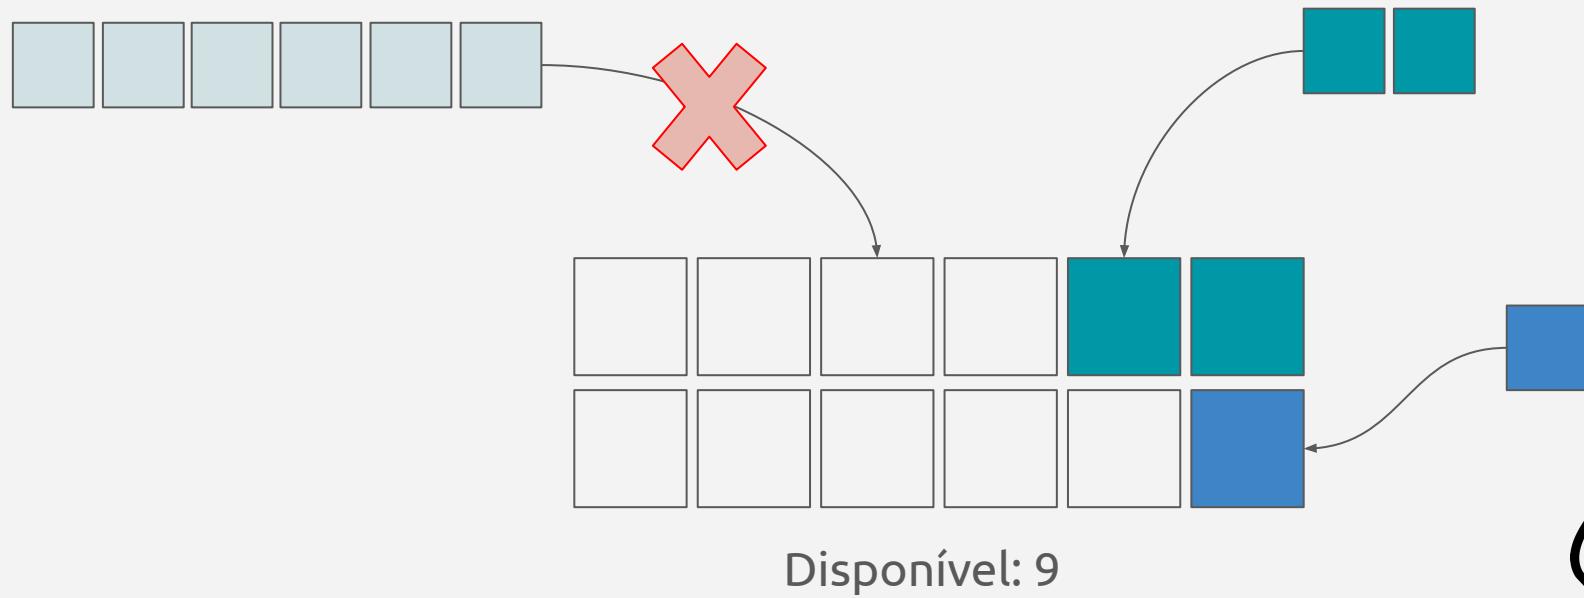
Disponível: 9

Alocando um novo valor

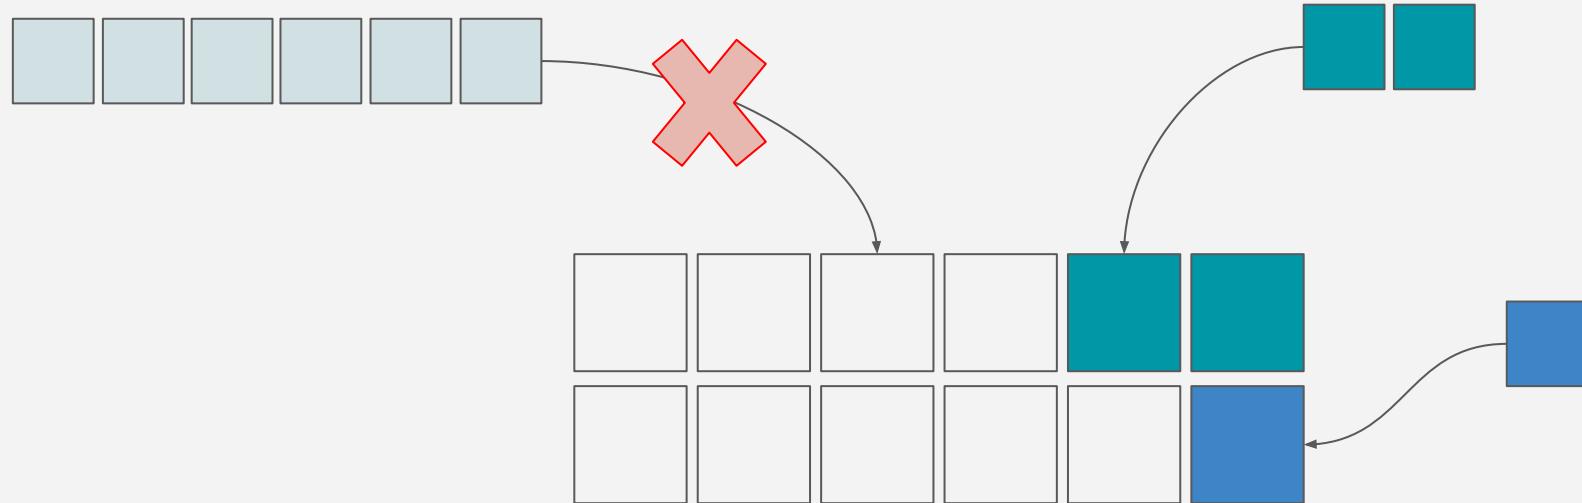


Disponível: 9

Alocando um novo valor



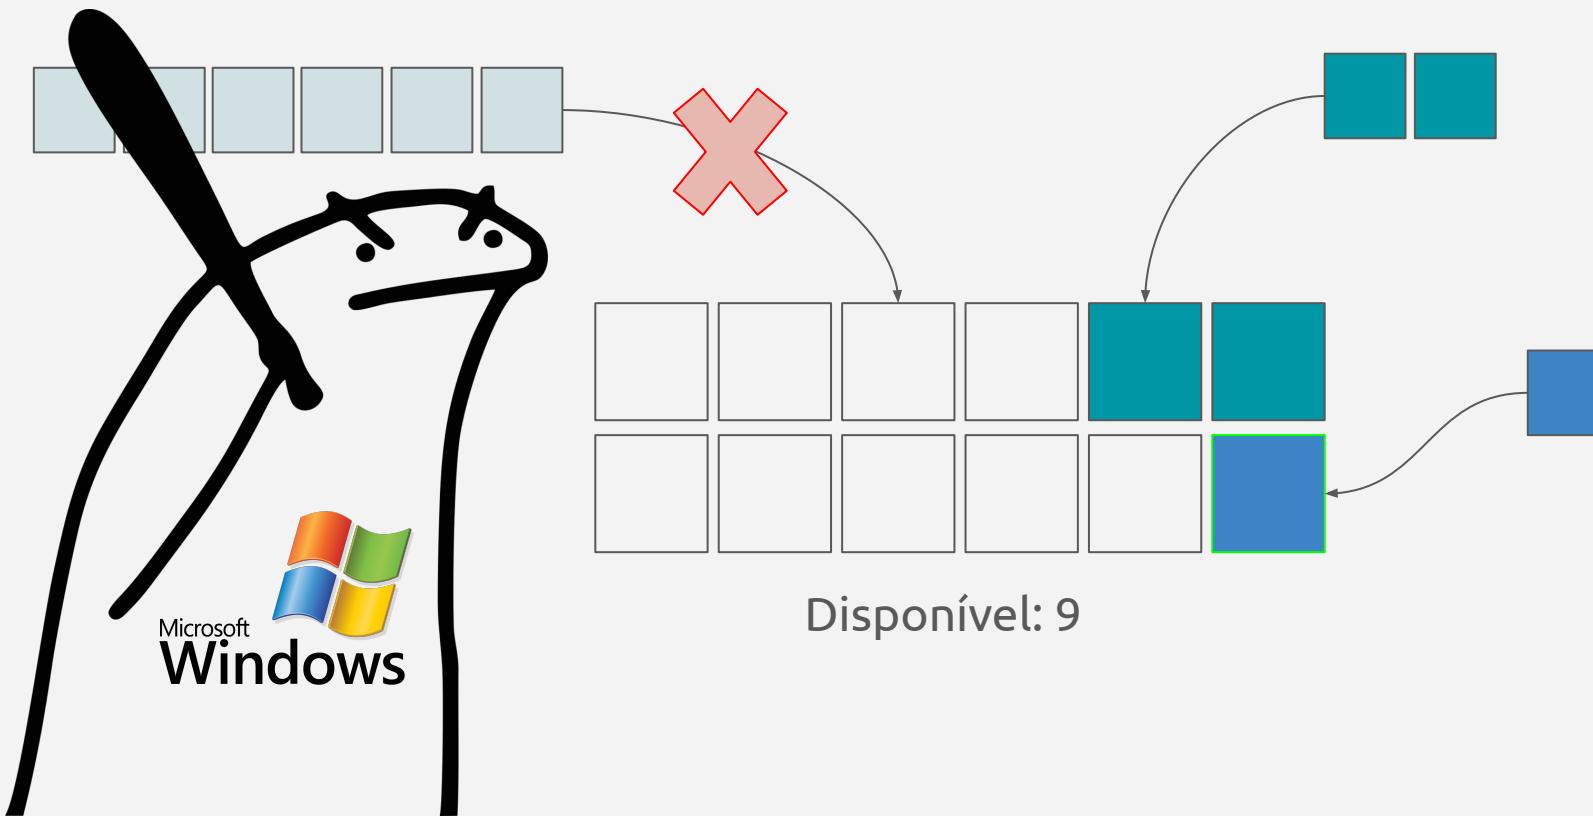
Fragmentação



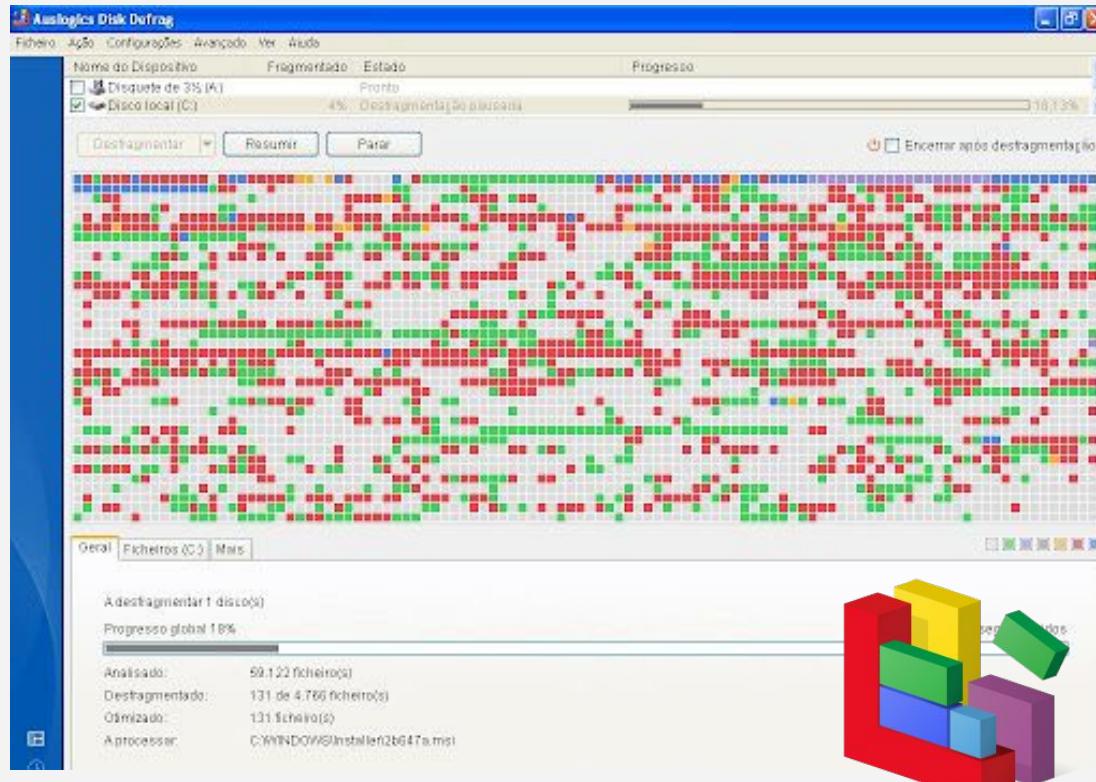
Disponível: 9



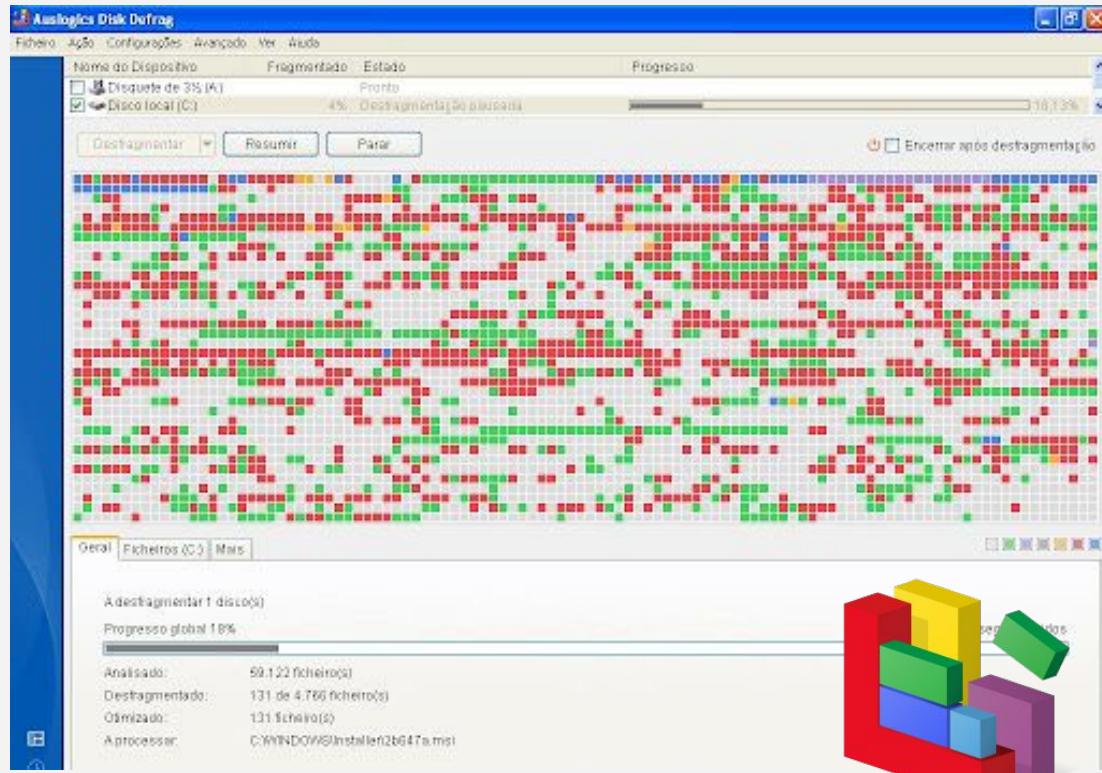
Fragmentação



(des)Fragmentação



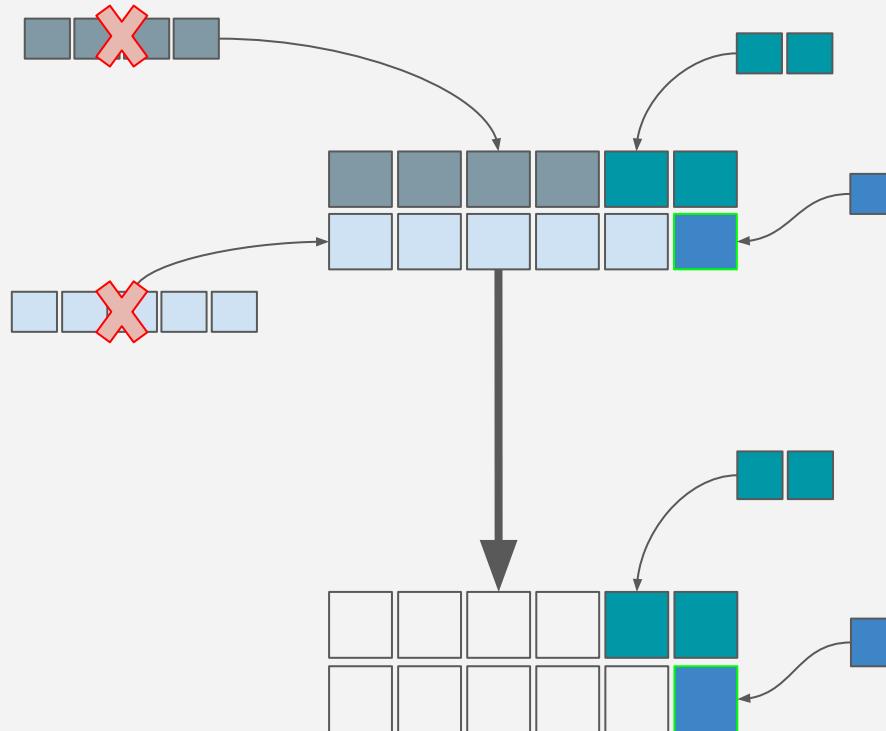
(des)Fragmentação



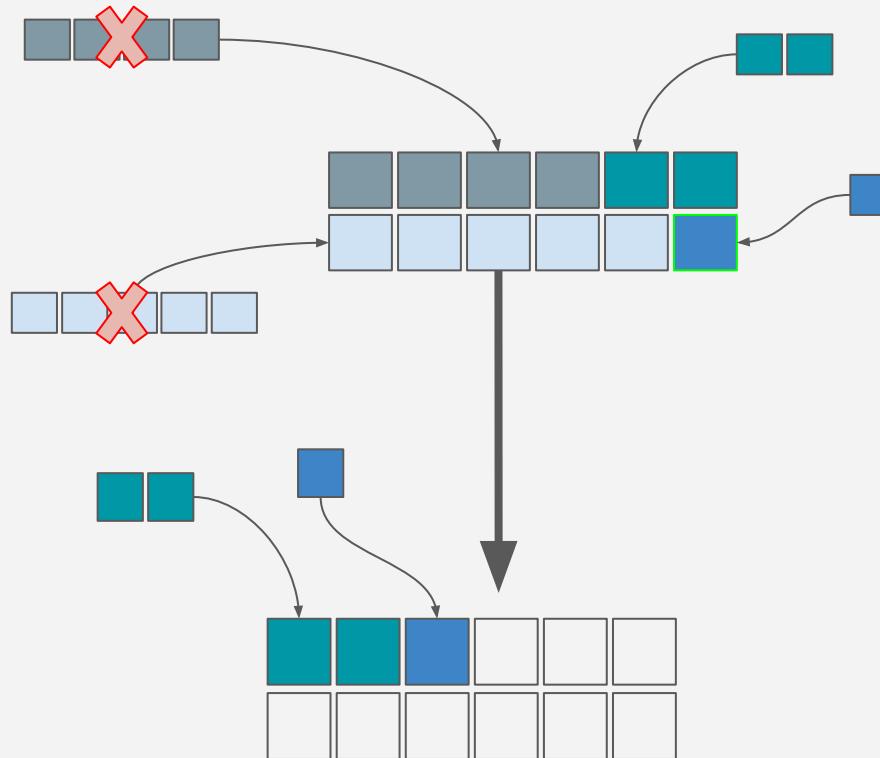
Move os dados para
regiões contíguas de
memória e elimina
espaços vazios.

Compactação

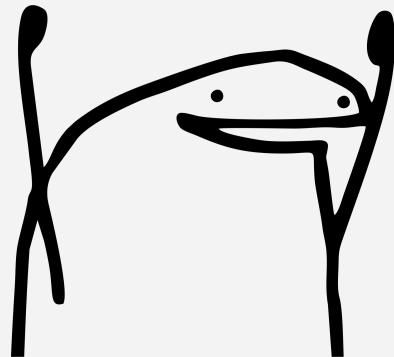
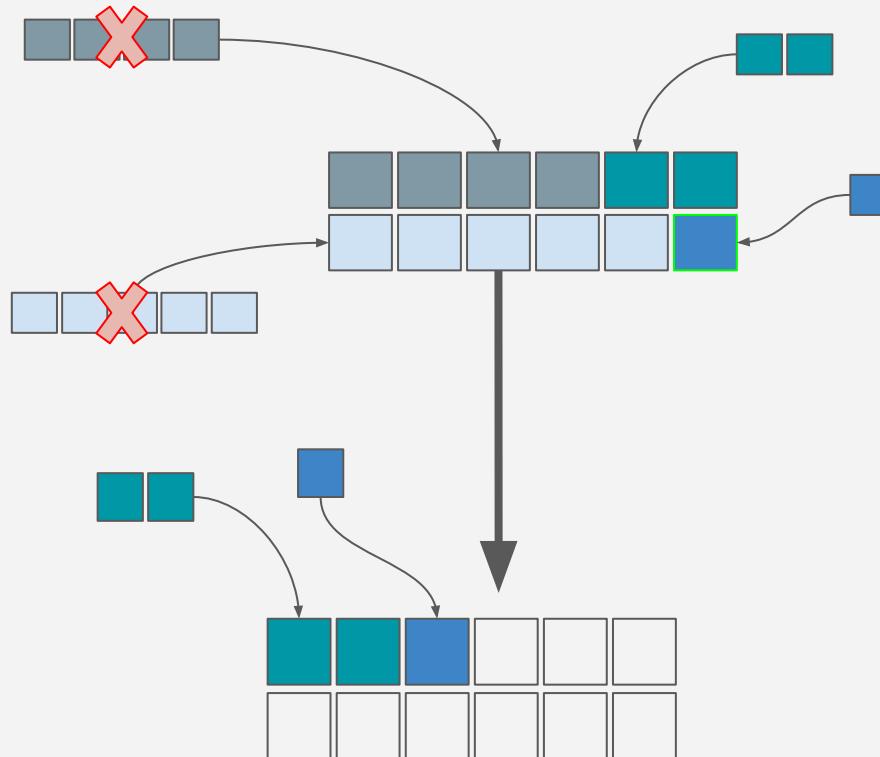
Compactação



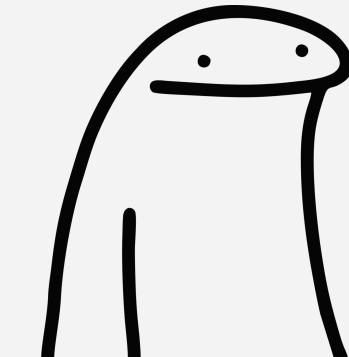
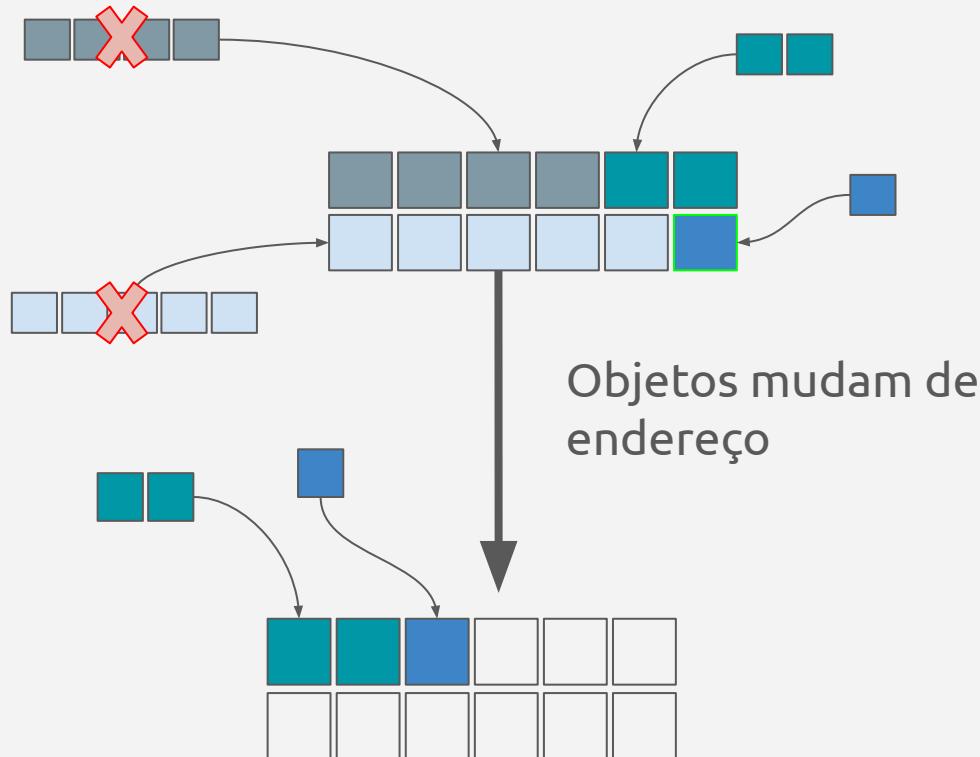
Compactação



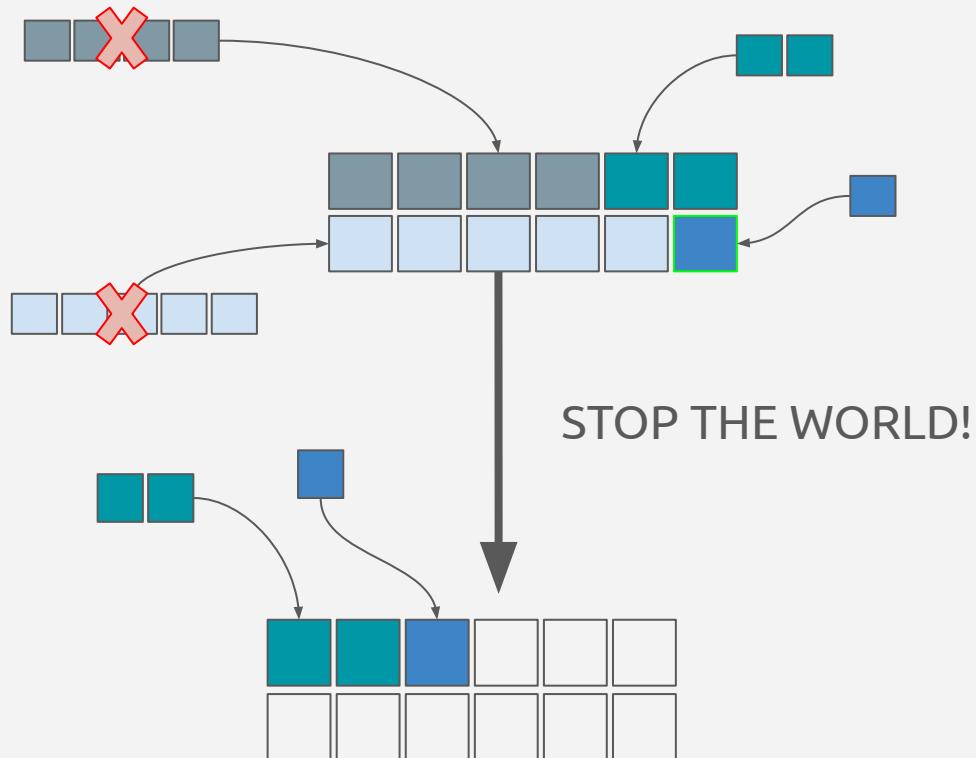
Compactação



Compactação



Compactação



Como tornar a compactação mais eficiente



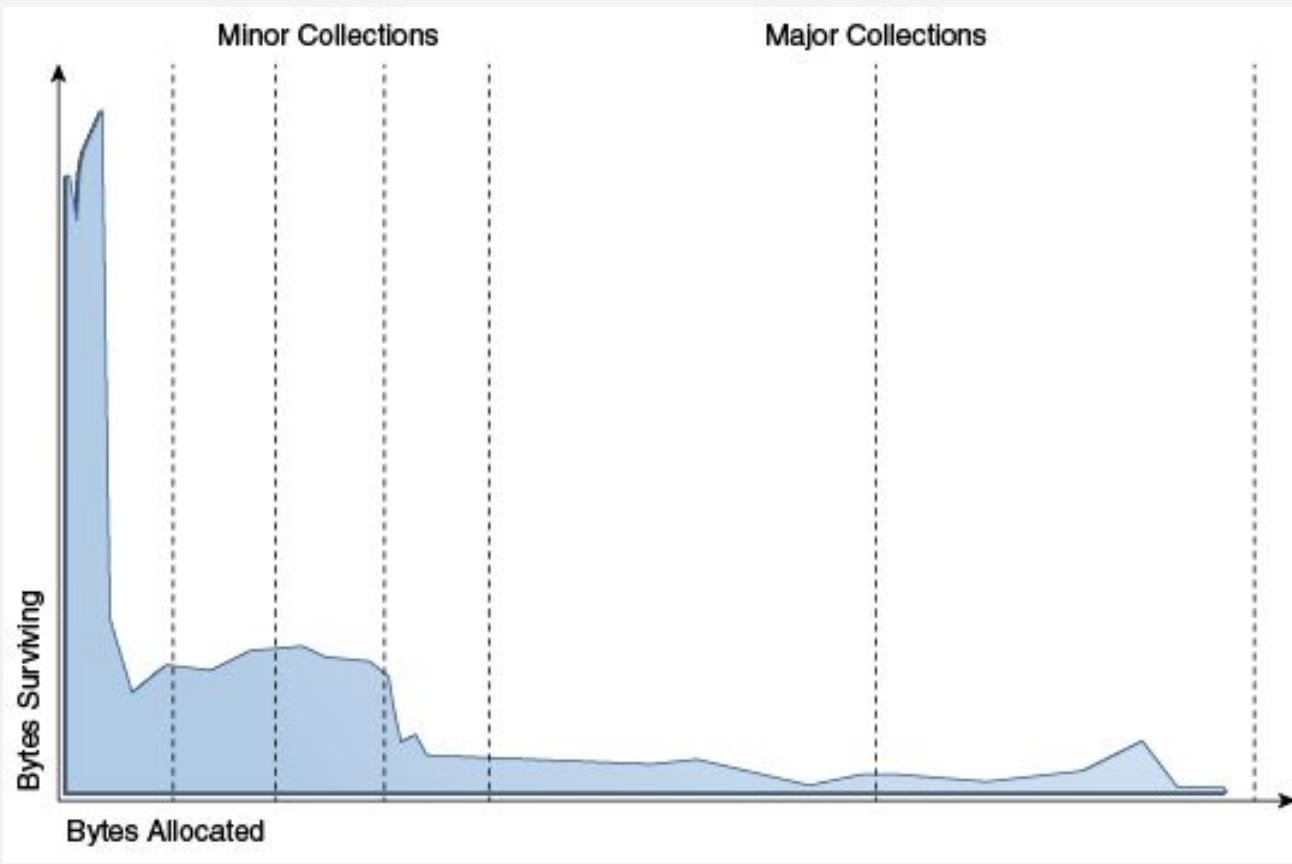
Weak Generational Hypothesis

**Generation Scavenging: A Non-disruptive High Performance
Storage Reclamation Algorithm**

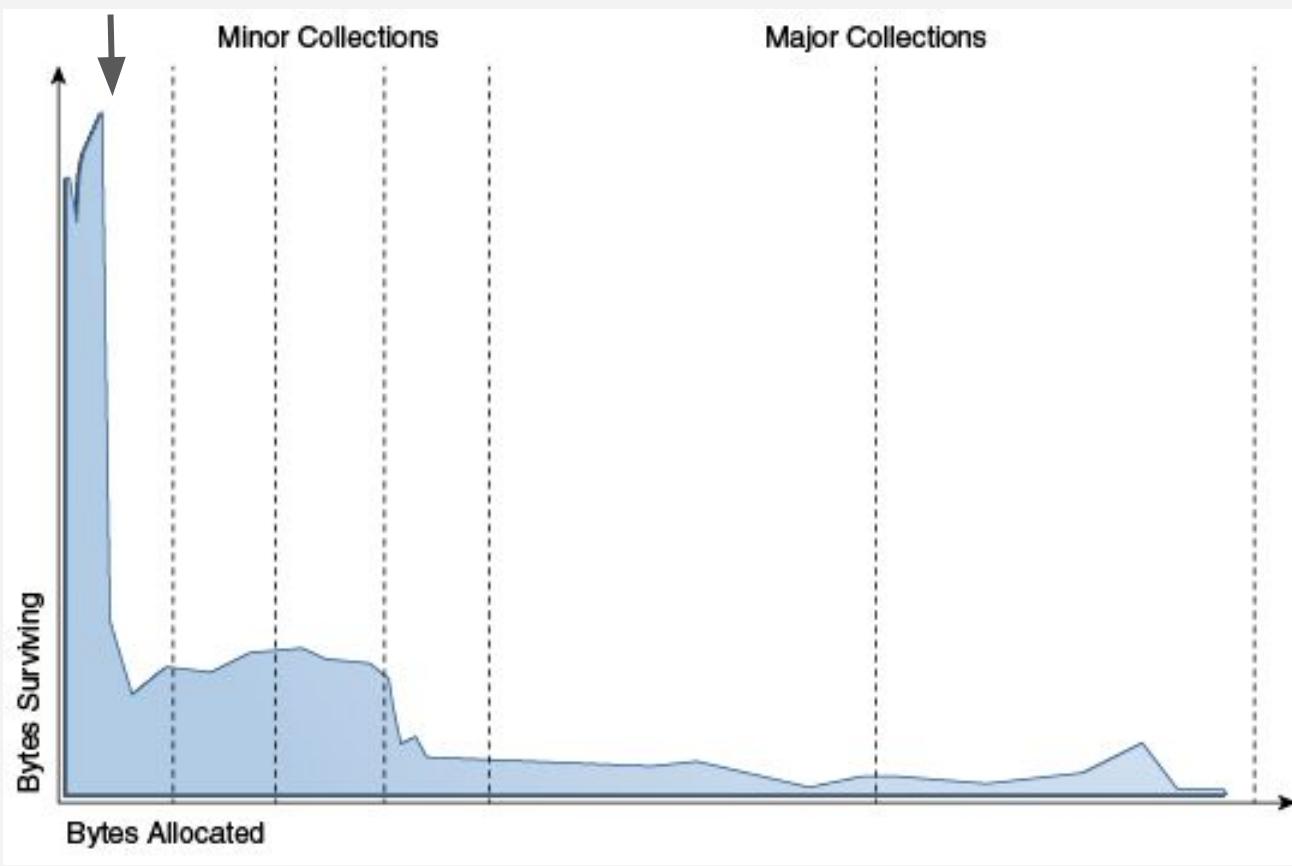
David Ungar

Computer Science Division
Department of Electrical Engineering and Computer Sciences
University of California
Berkeley, California 94720

Weak Generational Hypothesis

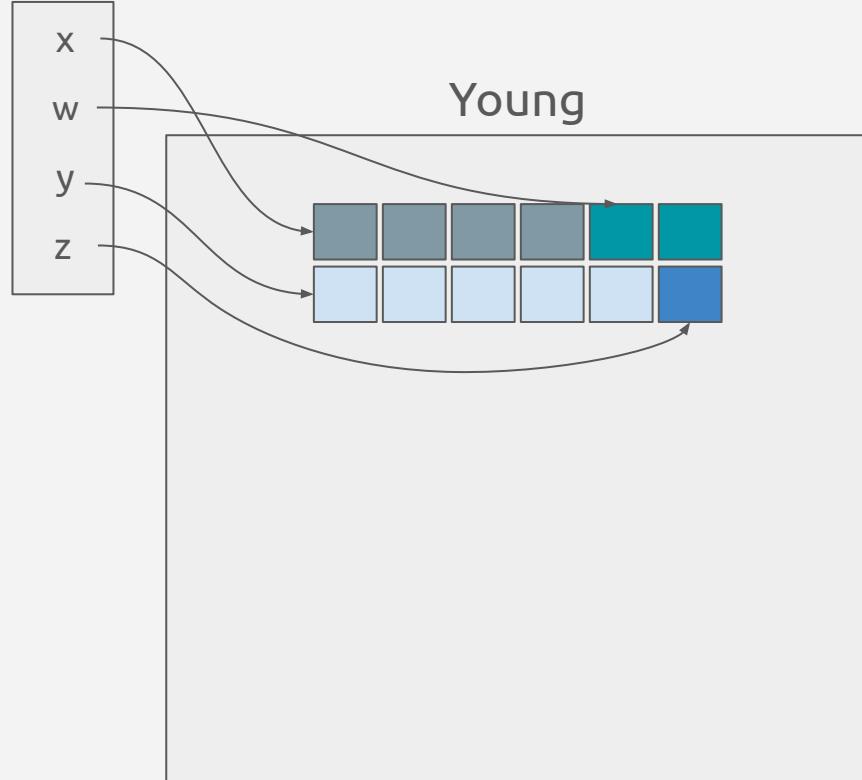


Weak Generational Hypothesis



GC geracional

Roots

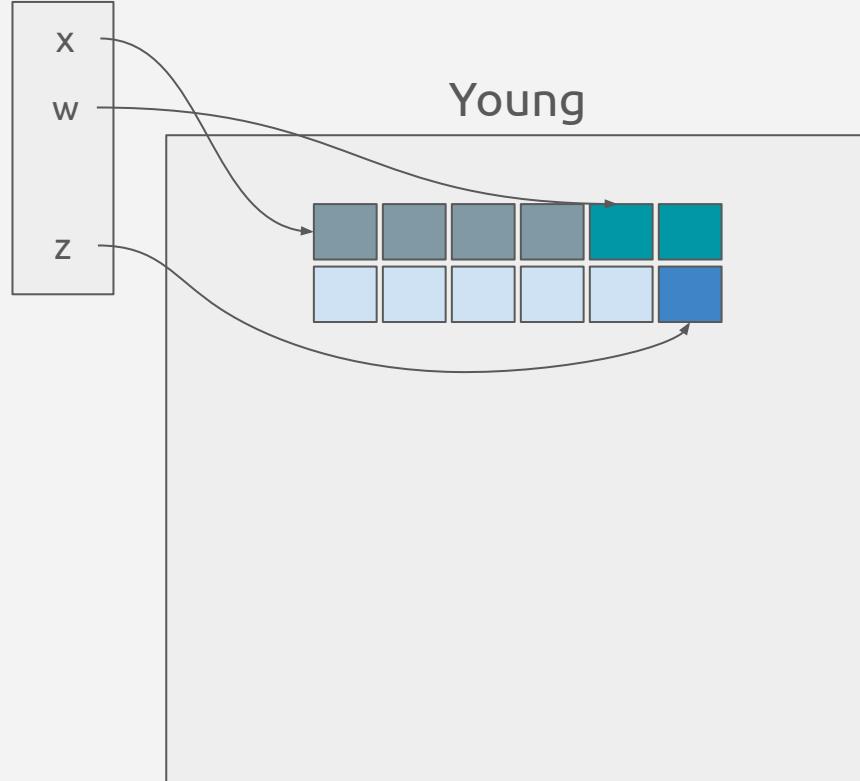


Young

Old

Variável y sai de escopo

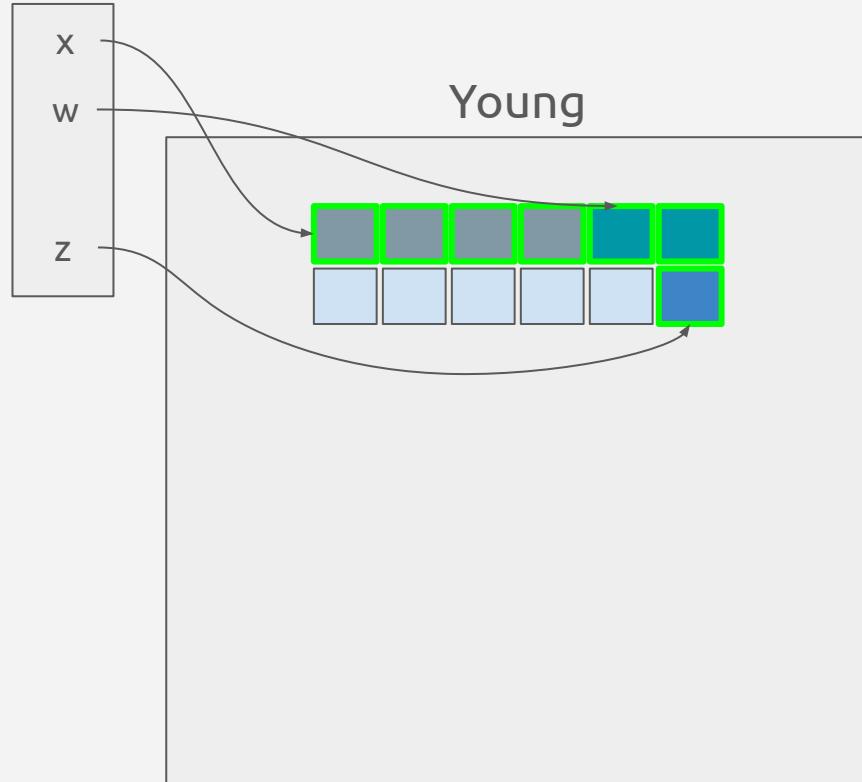
Roots



Old

Mark

Roots

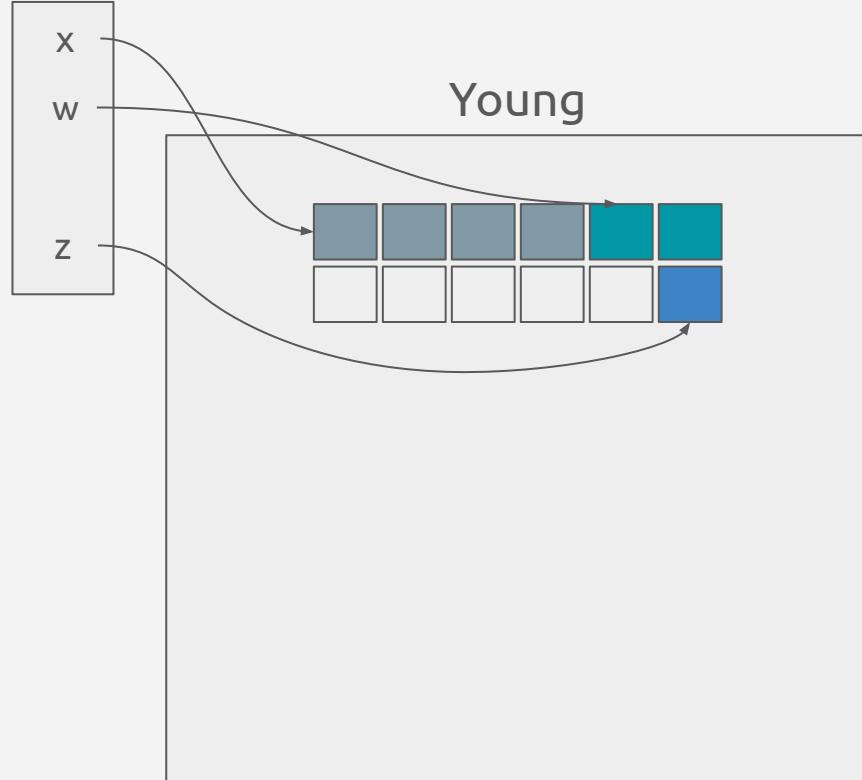


Young

Old

Sweep

Roots

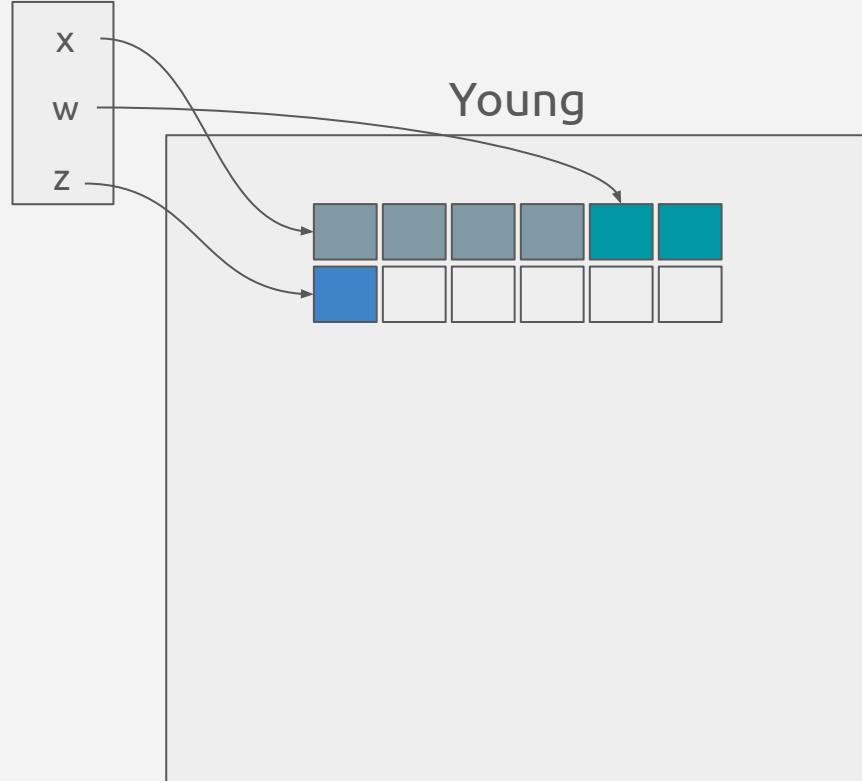


Young

Old

Compact

Roots

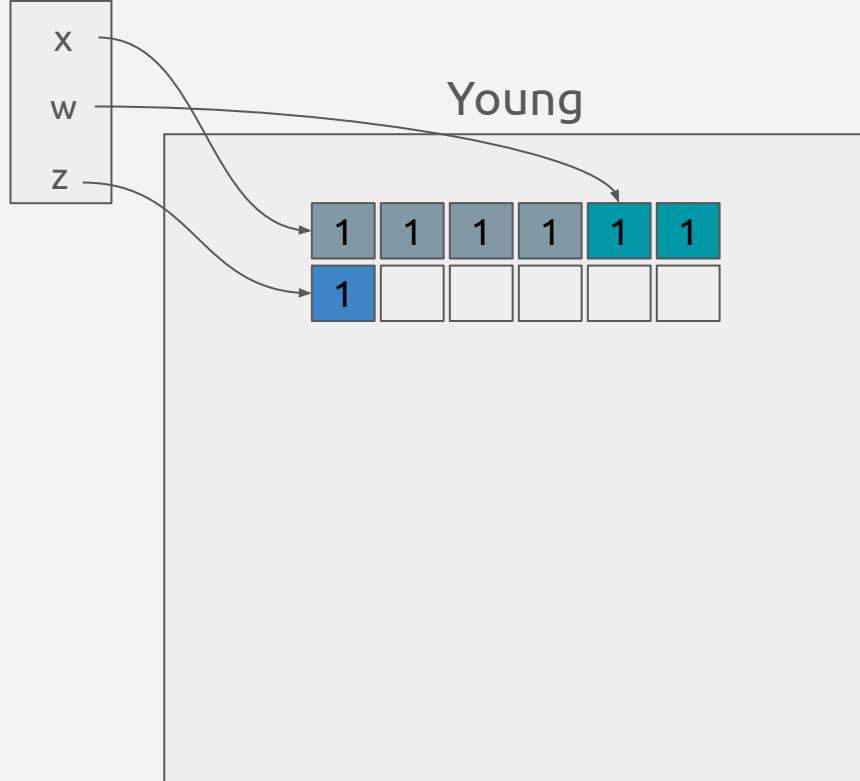


Young

Old

Marca geração

Roots

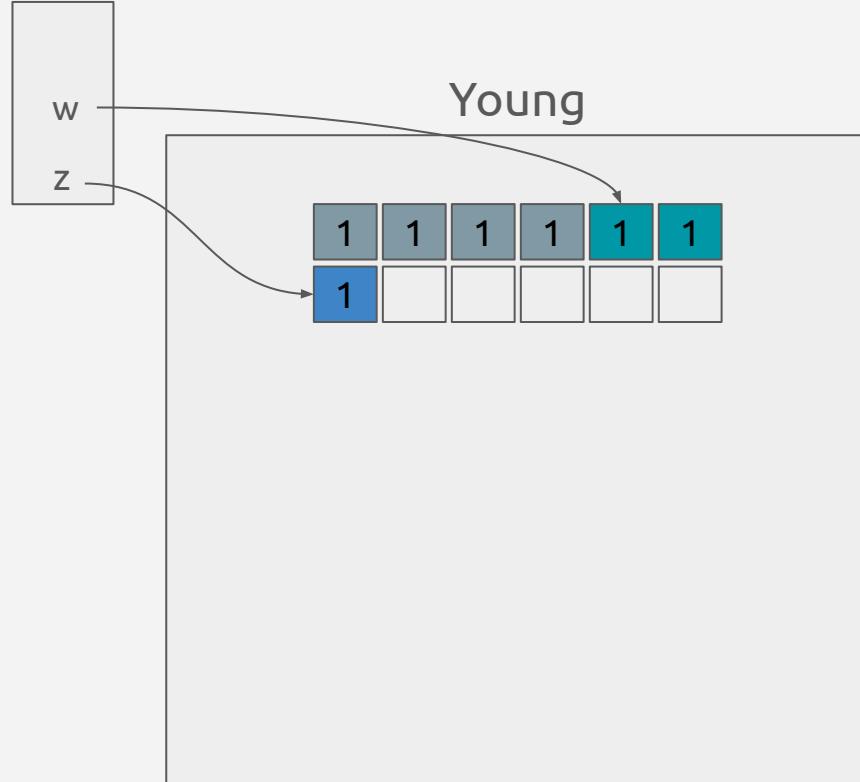


Young

Old

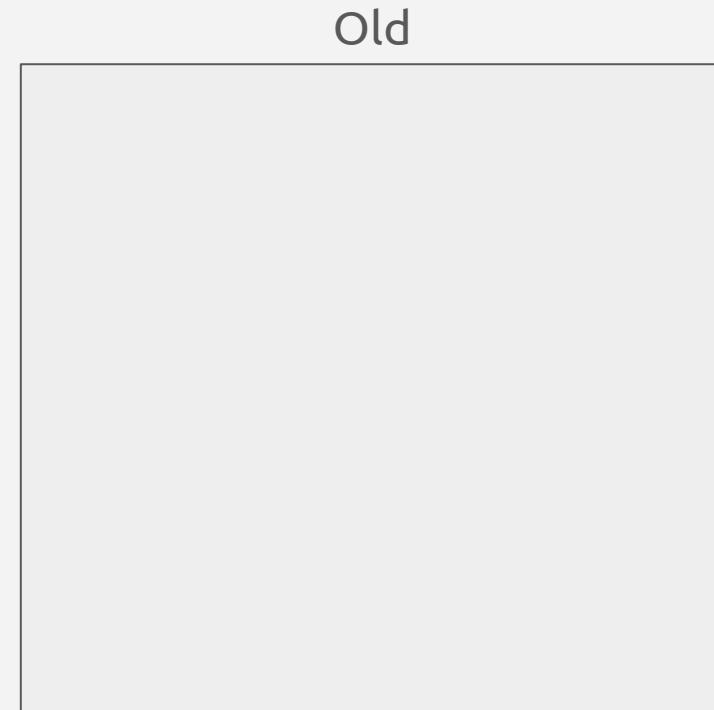
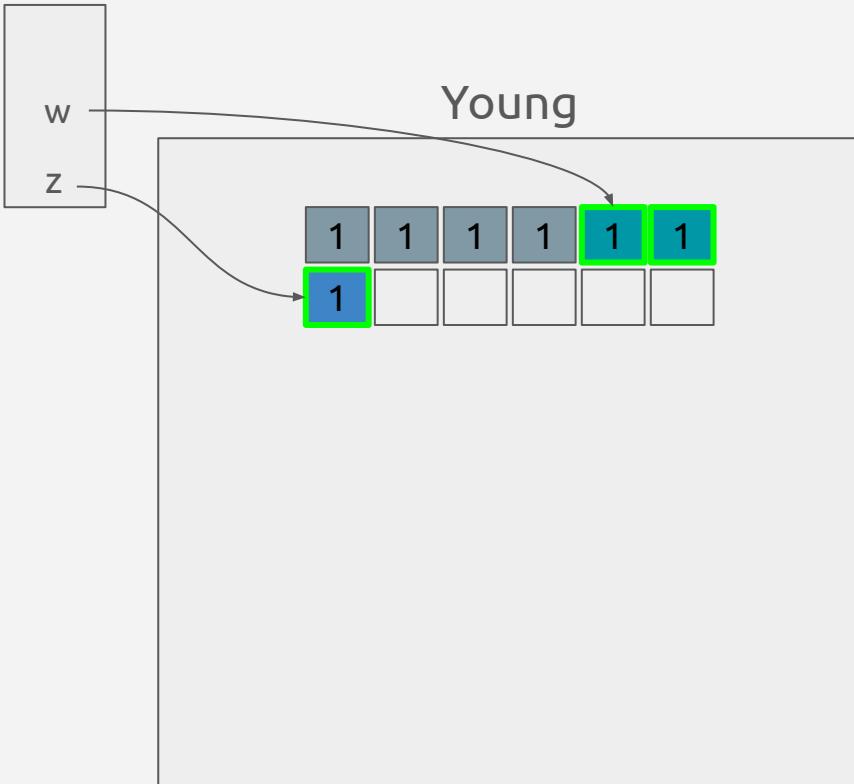
Variável x sai do escopo

Roots



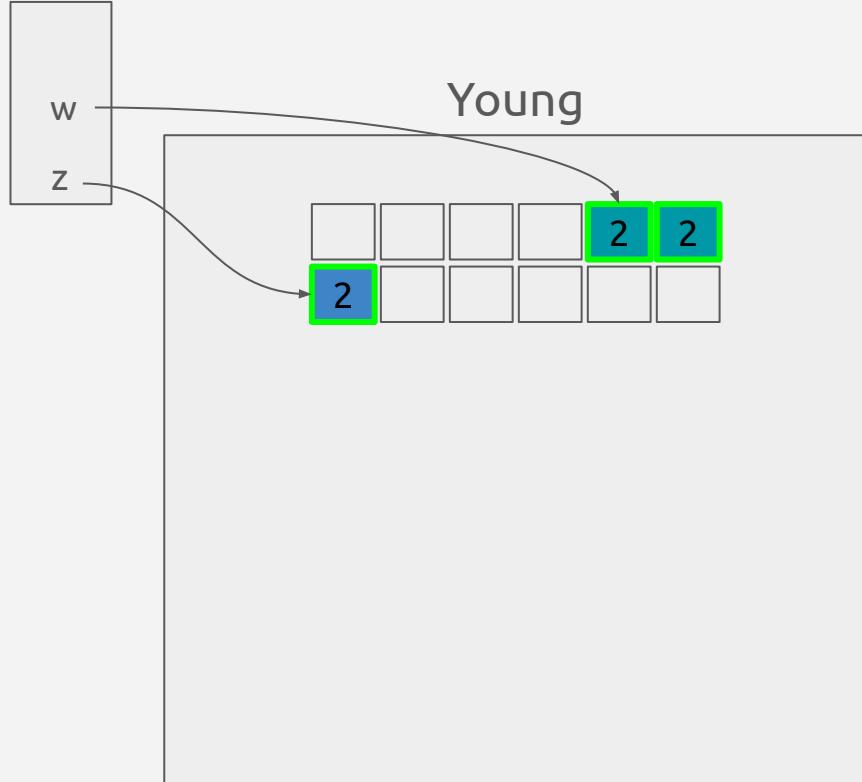
Roots

Mark

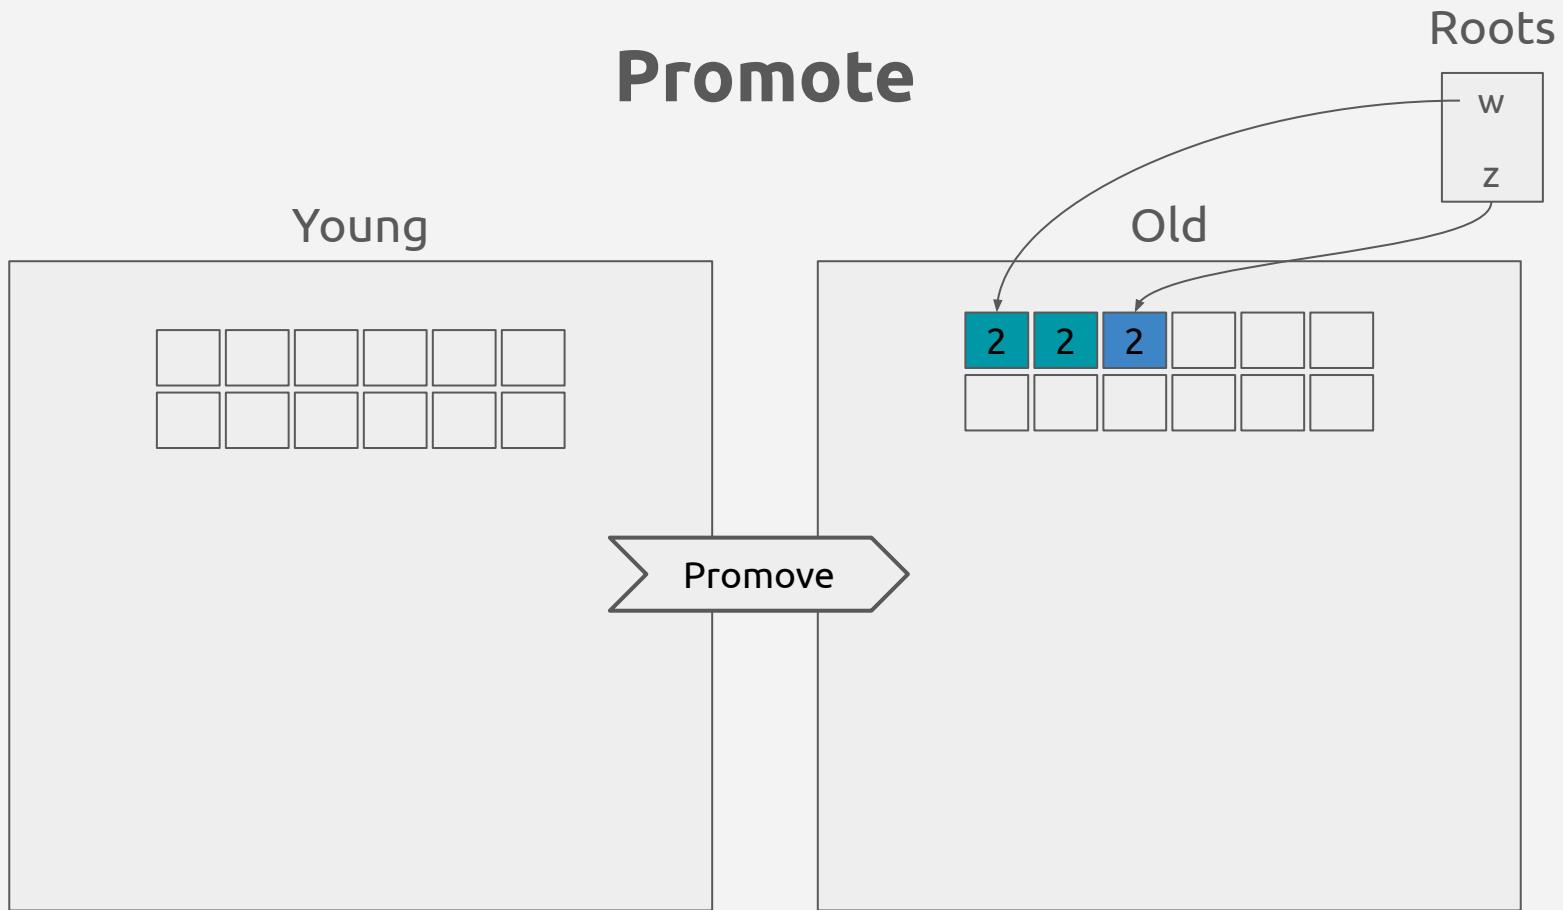


Sweep

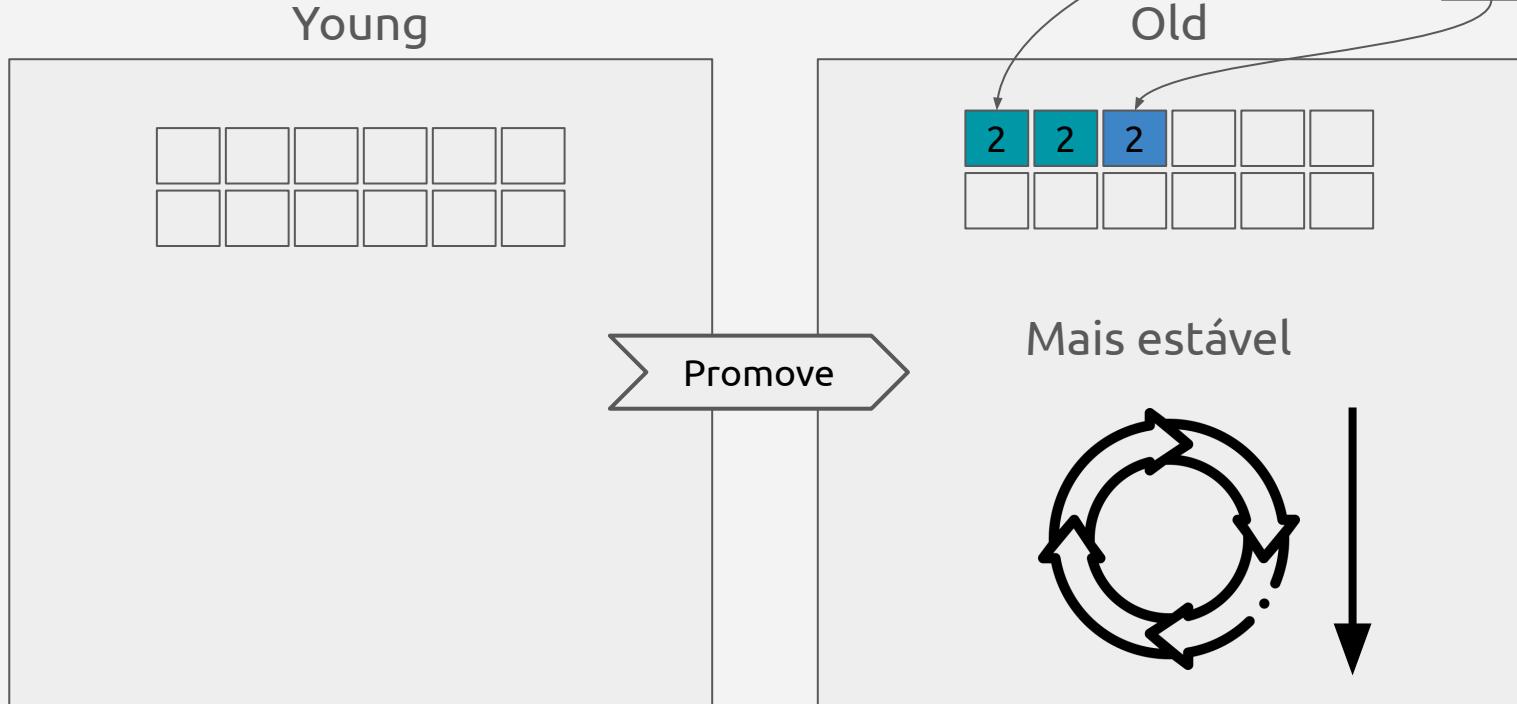
Roots



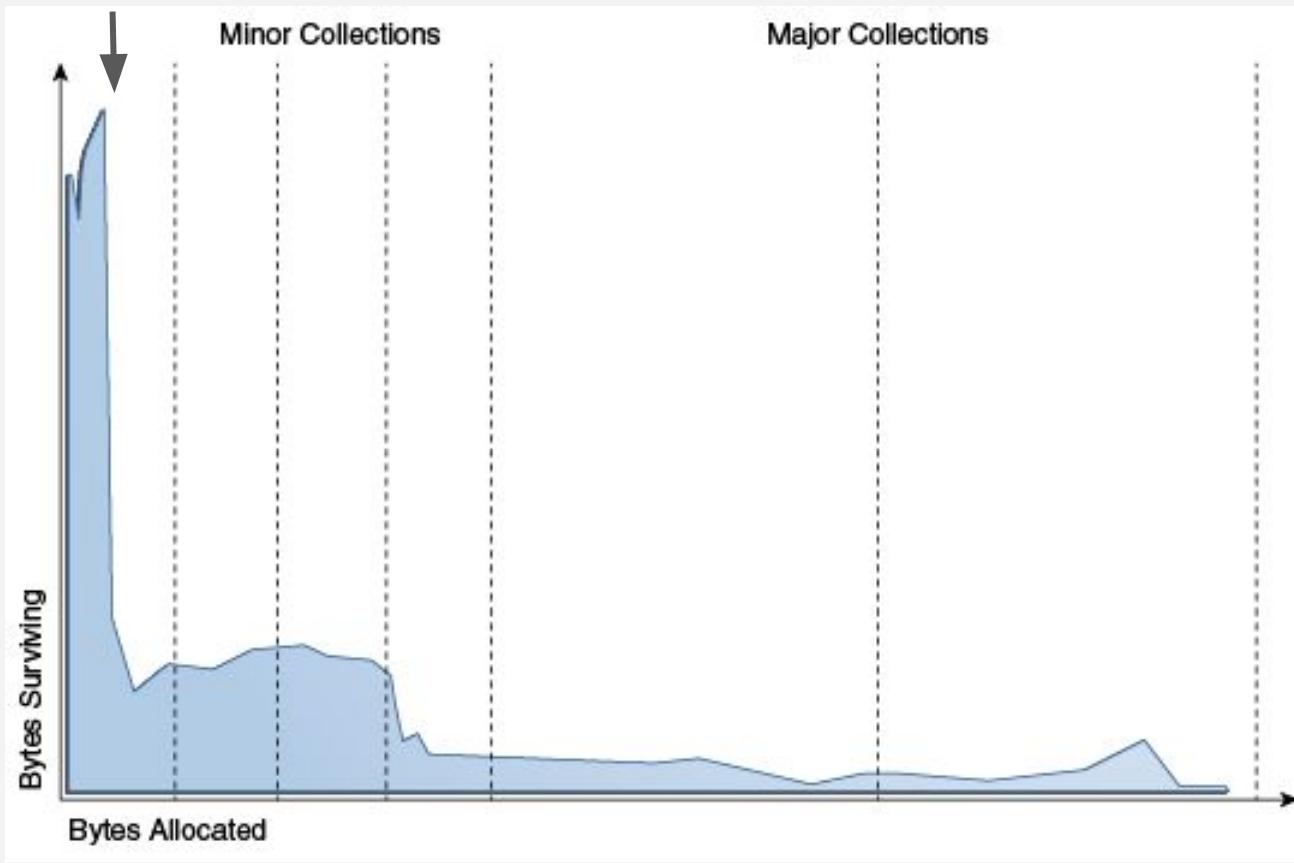
Promote



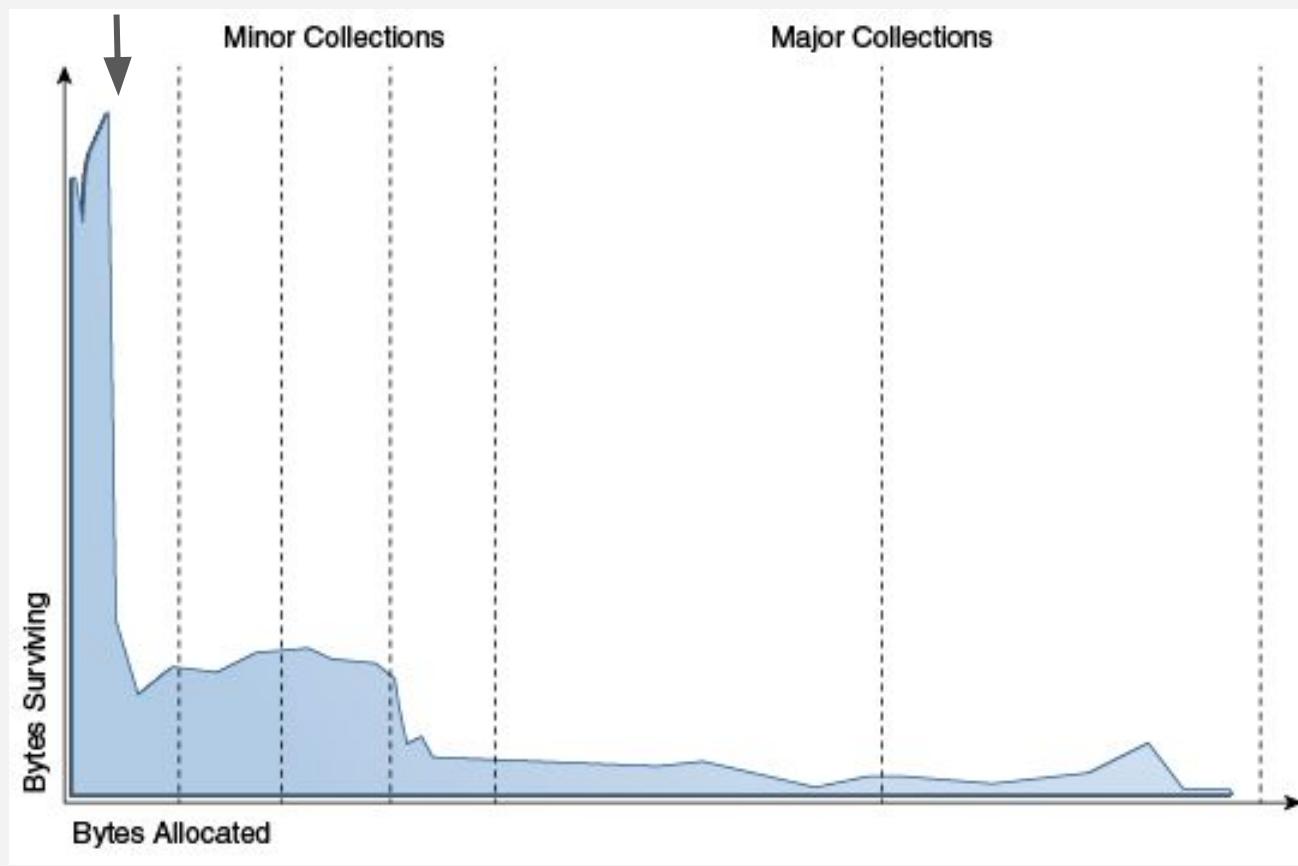
Promote



Objetos novos vivem pouco

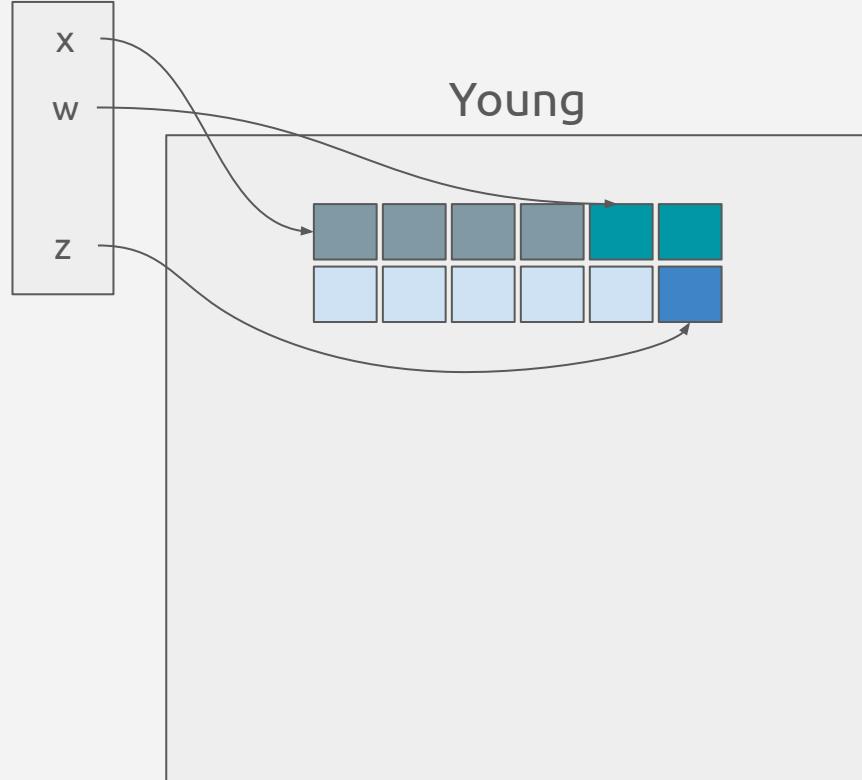


Mais removidos do que mantidos

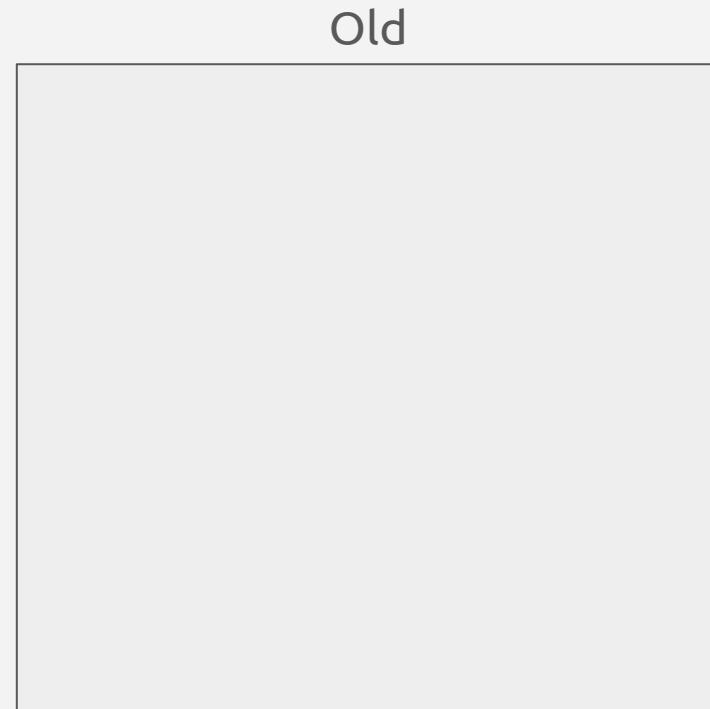


Ao invés de marcar e remover...

Roots



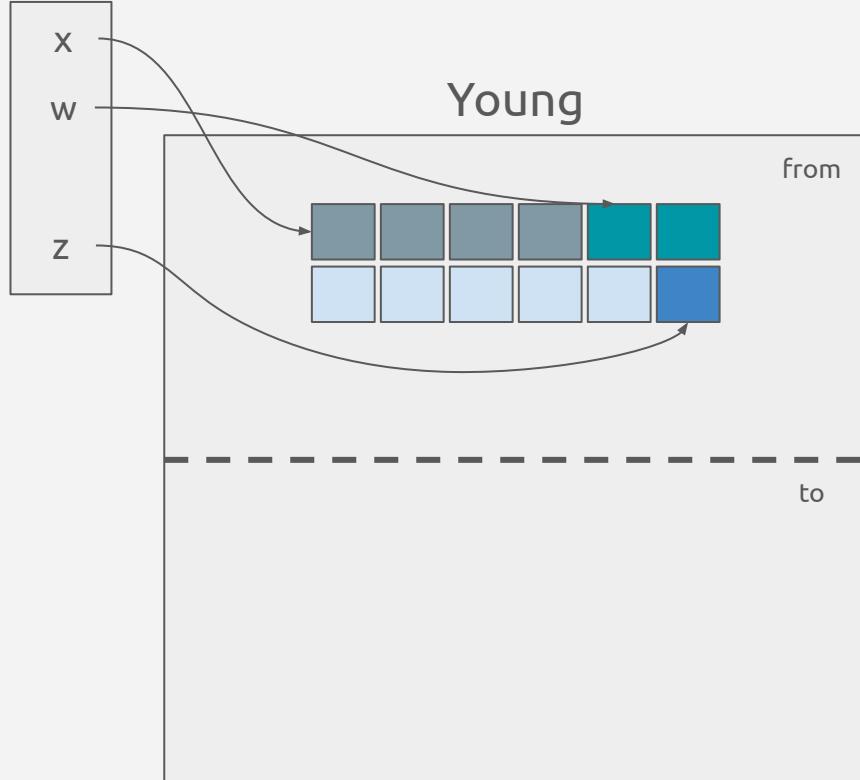
Young



Old

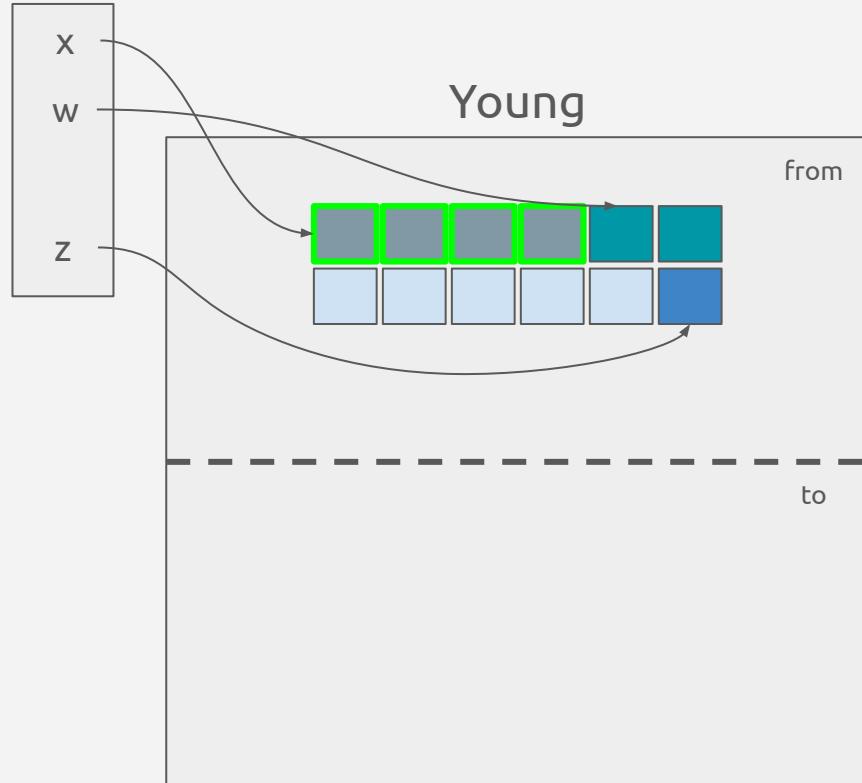
Utilizando espaço a mais

Roots



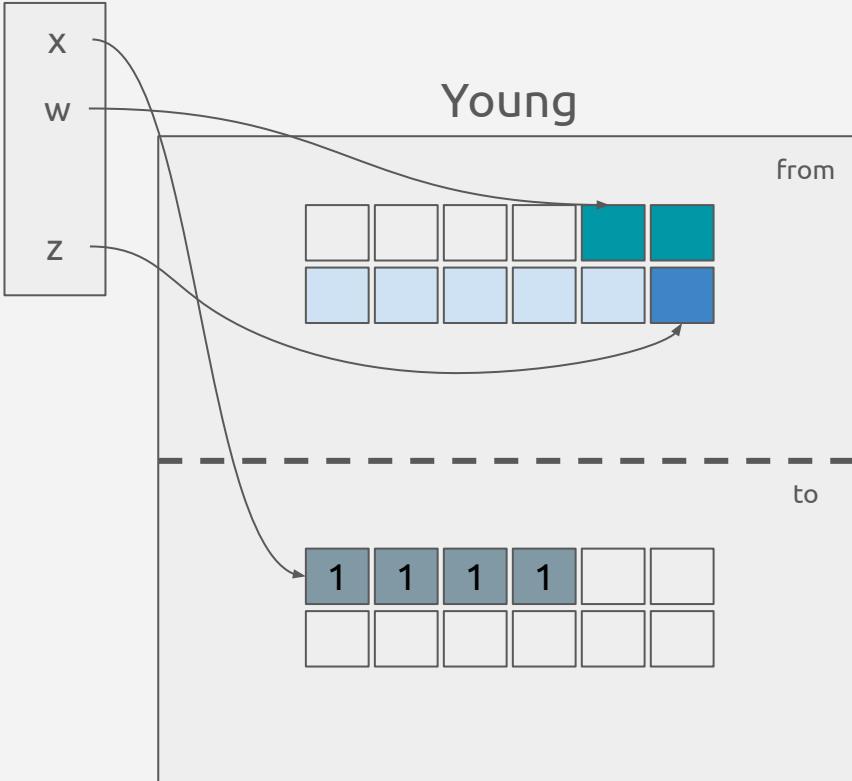
Mark

Roots



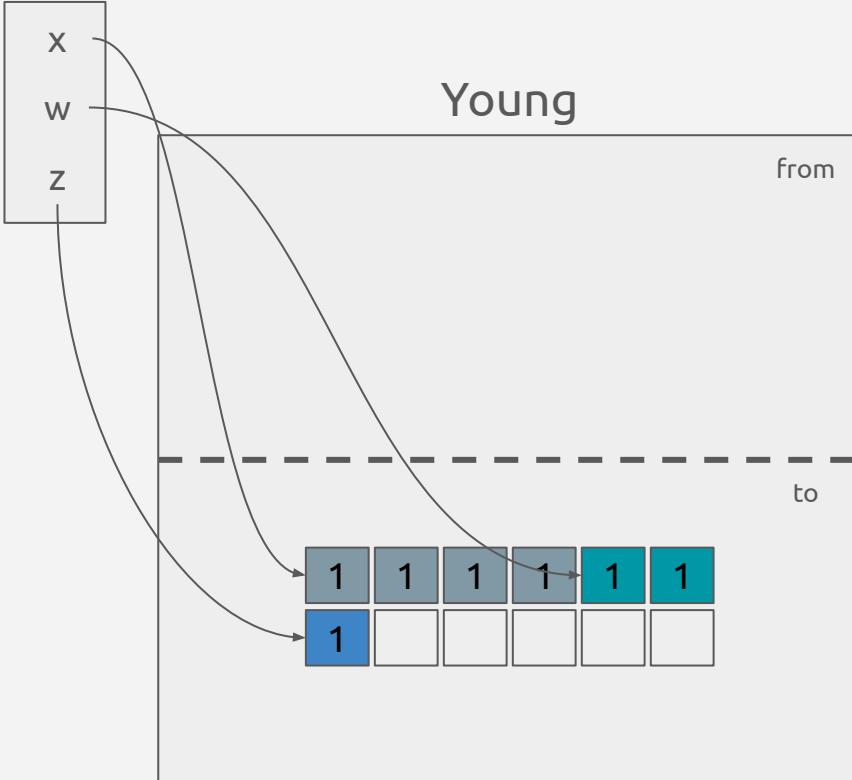
Roots

and... Copy



Mark-and-Copy

Roots



Visual GC

 Spaces Graphs Histogram

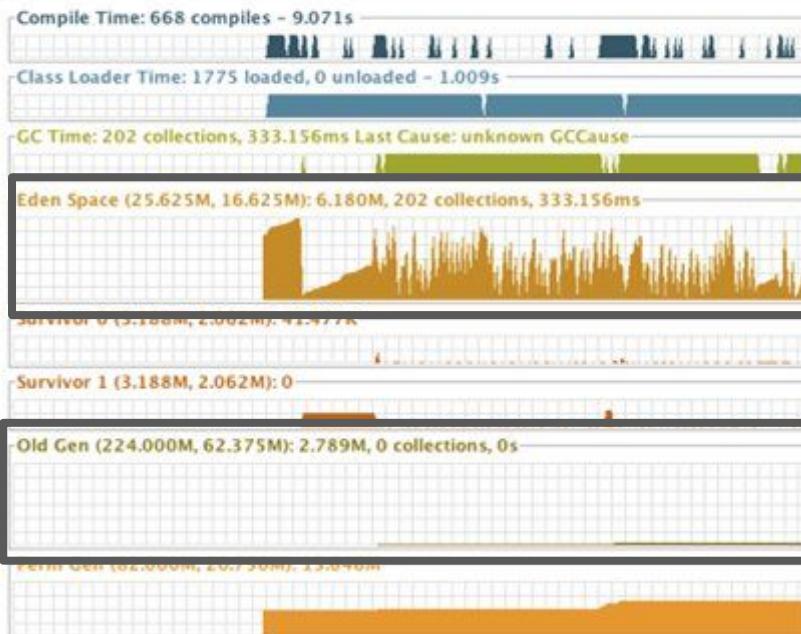
Refresh rate: Auto msec.

Spaces

Perm Old Eden



Graphs



202

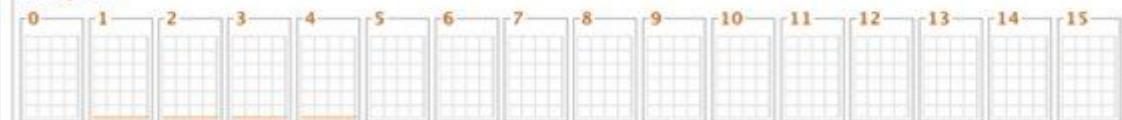
0

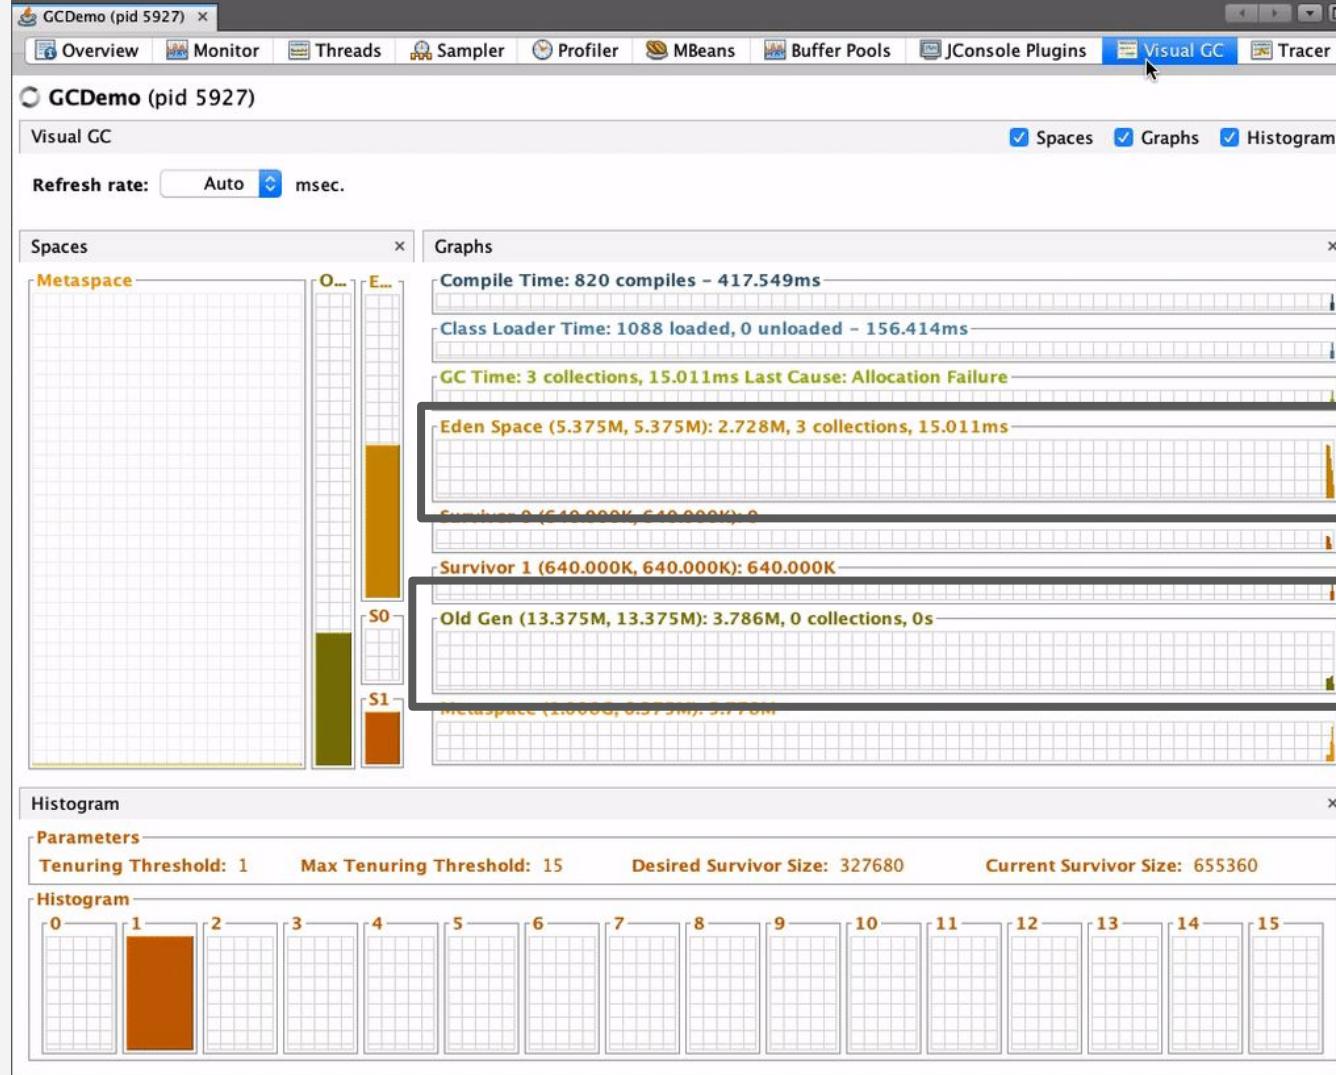
Histogram

Parameters

Tenuring Threshold: 4 Max Tenuring Threshold: 4 Desired Survivor S... 1081344 Current Survivor S... 2162688

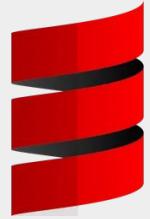
Histogram





Generational GC

JVM



CLR



JS



Mark-and-Copy (young)/Mark-and-compact (old) Generational GC

JVM

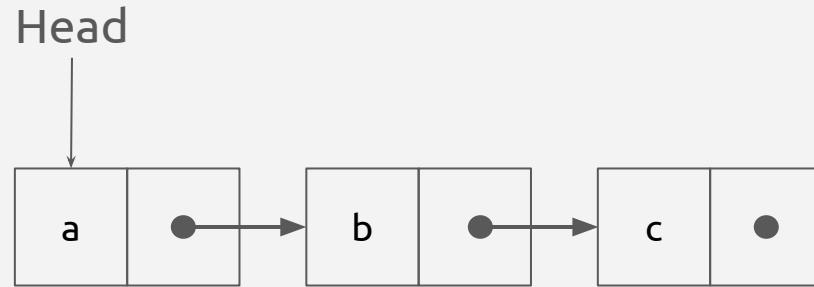


CLR

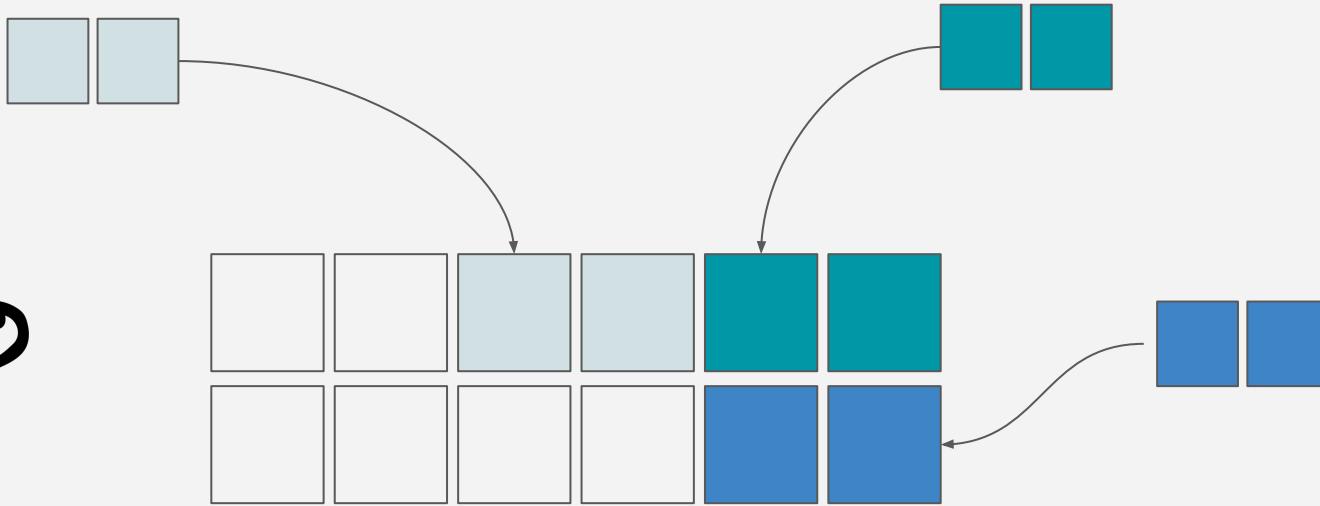


Cópias e compactação não são as únicas maneiras de lidar com fragmentação

Não havia fragmentação significativa em LISP

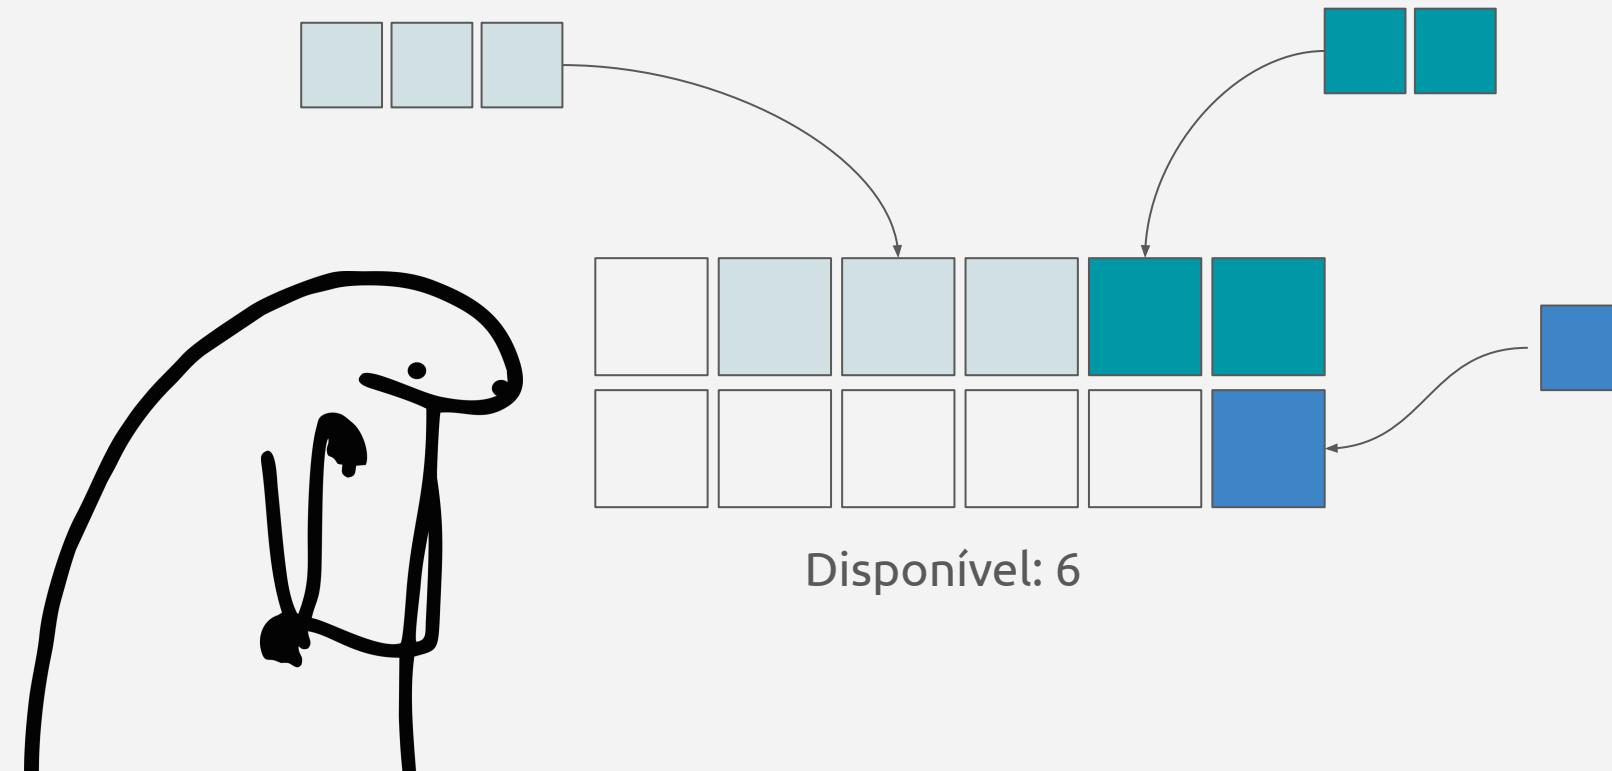


Alocações iguais == sem fragmentação



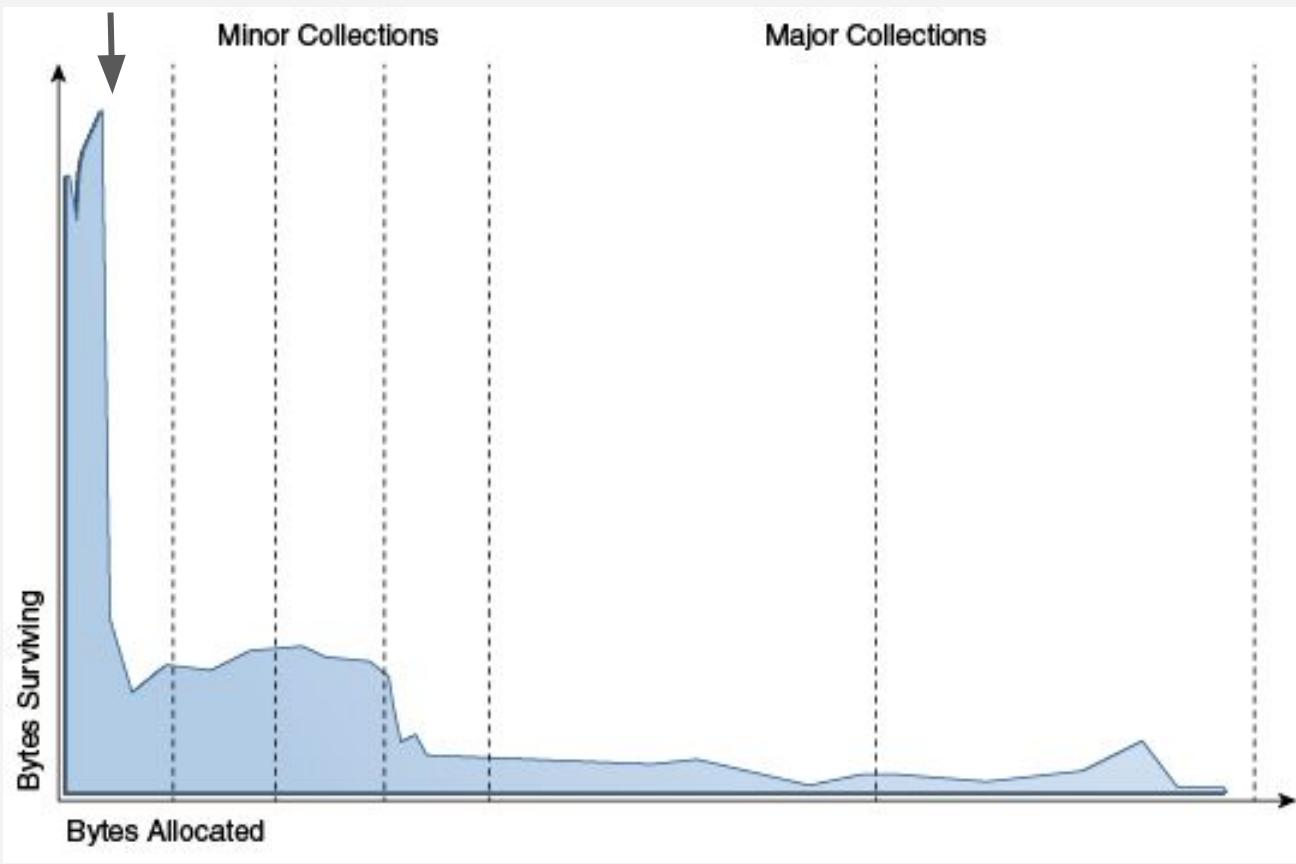
Disponível: 6

Mas o mundo real é diferente

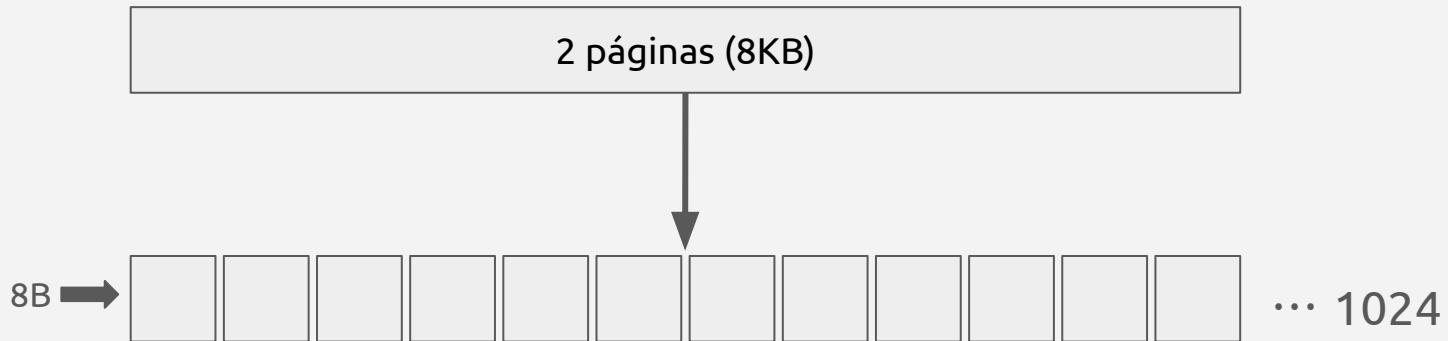


Modelo de memória do Go

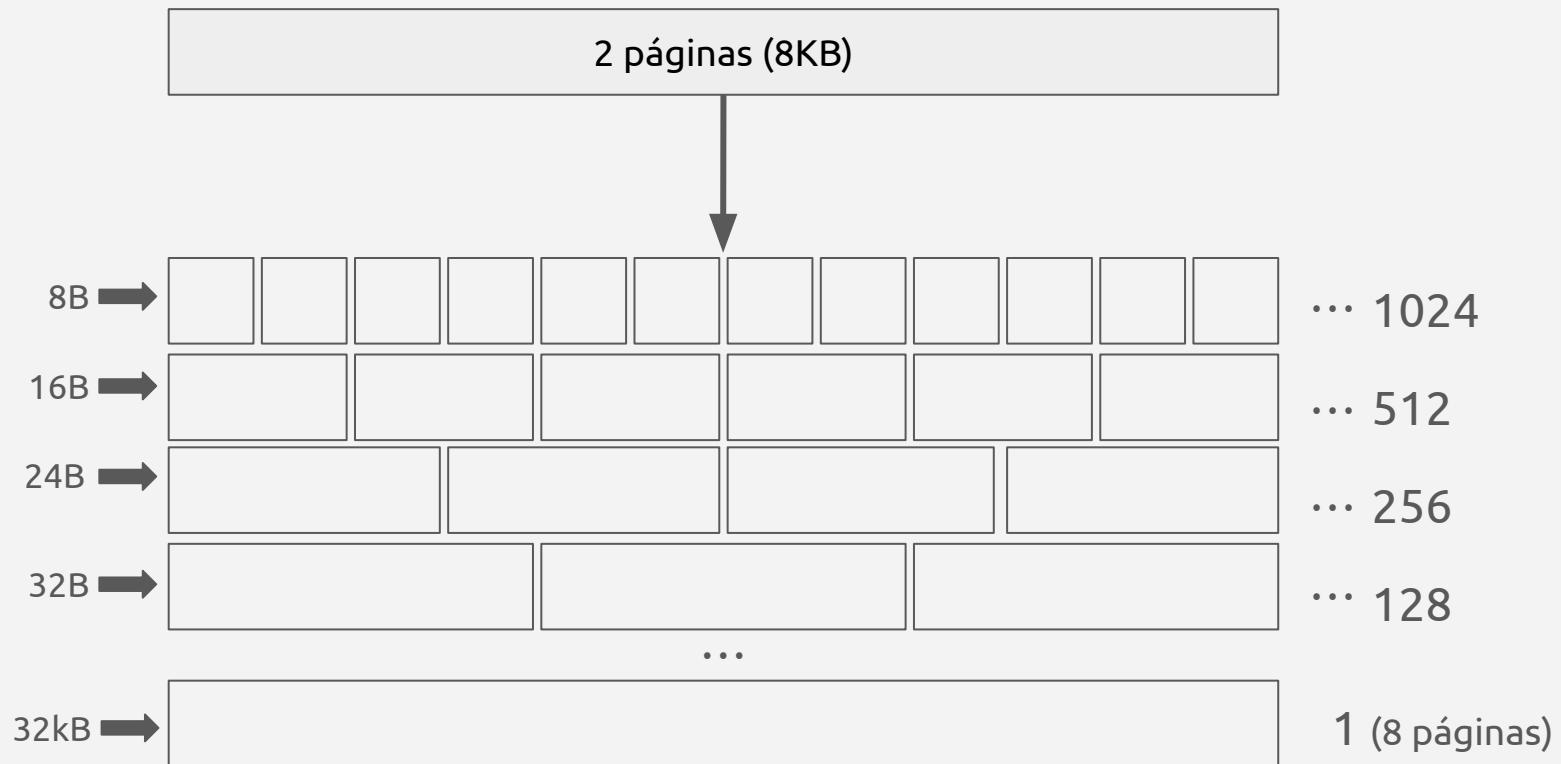
Novamente...



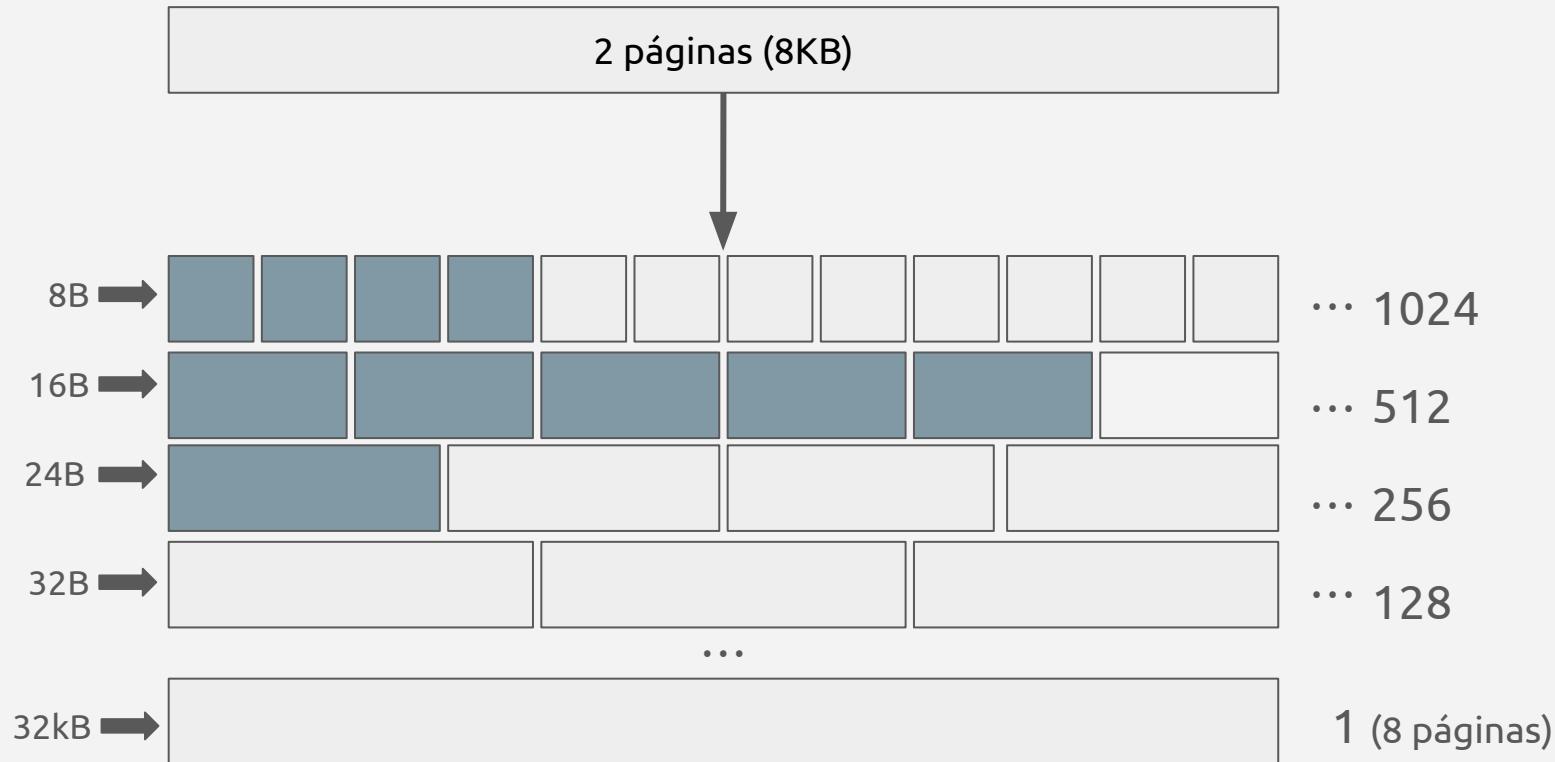
Quebra porções de memória em blocos iguais (classes)



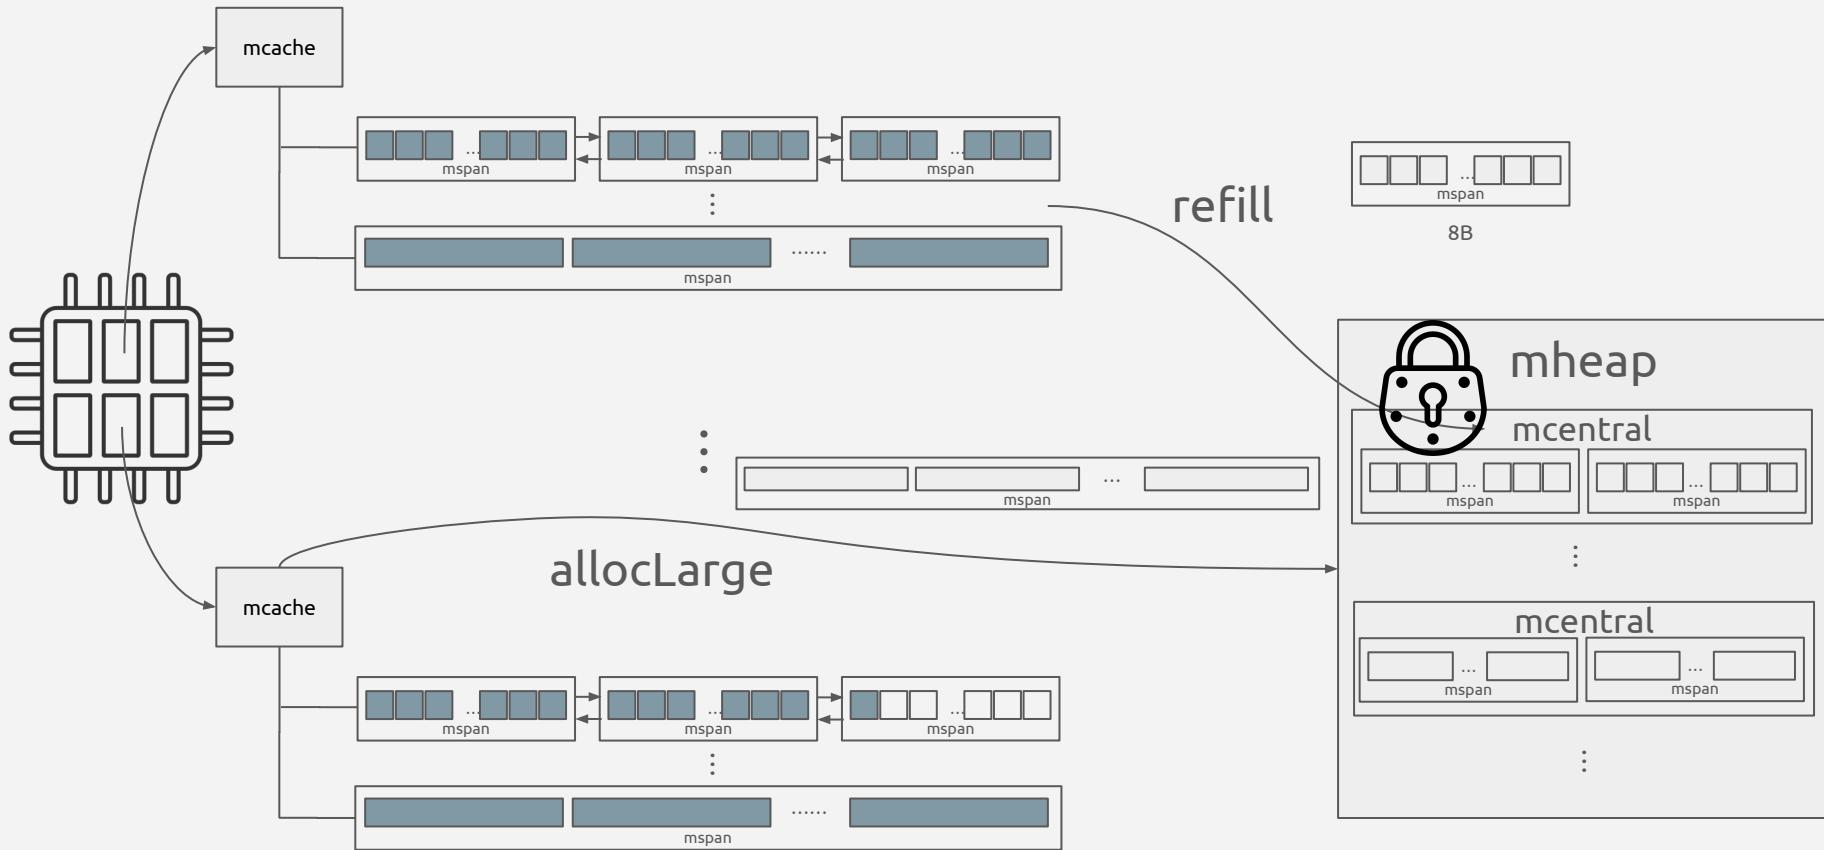
Cada classe de tamanho tem uma lista



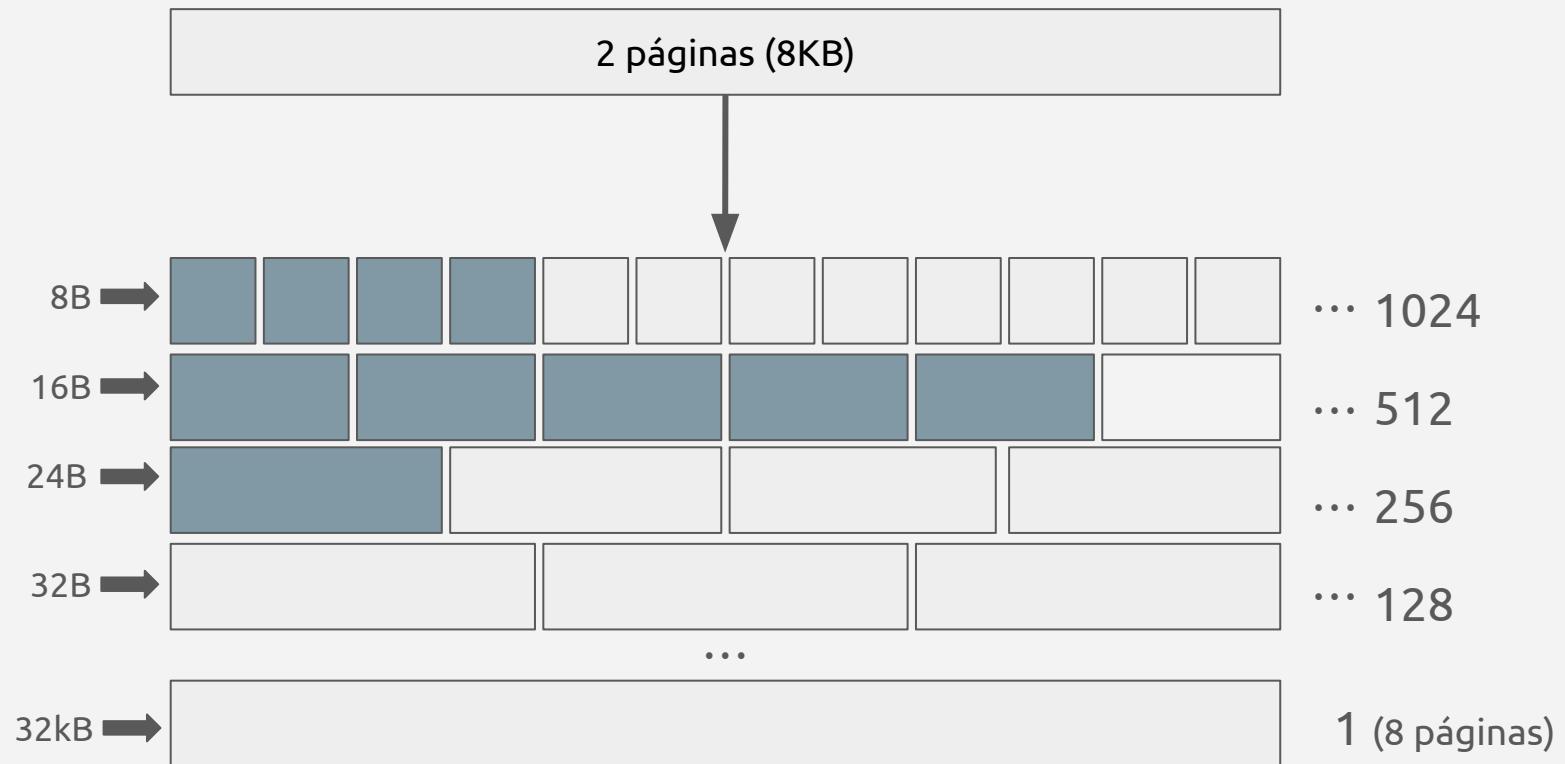
Objetos de tamanho similar são alocados e removidos juntos



Na realidade é bem mais complicado

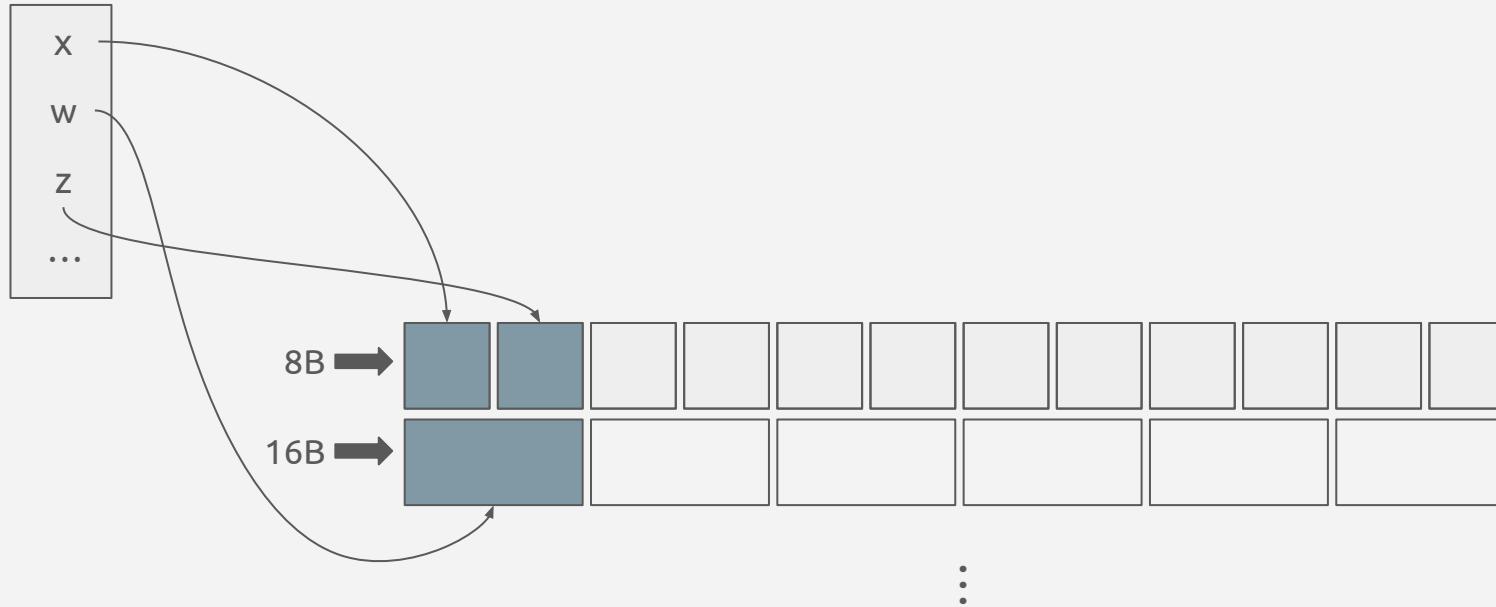


Mas a ideia original é a mesma



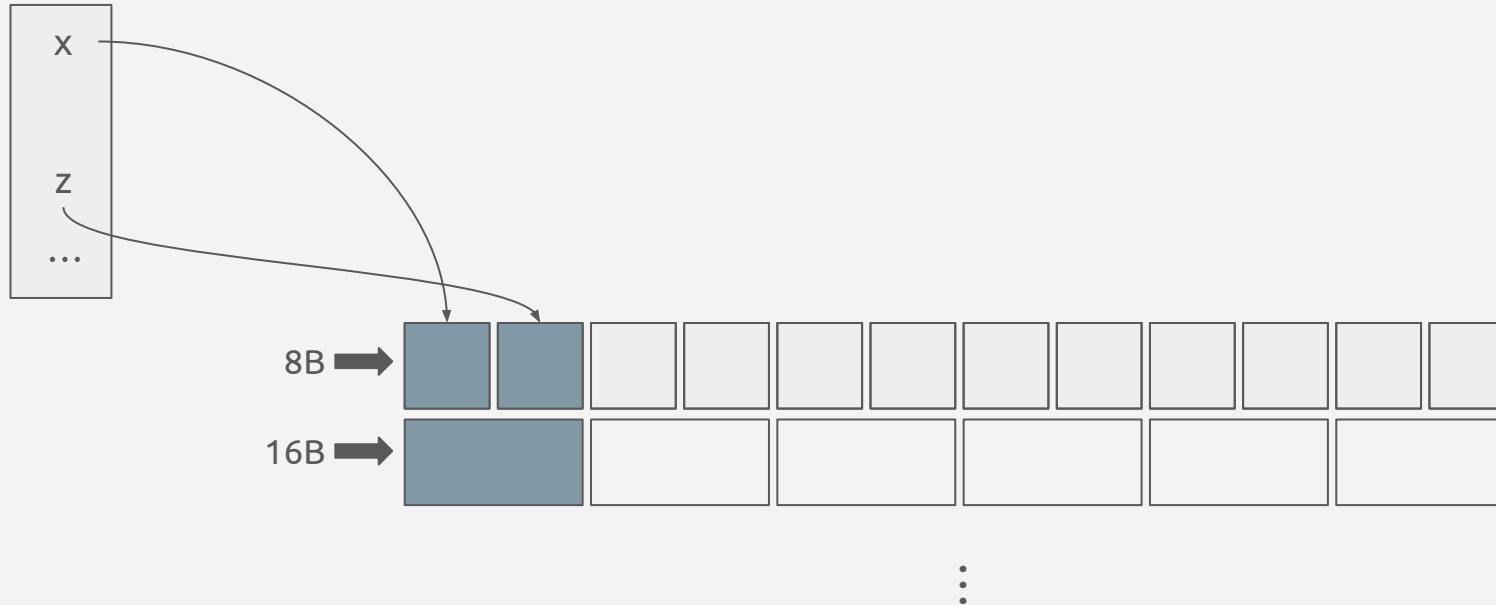
A lógica de Mark-and-Sweep é a mesma

Roots



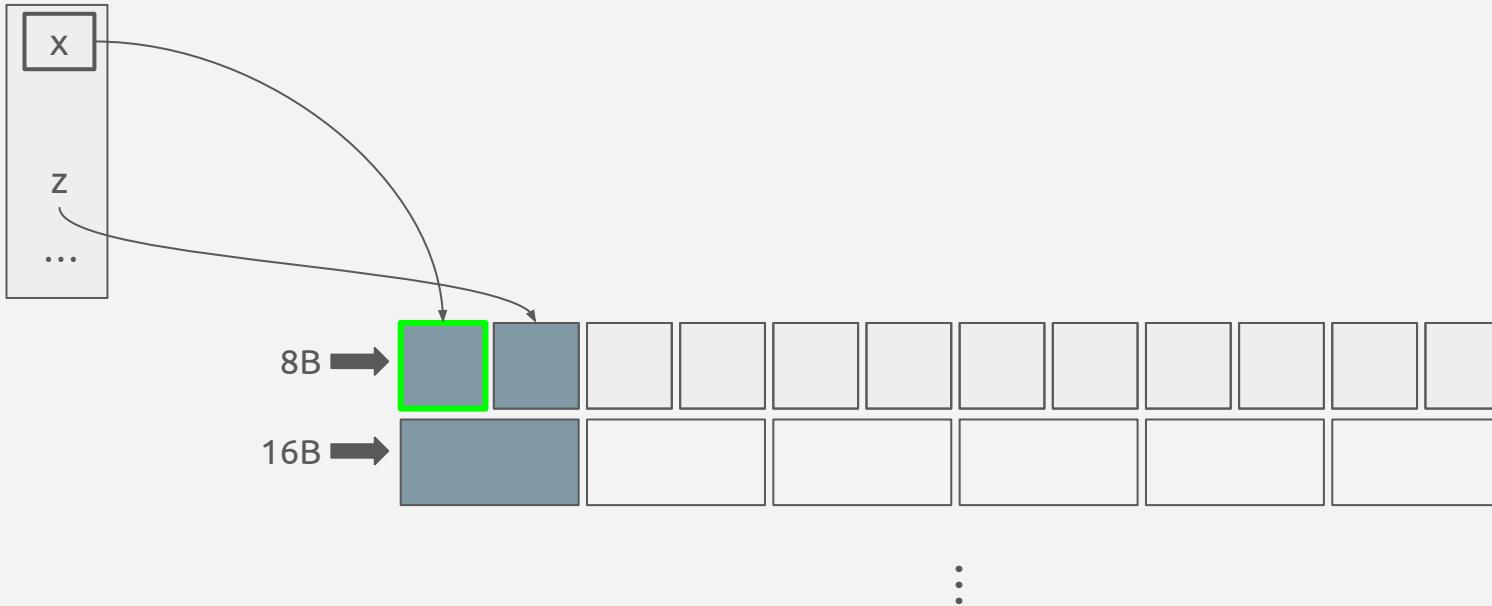
w sai de escopo

Roots



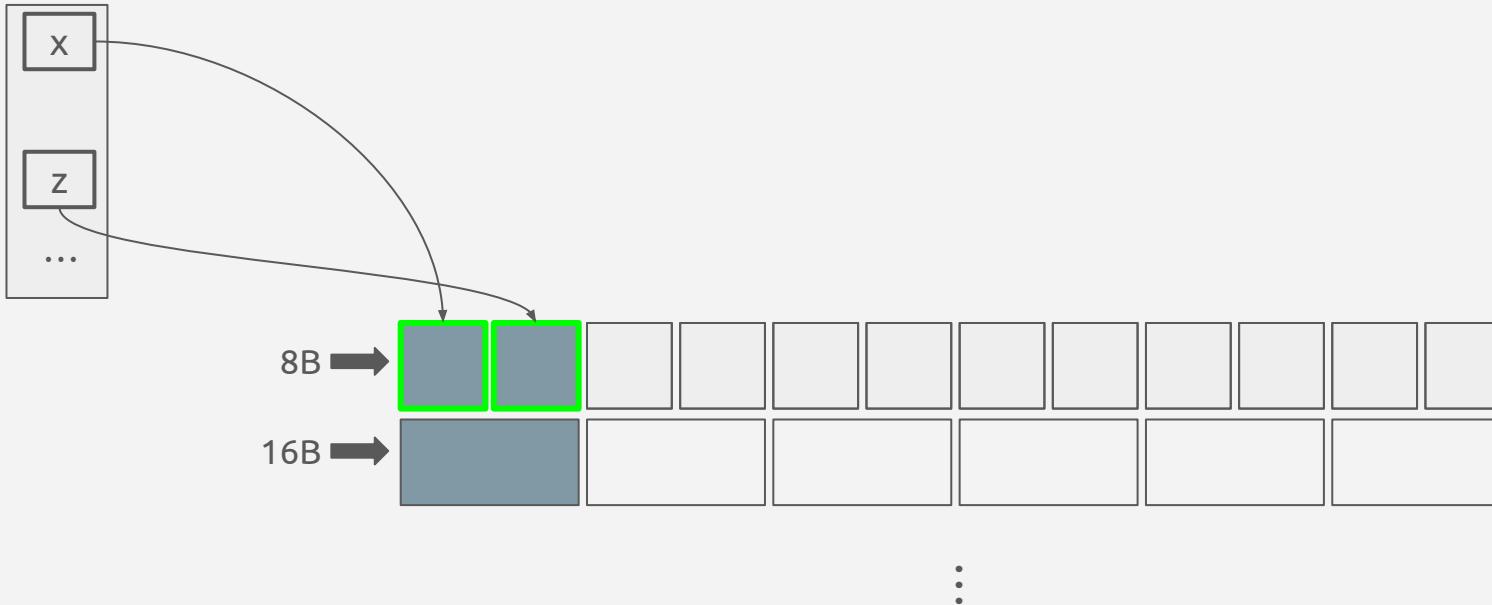
Mark

Roots



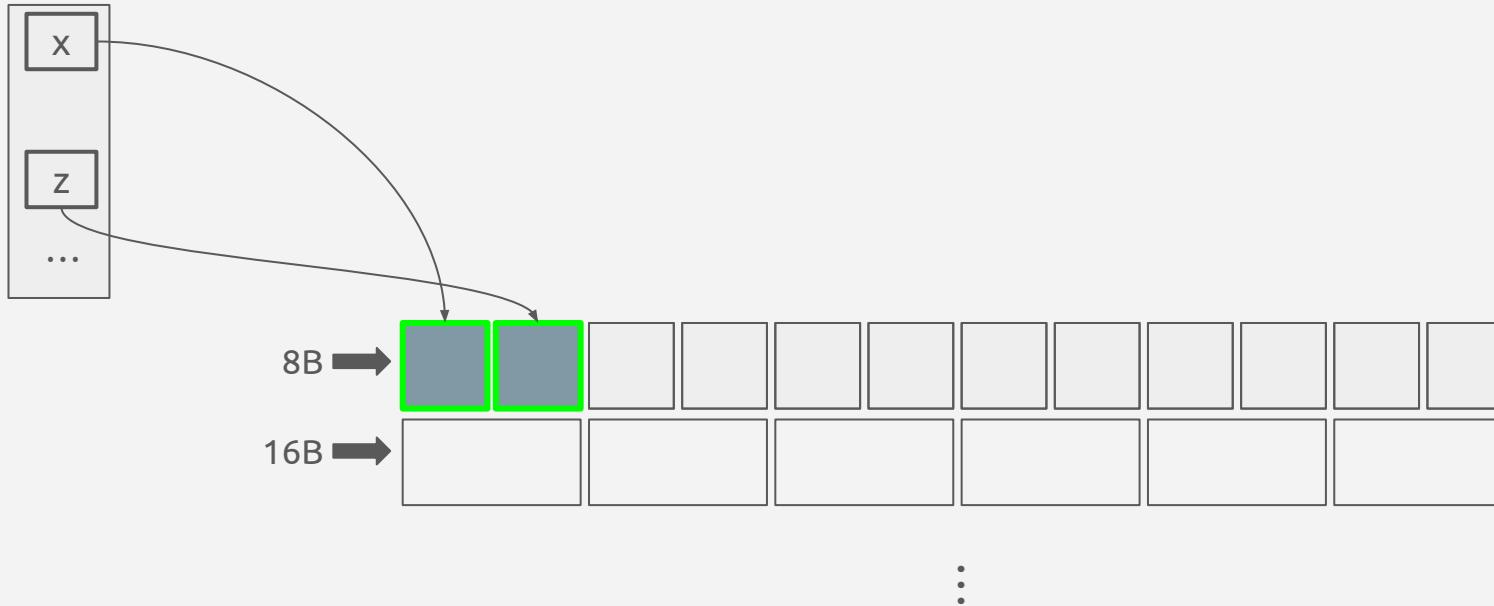
Mark

Roots



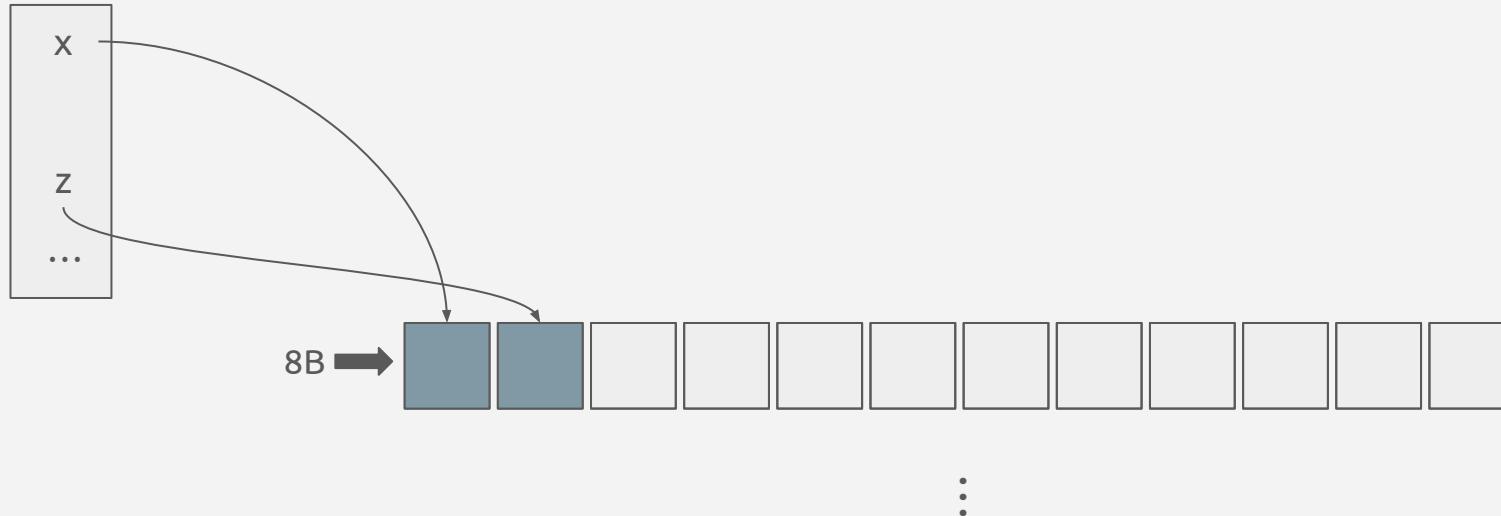
Sweep

Roots

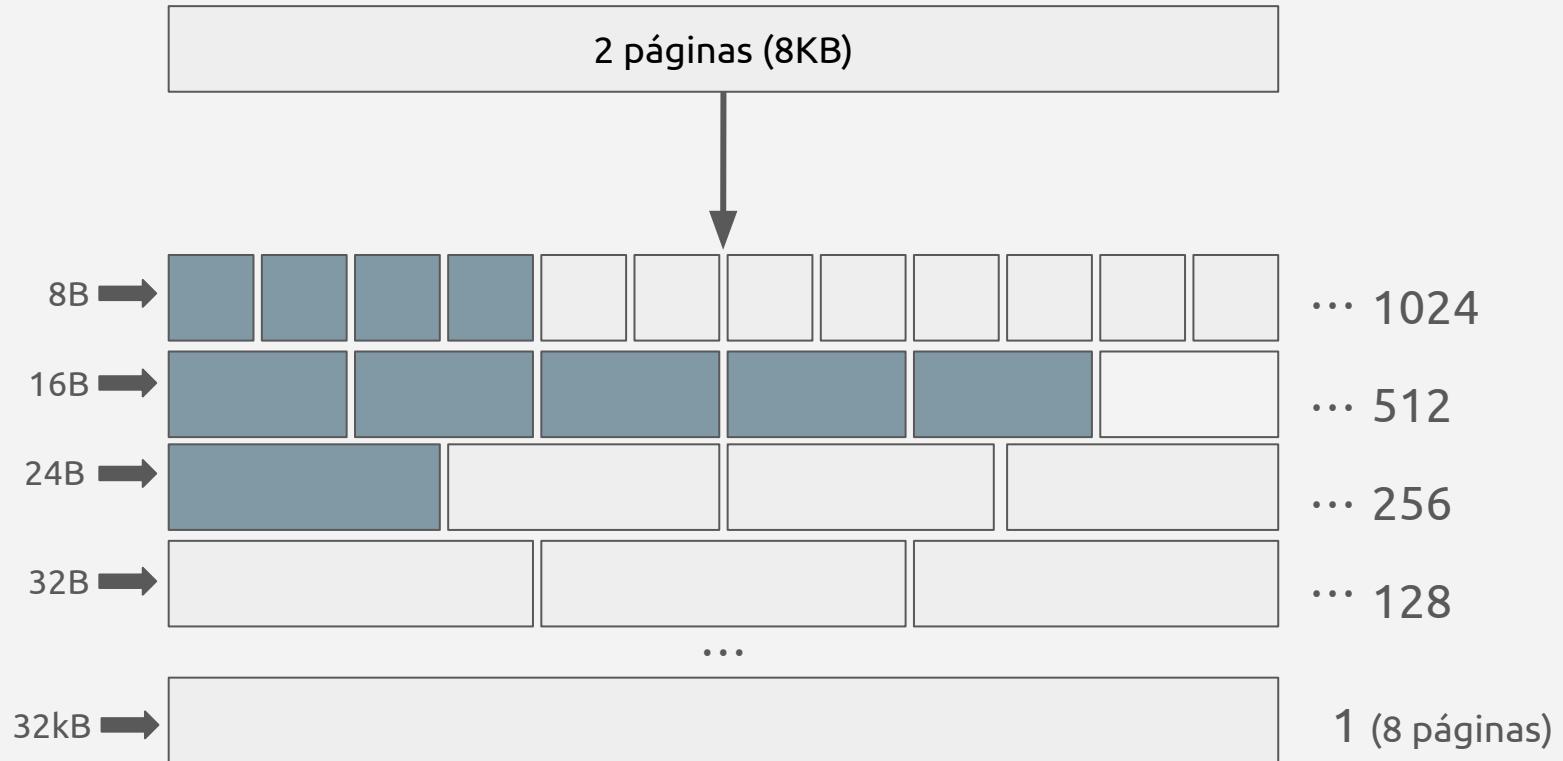


A lógica de Mark-and-Sweep é a mesma

Roots



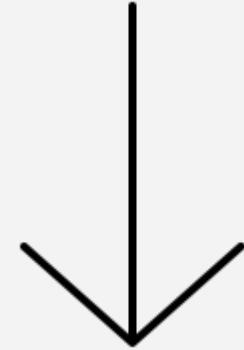
A diferença é como lidar com a fragmentação



Vantagens



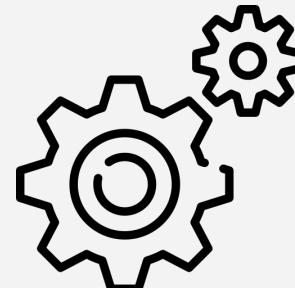
Vantagens



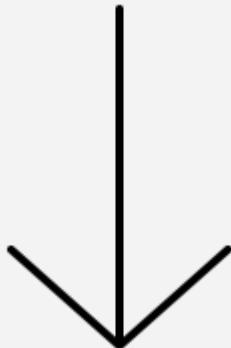
Vantagens



Latência



Otimização



Mais Fácil

Qual é melhor?



Depende

**PAPAGAIO APRENDE A DIZER
“DEPENDE” E VIRA DEV SÊNIOR**



Go vs Java

Característica	Go	Java
Latência	<1ms	~5-100ms (depende da JDK)

Go vs Java

Característica	Go	Java
Latência	<1ms	~5-100ms (depende da JDK)
Throughput (% CPU gasta em GC)	Moderado (~80-90%)	Alta (~95-98%)

Go vs Java

Característica	Go	Java
Latência	<1ms	~5-100ms (depende da JDK)
Throughput (% CPU gasta em GC)	Moderado (~80-90%)	Alta (~95-98%)
Previsibilidade (consistência)	Alta (Não geracional)	Menor (Depende da geração)

Go vs Java

Característica	Go	Java
Latência	<1ms	~5-100ms (depende da JDK)
Throughput (% CPU gasta em GC)	Moderado (~80-90%)	Alta (~95-98%)
Previsibilidade (consistência)	Alta (Não geracional)	Menor (Depende da geração)
Velocidade de alocação	Menor (Avita fragmentação na alocação)	Maior (Bump allocator)

Go vs Java

Característica	Go	Java
Latência	<1ms	~5-100ms (depende da JDK)
Throughput (% CPU gasta em GC)	Moderado (~80-90%)	Alta (~95-98%)
Previsibilidade (consistência)	Alta (Não geracional)	Menor (Depende da geração)
Velocidade de alocação	Menor (Avita fragmentação na alocação)	Maior (Bump allocator)

Os dois são concorrentes: Algoritmo Tri-Color

Otimização

Depende

Otimização

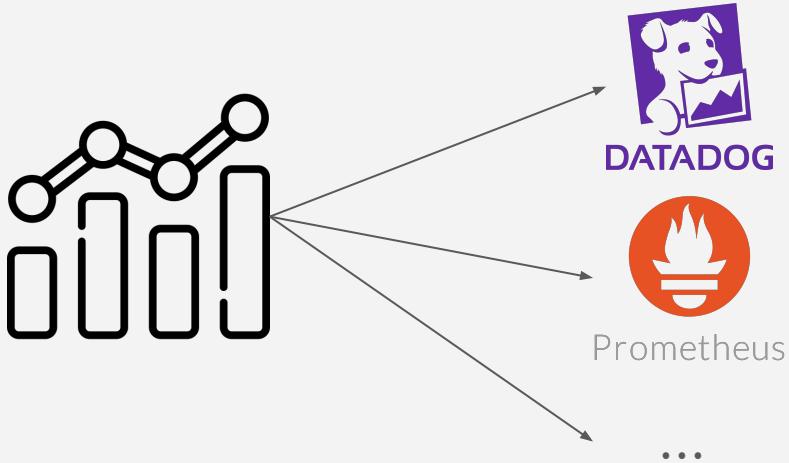
Entender como o GC funciona não é somente uma curiosidade acadêmica

Otimização

Cada combinação de workload com algoritmo de
GC gera algo diferente

Otimização

Cada combinação de workload com algoritmo de GC gera algo diferente



Como então bolsas de valores usam Java?

[B]³

Nasdaq



London
Stock Exchange Group

Existem JDKs especializadas em Real Time



London
Stock Exchange Group

azul

Zing

IBM WebSphere

WebsphereRT

JDKs comerciais

[B]³

Nasdaq



London
Stock Exchange Group



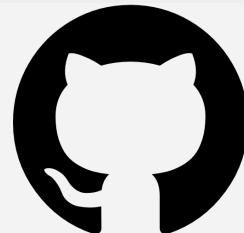
azul

Zing

IBM WebSphere

WebsphereRT

Redes sociais



[reneepc](#)



[/in/reneepc](#)

