
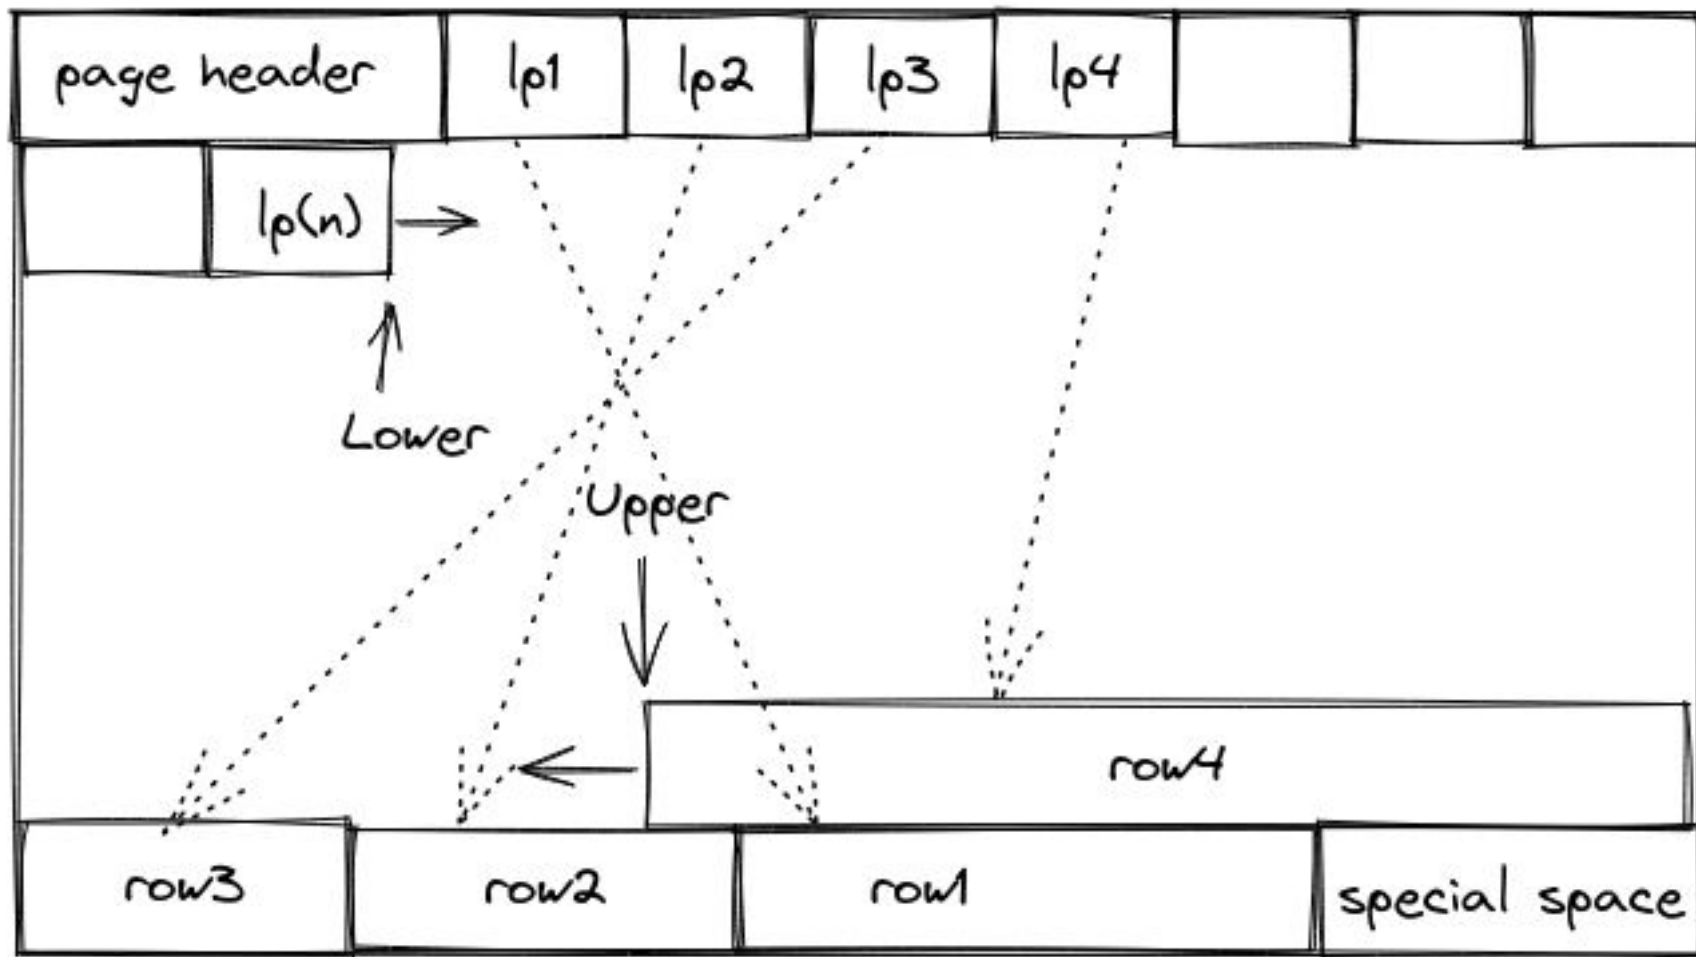


Objetivos

1. Características centrais do Scylla e como utilizá-lo
2. Replicação, consistência e queries
3. Adapter
4. Performance







CAP Theorem

Partition


Falha de comunicação
entre nós

Consistency

Nega requisição até que
recupere comunicação

Availability

Aceita requisição mas gera
inconsistência



version: "3"

services:

Cluster Name

Enter cluster name



Choose Your Cloud Provider

Where should we install your Cluster?

aws



Run Under Which Account?

Under which account should we work?



aws



\$100

Additional Services - Scylla Cloud clusters include additional instances: Scylla Manager (t3.medium) and Scylla Monitoring (m5.large)

Scylla Version

Scylla Enterprise 2022.2.6 - CQL Compatible API



Scylla Enterprise 2022.2.6 - CQL Compatible API



Scylla Alternator 2022.2.6 - DynamoDB™ compatible API

Regions

Choose your default region, you'll be able to add more regions later
(VPC must be enabled for multi-region)

Region

US East (N. Virginia)



Scylla cloud prices are billed hourly. For special pricing on an annual contract, please contact [Scylla Cloud Sales](#).



web:


driver: bridge

web:


driver: bridge

Cluster | Datacenters

DC1 - us-east-1



DC2 - sa-east-1



Cluster | Keys

PRIMARY KEY((rule_id, date), timestamp)

Partition key


Cluster key

Ordenação

Partitioner

Token


9143256562457711404



Cluster | Replication Factor


DC1 - us-east-1

RF: 3



DC2 - sa-east-1

RF: 2



RF: 5

Keyspace | Replication Factor

```
CREATE KEYSPACE swat  
WITH replication = {'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2': 2}
```



Keyspace | Create Table

```
CREATE KEYSPACE swat  
WITH replication = {'class': 'NetworkTopologyStrategy', 'DC1': 3, 'DC2': 2}
```

```
CREATE TABLE feature (  
  id uuid  
  name text,  
  internal boolean,  
  tags set<text>,  
  dependencies list<frozen<dependency>>,  
  date date,  
  created_at timestamp,  
  PRIMARY KEY((id), date, created_at))  
WITH CLUSTERING ORDER BY (date  
DESC);
```

```
CREATE TYPE dependency(  
  id uuid,  
  status_map map<text,  
  text>  
)
```




Operações | Insert

INSERT INTO feature

(id, name, internal, tags, dependencies, date, created_at)

VALUES (uuid(), "BOPE Logs", {'Bope', 'Logs'}, [{'1234': 'Insert Log'}, {'12345': 'Batch
Insert Log'}], '08:12:54', 1299038700000);




Operações | Consistency Level

```
INSERT INTO feature
(id, name, internal, tags, dependencies, date, created_at)
VALUES (uuid(), "BOPE Logs", {'Bope', 'Logs'}, [{'1234': 'Insert Log'}, {'12345': 'Batch
Insert Log'}], '08:12:54', 1299038700000);
```

CONSISTENCY ONE;
CONSISTENCY QUORUM;
CONSISTENCY ALL;

CONSISTENCY LOCAL_XXX




Operações | Select

TODA QUERY DEVE POSSUIR TODAS AS
COLUMNS DEFINIDAS NA PARTITION KEY

```
SELECT name, tags, dependencies
FROM feature
WHERE id =
'e94ef903-88ec-46d5-9739-219d02b5e8
d8' AND date <= '2023-01-23';
```


```
SELECT name, tags, dependencies
FROM feature
WHERE id =
'e94ef903-88ec-46d5-9739-219d02b5e8
d8' and name = 'Logs BOPE';
```



Operações | JSON

INSERT INTO feature

```
JSON '{  
  \"id\":  
    \"e94ef903-88ec-46d5-9739-219d02b5e8d8\",  
  \"name\": \"Logs BOPE\",  
  \"tags\": [\"Bope\", \"Logs\"],  
  \"dependencies\": [{\"1234\": \"Insert Log\"},  
    {\"12345\": \"Batch Insert Log\"},  
  \"date\": \"08:12:54\",  
  \"created_at\": 1299038700000}'
```




MVs | Materialized Views

```
CREATE MATERIALIZED VIEW internal_feature AS
SELECT id, name, tags, dependencies
WHERE id IS NOT NULL AND name IS NOT NULL AND internal = true
PRIMARY KEY (name, id);
```

id	name	tags	dependencies
...
...


CRIA OUTRA TABELA
OUTRA PARTITION KEY



Índices | Global Secondary Index

```
CREATE INDEX feature_name ON feature(name);
```


```
SELECT * FROM feature  
WHERE name = 'Logs BOPE';
```



Indices | Local Secondary Index

```
CREATE INDEX ON feature((id), name);
```


```
SELECT * FROM feature  
WHERE  
id = 'e94ef903-88ec-46d5-9739-219d02b5e8d8' AND  
name = 'Logs BOPE';
```



Driver | Token Aware Policy


```
CREATE INDEX ON feature((id), name);
```

```
SELECT * FROM feature  
WHERE  
id = 'e94ef903-88ec-46d5-9739-219d02b5e8d8' AND  
name = 'Logs BOPE';
```



Filter | Filtering


```
SELECT * FROM feature  
WHERE name = 'Logs BOPE'  
AND created_at >= 1299038700000  
ALLOW FILTERING;
```



Filter | Filtering

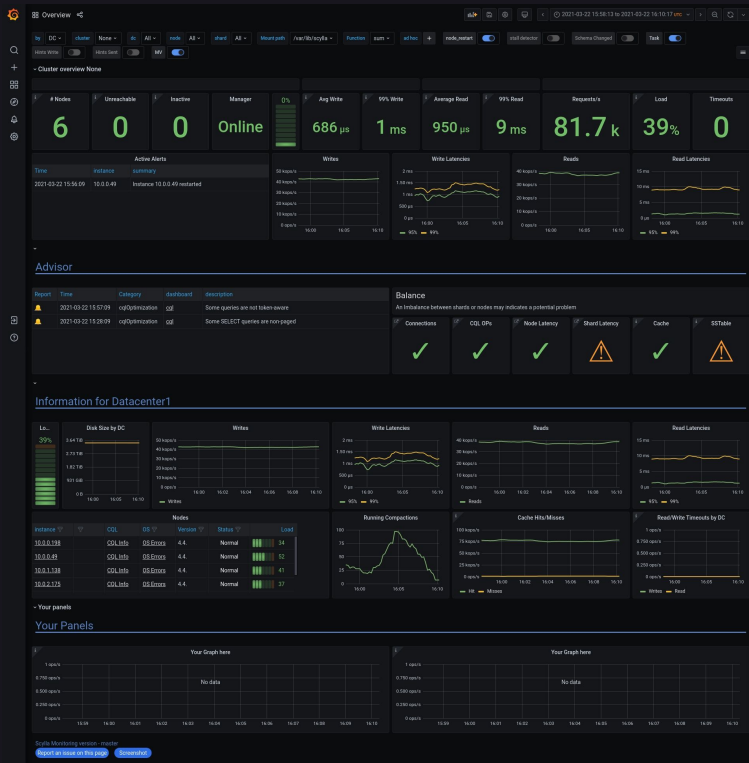
```
SELECT * FROM feature  
WHERE  
created_at >= 1299038700000  
ALLOW FILTERING;
```

UTILIZAR FILTERING
PODE SER MELHOR
QUANDO A SELETIVIDADE
É BAIXA





Monitor



SCYLLA.





Integrações



JanusGraph

Modelagem de grafos

Usando o Scylla como persistência
(Interessante para features do SWAT)



elasticsearch

Pesquisas por texto

Facilidade para front-end



kafka

Streaming de mensagens

Pode funcionar como sink ou para emitir
mensagens com mudanças do Scylla



Engenharia de Dados

Agregador de queries, batch
processing, data science e machine
learning



Adapter

Hosts
Keyspace
DefaultConsistency
MigrationDir

Adapter

Migrate()
NewTable()

```
./cql
- migration-00001.cql
- migration-00002.cql
```

TableBuilder

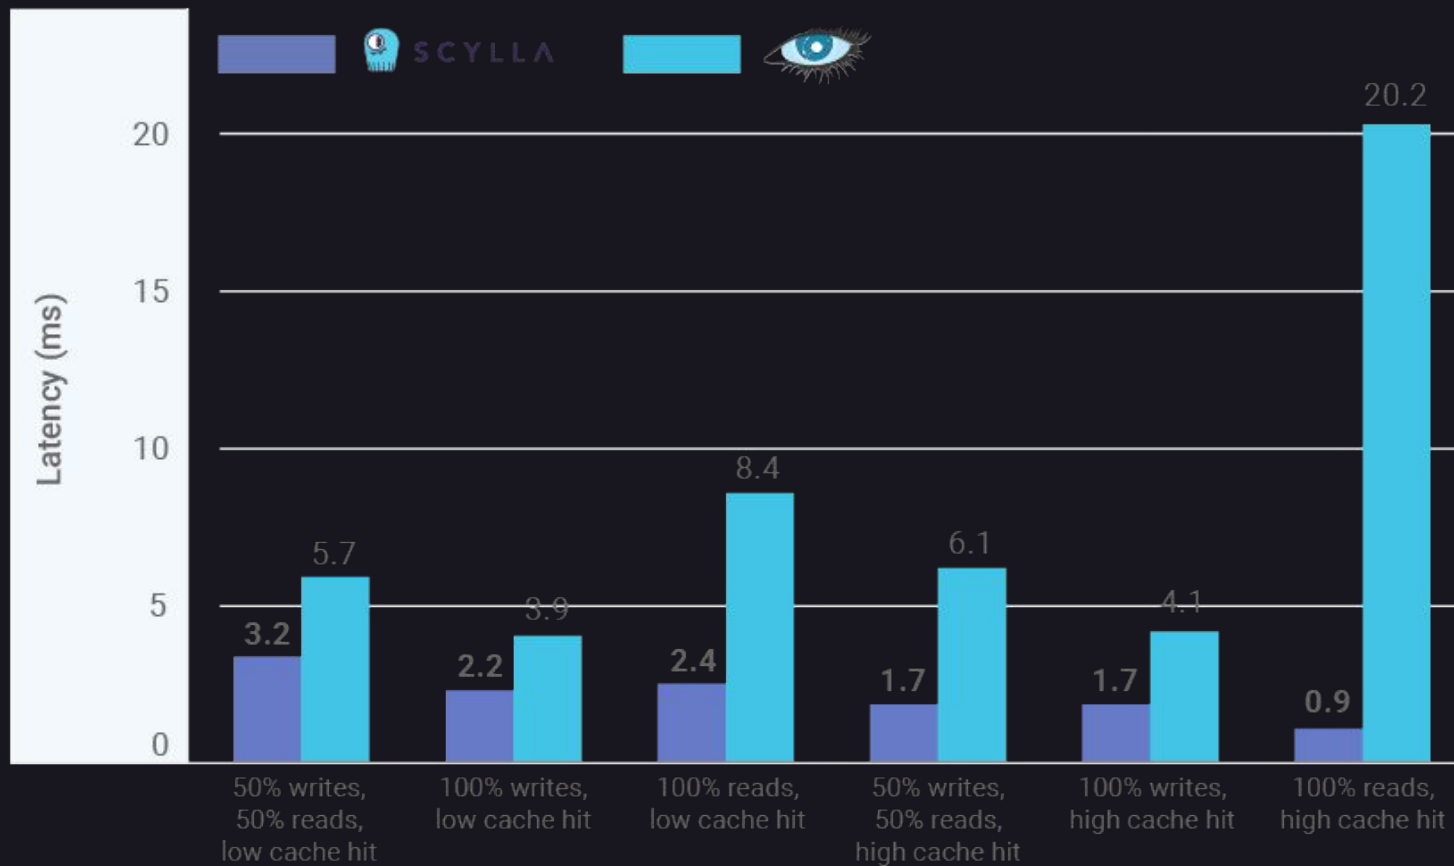
Name()
PartitionKeys()
ClusterKeys()
Columns()
Do()

Table


Insert(obj)
Update(obj)
Delete(obj)
Select(columns...)
XXXBuilder(columns...) -> gocqlx/qb

SetupScylla
(resources)

Performance | Cassandra




Scylla benchmark by Samsung




Performance | DynamoDB

3 x i3.8xlarge vs 160K write | 80K Read

Scylla vs DynamoDB update latency comparison using YCSB workloads



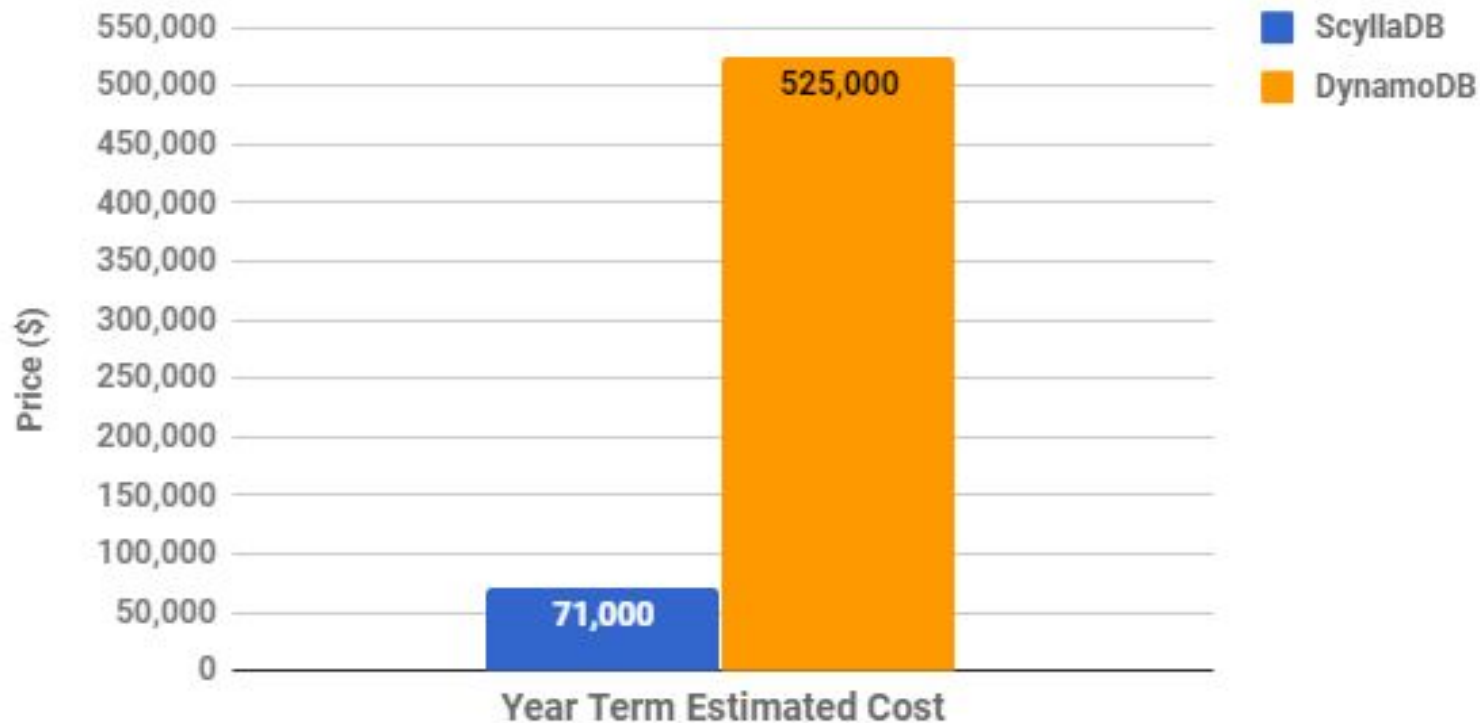
Scylla vs DynamoDB read latency comparison using YCSB workloads



Performance | DynamoDB


3 x i3.8xlarge vs 160K write | 80K Read

TCO: ScyllaDB vs. DynamoDB




Performance | MongoDB

Mongo P90 Latency (Before)




Scylla P90 Latency (After)




Performance | MongoDB

Mongo Request Rate (Before)



Scylla Request Rate (After)



MAS BÁTIMA

DE ONDE VOCÊ TIROU ESSE ESCUDO?

YCSB



Traduzido do inglês - O Yahoo! O Cloud Serving Benchmark é uma especificação de código aberto e um conjunto de programas para avaliar os recursos de recuperação e manutenção de programas de computador. Geralmente é usado para comparar o desempenho relativo dos sistemas de gerenciamento de banco de dados NoSQL. [Wikipedia \(inglês\)](#)

Ver descrição original ▼

VOCÊ TÁ MUITO ENGRAÇADINHO HEIN ROBIN



LÓGICO QUE FOI DO CU! DE ONDE MAIS?

made on imgur